CHAPTER 13

# The Anatomy of a Conversation

Communities are all about conversations.

This isn't a new idea, and it isn't ours—*The Cluetrain Manifesto* (*www.cluetrain.com*), talks about this concept extensively. It's just taken a long time for mainstream businesses to embrace it. But embrace it they have, and today's marketing organization is a veritable group hug with communities in the farthest reaches of the Web.

Having looked at some of the reasons that you need to engage in a dialogue with your audience—both on platforms that you run yourself and those run by others—it's time to turn to the conversations themselves.

All conversations consist of three distinct components: who's talking, what they're talking about, and how they're talking. The same is true of communities, and these three components can help you understand the dynamics of community monitoring.

## The Participants: Who's Talking?

Some of the community participants are members of your own organization, helping to grow and moderate it, but the vast majority of community members are outsiders.

### Internal Community Advocates

Communities thrive on activity. As with conversations, an awkward silence makes everyone uncomfortable. If you're organizing and moderating a community, it's your job to fill in the gaps in the conversation so that the other participants stick around.

Nurturing and engaging communities requires the involvement of employees from the mailroom to the boardroom. If you're running a community by yourself, such as a customer support forum or a documentation wiki, you'll need a lot of internal contributors and advocates. On the other hand, if you're just joining communities that already exist elsewhere, you won't need as much internal support, because the community's already self-sustaining and there's less need to administer or police things.

Even if you're only concerned with communities that you join, it's still wise to keep upper management apprised of what you're up to. This way, you'll surface any concerns they may have about discussing your company in a public forum.

Here are some of the internal advocates you'll need on your side.

#### Executive sponsor

Online communities rarely succeed without executive buy-in. A senior member of the company has to be convinced that the community is good for business. She needs to know the differences between traditional marketing and community marketing, or the cost benefits of self-service portals.

Many executives will pay lip service to community initiatives—after all, it's political suicide to say you don't want to talk to customers—but far fewer of them will really support the effort. There are simply too many things that could potentially go wrong with an unfettered discussion for most executives, who are notoriously risk-averse.

To get executive sponsorship, you'll have to answer one of the toughest questions of all: *why are we investing in a community?* We've given you some of the reasons, from better marketing to improved support to risk mitigation, but you need to actively communicate and get buy-in for this motivation. Many community managers go straight to community building without knowing why they're doing it. This is a big mistake. You won't know what to measure or how to define success. Don't skip this step in your community monitoring process, or you'll pay for it later when the political winds change amidst palace intrigues and someone second-guesses your work.

Once you have a genuinely supportive executive sponsor, you'll get the budget, resources, runway, and expectations to make the community flourish. In time, you'll repay your sponsor with hard evidence that the community effort is paying off—increased word-of-mouth sales, reduced call center traffic, lower hiring costs, and so on. This will give your sponsor the ammunition she needs to justify the gamble to less forward-thinking peers and marketing Luddites.

Ultimately, the executive sponsor's role is to help determine the business case for the community, to legitimize the effort, to overcome political roadblocks, and to decide which metrics and milestones will be used to determine success.

#### Administrators

Community administrators ("admins") link internal community members and external resources such as contractors or developers. They organize the structure of the community by creating topics, modifying areas, leading teams of moderators, and arbitrating disputes.

Administrators also adapt the community's terms of use when unexpected events happen, turning individual incidents like a flame war or the posting of inappropriate

content into community-wide policies and providing a platform that allows the community to grow organically.

Finally, administrators are responsible for collecting community data, either from internal platforms or from whatever tools are available on external sites, and aggregating it into reports that the executive sponsor and the community team can understand.

Any community that you moderate or run yourself needs an authoritative administrator. The administrator will have the final word on contentious disputes that will inevitably arise. This person needs to be level-headed and professional, and must be able to anticipate flame wars and settle issues quickly, consistently, and decisively.

#### Moderators

Community moderators ("mods") are responsible for specific sections of the community. They read, maintain, and moderate platforms that are often related to their own interests or their roles in the organization. They answer basic questions that community members have, and make sure that new participants are welcomed.

Moderators' interactions happen on a public platform, so their work helping others is searchable. This can dramatically lower the support volume of your organization, because repeated questions can be answered with a search, which lets support scale more effectively.

Moderators also provide rough quantitative and qualitative measurements to administrators. They are your first line of defense against spam and trolls, and act as an early warning system for potential problems. Since community moderation is extremely subjective, mods are intimately familiar with your terms of service and are able to enforce them gently before escalating.

### Don't Feed the Trolls

A troll is a community member who contributes controversial, inflammatory, or off-topic content to a community in the hopes of getting attention or provoking reactions. Trolls are one of the most difficult aspects of community management, since any response to their behavior is precisely the effect they're hoping to have.

In certain cases, legitimate and established community members can be seen trolling other users to get reactions out of them, but these incidents are often short-lived, done in jest, and superseded by the overall quality of content that these members provide. In other words, the more valuable a member is to the community, the more likely he will be able to troll without community backlash.

Wikipedia has dealt surprisingly effectively with trolls, considering the amount of subjective information on the site. Its primary rule, *do not feed the trolls*, simply states that trolling behavior should not be acknowledged. Some other techniques Wikipedia suggests include deleting offensive or illegal content quickly, enlisting the community's help to report trolling, and not undoing a troll's work immediately (provided that it's

not illegal) so that the troll loses interest. You can read more about Wikipedia's trolling policy at *http://meta.wikimedia.org/wiki/What_is_a_troll%3F#Dealing_with_trolls*.

Some social news aggregators hide responses that are rated below a certain level from everyone but the poster, so trolls don't know they're being ignored. Boing Boing has another approach: disemvoweling, in which offensive comments have their vowels removed, neutering them without deleting them entirely. Check out Cory Doctorow's web article, "How to Keep Hostile Jerks from Taking Over Your Online Community" at *www.informationweek.com/shared/printableArticle.jhtml?articleID=199600005*.

If you're running your own community platform, invite your best, most enthusiastic customers to help you run things. Strike a healthy balance between in-house moderators (employees who have access to moderation tools and administrative controls) and external moderators (third-party volunteers who care and are considered more "genuine").

#### Subject matter experts

Subject matter experts are occasional contributors who provide specialized knowledge, such as lists of frequently asked questions and answers, snippets of code, recipes, or technical configurations. They're effectively seeding this information into the community. Often, subject matter experts are product managers or certified professionals, such as doctors, who have in-depth experience dealing with the topics around which the community is organized, but have to be hounded for content or participation.

## External Community Members

The community participants who really matter are those from outside. They're your content producers and consumers, your performers and your audience, your feedback, and your friends and your critics. They have different levels of involvement with your community and want different things from it.

There are many ways to classify the external members of your community. One approach is to look at their online behaviors. While many community members consume content, fewer share it, fewer still comment on it, and only a small portion produce or maintain it. Figure 13-1 shows one model of community participants and the roles they play.

There are many other ways to segment online contributors. Forrester Research, for example, splits them into creators, critics, collectors, joiners, spectators, and inactives, according to the actions they perform online. For the purpose of monitoring and responding to community members, we've separated them into six distinct groups according to their rates of contribution and engagement:

- The hyperinvolved, prolific few (what we call the *long tail of freaks*, or LTOF)
- The totally engaged, but less pathological, *fans and contributors*

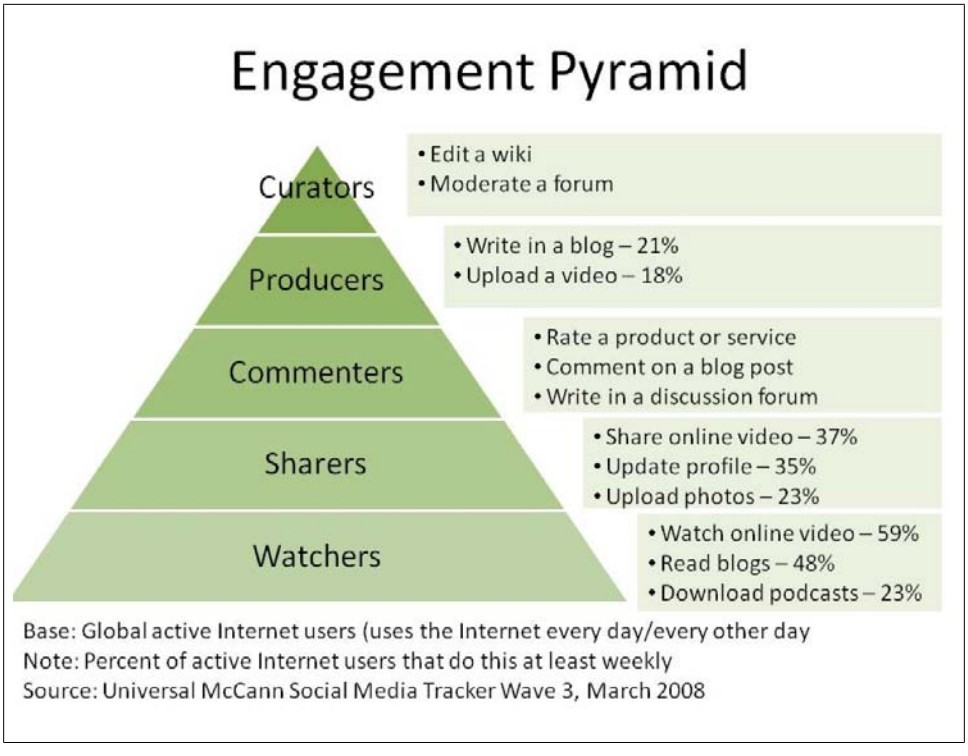


*Figure 13-1. The Altimeter Group's model of community participants based on engagement; Altimeter's model counts community members by behavioral segments in the Pew Internet and American Life survey on Internet usage, where available*

- *Bursty users*, who vary significantly in their contribution rates
- *Occasional participants*, who only interact in small ways
- *Lurkers*, who watch from the sidelines
- The *disengaged*, who you count as part of the community, but who no longer consume nor contribute content

To understand how we've separated these groups, let's look at a fundamental concept of information systems: power laws.

#### Power laws and community members

Most things in the information world follow a power law. A power law simply means that a system has a few very common things, and many, many rare ones. One power law you've probably heard of is the 80/20 rule (known more formally as the *Pareto distribution*), which says, for example, that 20 percent of people control 80 percent of the wealth.

Power laws show up in everything from bookstores (a few popular books sell millions, while many niche publications sell only a few copies) to last names (there are a few very common names and many rare names, as shown in Figure 13-2). On the Web, a small number of websites have nearly all of the Internet traffic, while there are millions of seldom-visited ones.

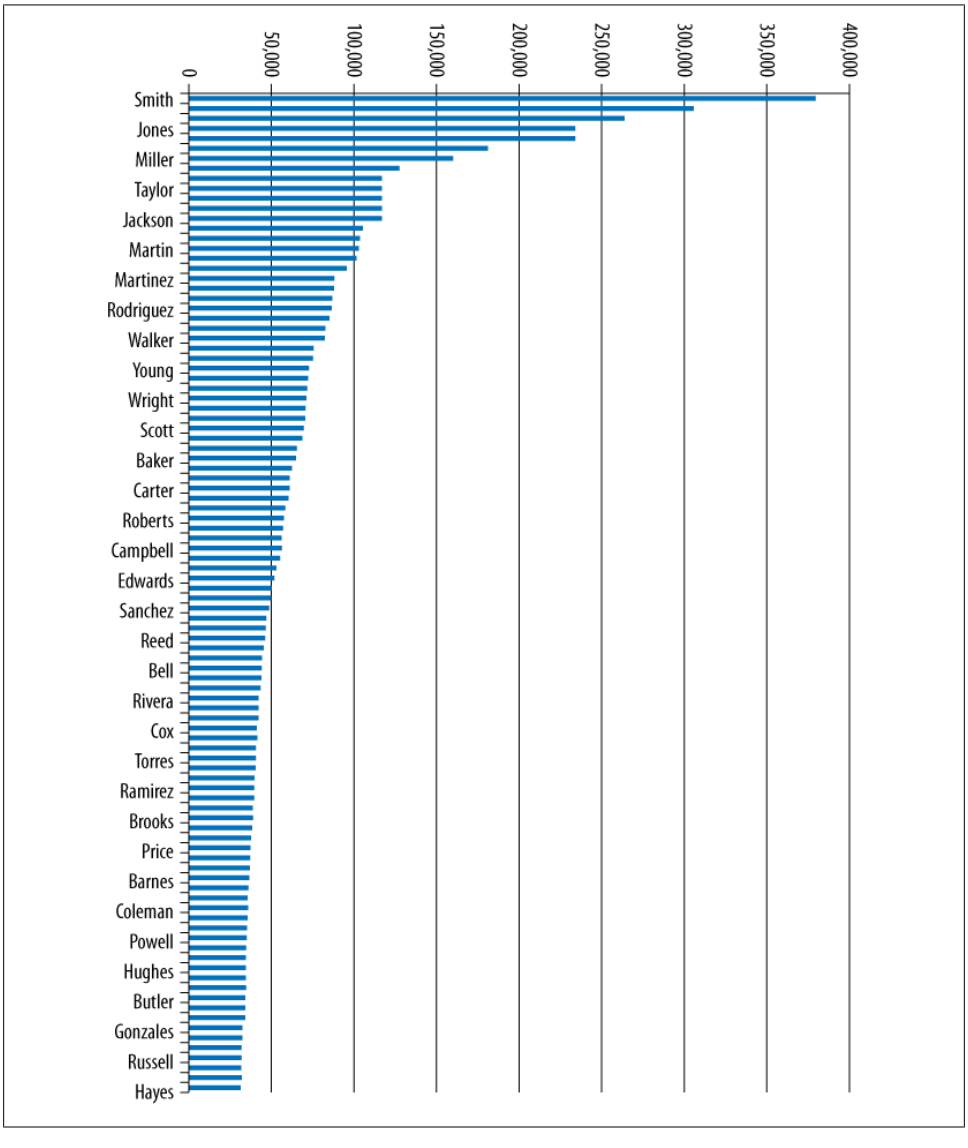


*Figure 13-2. A graph of last names by popularity shows a characteristic power law distribution*

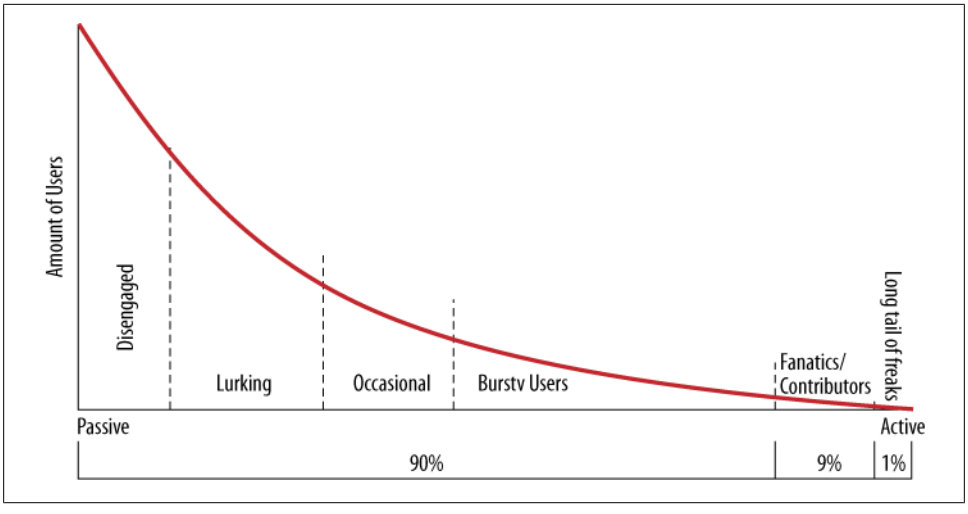


*Figure 13-3. An analysis of community members by level of activity shows many disengaged members and a long tail of hyperactive participants*

The long tail of power laws has been the subject of much discussion on the Internet, as information systems make it easier to find and analyze them (see "Power Laws, Weblogs and Inequality" at *www.shirky.com/writings/powerlaw_weblog.html* for an excellent discussion of web content and power laws in action). One place they turn up is within communities.

Many of your community members will be completely disengaged, having participated once and forgotten you (don't take it personally—we still love you). Some will be lurking, and still others will occasionally participate in discussions. You'll see one group of members become very active for brief periods of time, often on subjects they care a lot about. About 10 percent of your visitors will be heavy contributors, while a small subset will be so overactive they'll need to be politely but firmly restrained.

These numbers seem to be universal. Wikipedia has an extremely small number of active contributors—75,000, which is just over 0.01 percent of all its visitors—yet that tiny group has created over 10 million articles. Jakob Nielsen's research supports this distribution: "In most online communities, 90% of users are lurkers who never contribute, 9% of users contribute a little, and 1% of users account for almost all the action."

Figure 13-3 shows the relative distribution of these people across a community.

If you flip the axes (as we've done in Figure 13-4) and look at the volume of contributions from these users, nearly all content comes from the fanatics, freaks, and contributors. Think we're exaggerating? Amazon.com's top reviewer at the end of 2008, Harriett Klausner, had contributed over 17,900 book reviews.

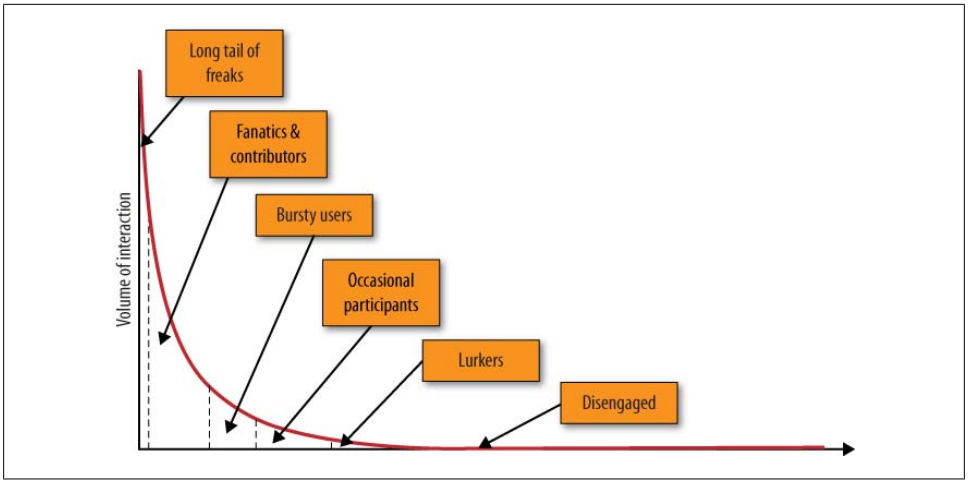


*Figure 13-4. A histogram showing the number of posts by a contributor reveals that a very small number of participants are responsible for nearly all of the content in a community*

We're not implying that Ms. Klausner is a freak—her reviews have been rated "helpful" by over 70,000 customers. But to read 17,900 books, a 70-year old person would have to finish a book every 34.2 hours since birth, which seems a bit out of the ordinary. Amazon, perhaps reacting to this and other users who reviewed in volumes, changed their practices to rank based on the percentage of positive or "helpful" votes by Amazon.com users. For more information, see *http://en.wikipedia.org/wiki/Harriet_Klausner*.

The boundaries of the community change contribution rate. On less hierarchical platforms like Twitter, for example, there are millions of *emergent communities* that might coalesce around a hashtag or a group of friends. But there's still a power law at work: a few celebrities and highly followed personalities account for most of the content that's repeated and amplified.

For communities that you run, your goal is to try to get as many members of the community as possible into the "regular contributors" category by engaging the disengaged, coaxing the lurkers out of hiding, discussing topics that the occasional participants feel compelled to join in on, and dealing gently but firmly with the freaks and abusers.

See Ross Mayfield's web article "Power Law of Participation" (*http://ross.typepad.com/blog/2006/04/power_law_of_pa.html*) for a discussion of contribution across population. Mayfield references the *Cornucopia of the Commons*, Dan Bricklin's term for open systems that grow because of their communities' contributions (*www.bricklin.com/cornucopia.htm*).

### Long tail of freaks

The LTOF is made up of the few users who will do their best to be active within the community, *despite any barriers to entry, no matter how high*. These members may create, modify, and correct content, even to the point of undermining other contributors. They will report bugs and suggest feedback often and liberally to anyone who will listen.

The LTOF is a mixed blessing. They'll start to derive a sense of self-worth from your community, and they may be seen as authorities by their peers, but they can also make newcomers feel unwelcome and challenge the authority of legitimate administrators. On Wikipedia, for example, debates constantly rage between the core team of editors—who feel personally responsible for every entry on the site—and less fanatical writers, whose occasional entries are often removed by those editors, discouraging them from becoming more engaged.

Dealing tactfully with the LTOF is a big part of a moderator's job. These members want attention, praise, and sometimes rewards, but they can also be brittle and fickle, and will air their grievances openly. When they do, their familiarity with the community makes them hard to silence. They're one of the main reasons, outside of spam, that a well-understood and decisively enforced terms of service agreement is vital to the well-being of any community.

The first line of defense against the LTOF is to monitor for them. Set thresholds for rates of contribution, number of votes in a visit, volume of posts, and number of links in a message, for example, and you'll quickly identify the members who exhibit an overabundance of personality.

### Fans and contributors

Fans and contributors are the ideal community members. They'll spend their time nurturing and pruning your community's content, and may even create new material that grows your audience and funnels members toward conversions. They might appoint themselves as moderators, spam catchers, or new user welcoming committees.

Their contributions often go beyond just posts. Contributors may use a website's APIs to develop tools that extend what your site can do, or they may find new ways of using your products and services that you hadn't intended to—or had time to—develop. They're your regular contributors and your grassroots support across the Web. They represent a small but valuable portion of your community. Love them and nurture them, and give them the authority they deserve.

From a monitoring standpoint, you need to analyze their behavior and use it as a "control group" against which to baseline other segments. For how long do they visit? How many comments a day do they make? Where do they link from? These measurements will give you clues about how to turn other segments of your community into the contributors you crave.

### Lurking, occasional, and bursty users

The vast majority of your community's members fall into three groups:

- Bursty users, who only post when they feel passionate about something being discussed, but do so furiously for a brief period and disappear shortly thereafter
- Occasional contributors, who chime in from time to time, but don't feel a sense of ownership
- Lurkers, who occasionally watch the proceedings without contributing

Even relatively passive community members contribute to the community to a small degree by amplifying and selecting good content, or by flagging abuse. This may come in the form of reTweeting, digging, or upvoting, which reinforces the popularity of superior material and helps it rise to the top of search rankings or message lists.

Many community sites go to great lengths to make it as easy as possible to contribute. As a result, when a community feels passionate about something, lurkers quickly become active participants.

Consider Electronic Arts' release of the game *Spore*, which contained strong copyright protection. The game's Amazon.com page quickly yielded over 3,000 comments and ratings in just a few weeks, as shown in Figure 13-5. By comparison, every version of publisher Maxis' bestselling *SimCity* game, dating back over 10 years, has less than 1,000 reviews.

### Disengaged users

Disengaged users are those who spend very little time within your community portal, entering and leaving almost immediately. They were either the wrong visitors in the first place—having arrived by accident—or have simply grown uninterested. Perhaps they signed up to get access to software or one-time content. Now, your email messages to them are bouncing.

If you're running a community, your goal is to optimize your points of entry so that users stay, subscribe, enroll, and bookmark the site, minimizing disengagement. Remember, too, that disengaged users can skew your metrics: you need to calculate your reach as the percentage of *active* community members who saw your message, not a percentage of all enrolled community members who saw it, to avoid counting disengaged users.

## The Topics: What Are They Talking About?

Communities emerge to discuss something, whether within a Facebook group, a mailing list, an IRC channel, a blog's comment thread, or a Twitter hashtag. You want to be aware of any topics that concern you.

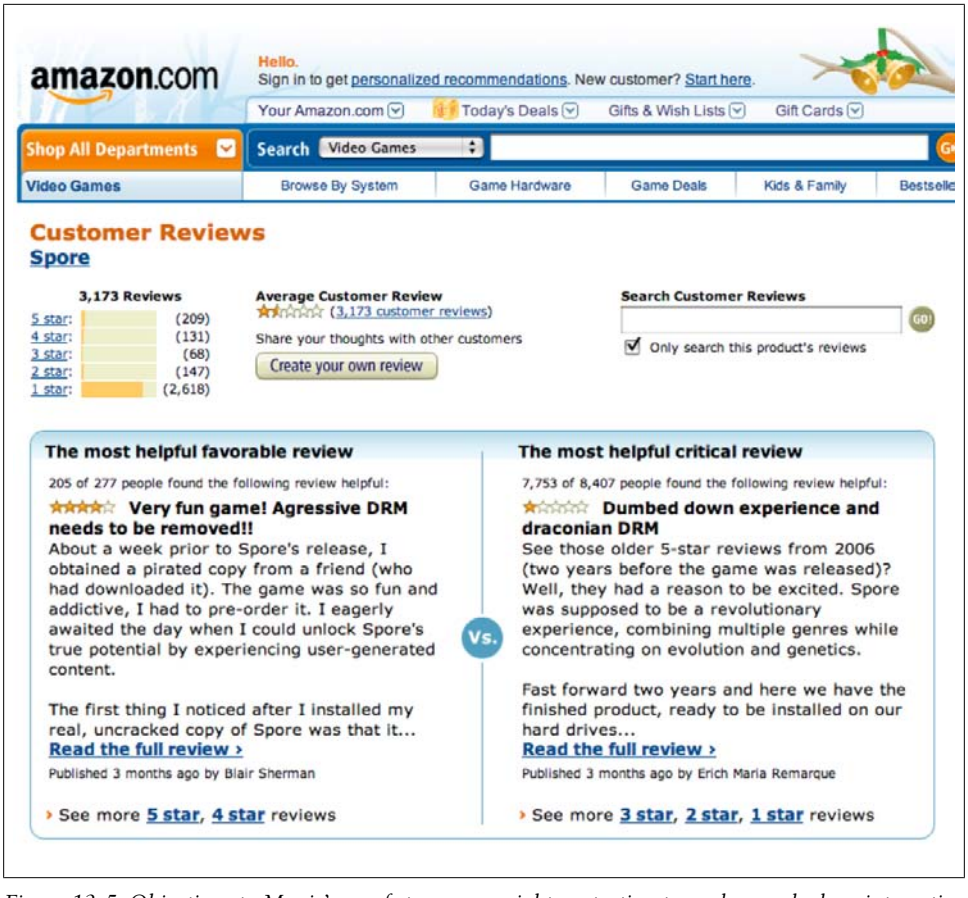


*Figure 13-5. Objections to Maxis' use of strong copyright protection turned many lurkers into active community members for a brief period*

You find these topics by searching for them in the aggregate (most popular Twitter hash tags or most active Facebook groups, for example) and in the individual (who's discussing a topic on Twitter or which Facebook users are members of a group). By tracking topics in the aggregate, you understand trends, community size, and growth. By tracking individual threads, you find evangelists and brand promoters who can help encourage the spread of your message and make others aware of the community to expand membership.

Platforms that you run yourself often have built-in tools to help identify popular subjects, and as an administrator you'll have access to user account information, so you can follow up directly with commenters. But if you're joining conversations elsewhere, you need to do some work to find out which topics are of interest to your business and your audience. Some obvious topics you probably want to track include company and product names—yours and your competitors'—and any topics that apply to you or

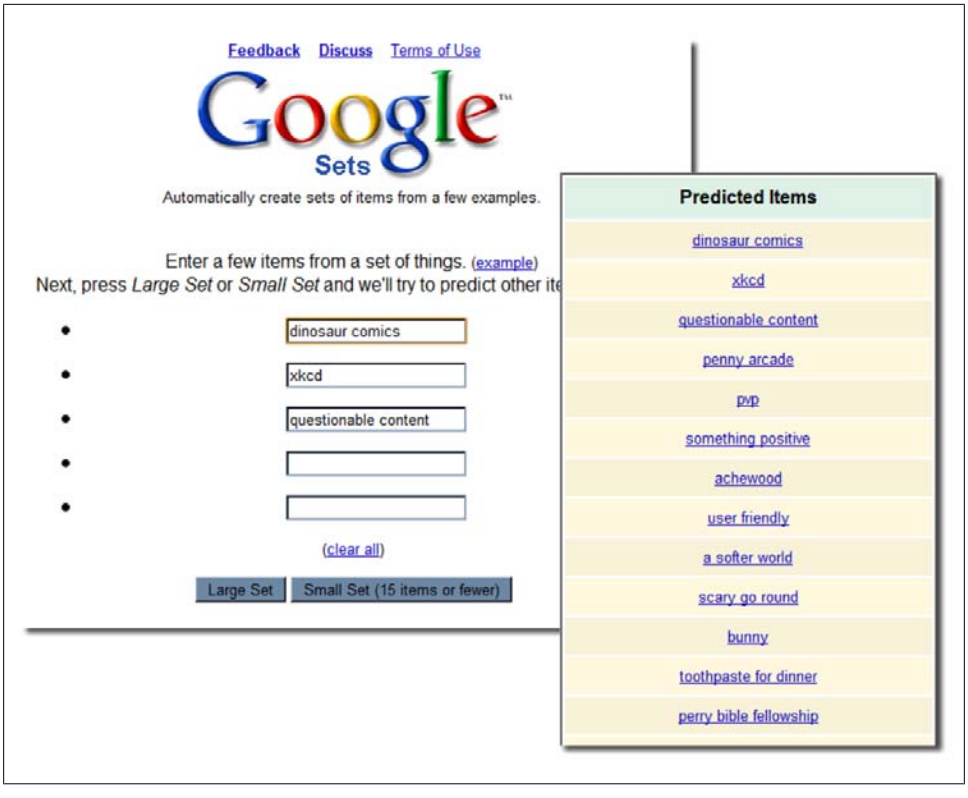


*Figure 13-6. Goggle Sets suggest sets of terms that are related to those you supply*

that you think are related, such as product categories, job descriptions, or geographic mentions.

If you're looking for inspiration about what to track, consider tools such as Google Labs' Google Sets tool (*http://labs.google.com/sets*), shown in Figure 13-6, which suggests related terms, companies, or products when you provide a few you know.

You should also check what terms others use to find you. A good list comes from organic search results within your analytics application, such as those shown in Figure 13-7.

It's not enough to look at the terms that are bringing *you* users—you should also look at companies that are competing with you for the mindshare of those terms. To do this, first find out which companies are also bidding on keywords that drive traffic to you. They may not be direct competitors, but they're competing for mindshare and attention.

Paid search uses an auction bidding model, so services like Google's AdWords offer a lot of detail about who's buying which terms. As a result, services like Spyfu, shown in Figure 13-8, can aggregate this kind of information.

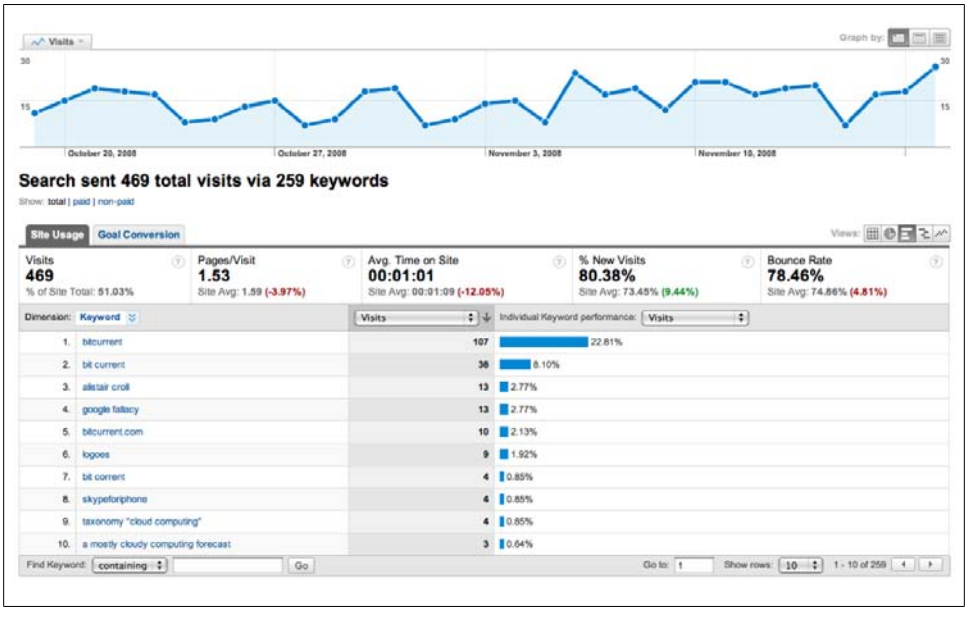


*Figure 13-7. Organic search terms in Google Analytics*

You can work the other way: starting with a known competitor (perhaps one that's stealing your mindshare), find out which keywords are sending traffic to that site. Figure 13-9 shows a list of organic and paid search keywords for analyst firm Gartner, generated by Spyfu.

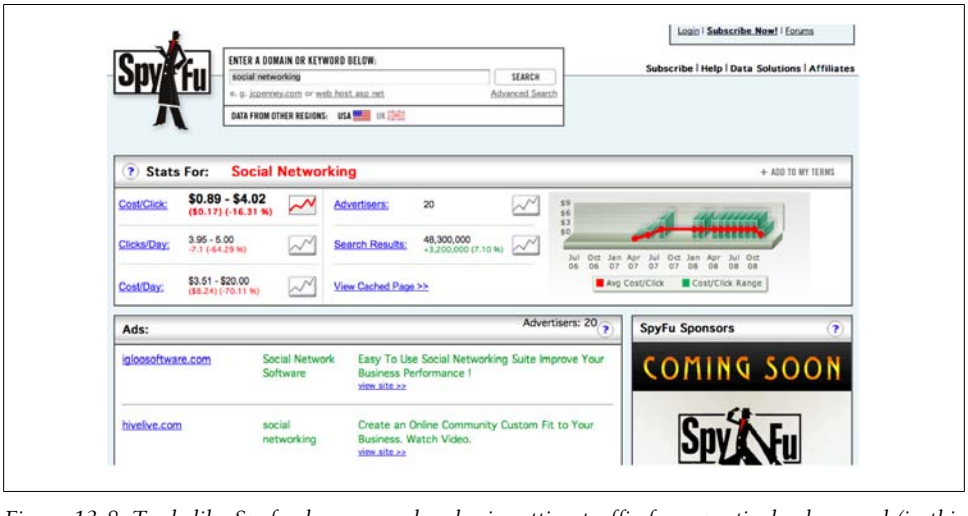


*Figure 13-8. Tools like Spyfu show you who else is getting traffic for a particular keyword (in this case, "social networking")*

You can then find groups and communities in which these keywords are being discussed. For example, you might enter a keyword into Twitter search to see what people think of it at the moment.

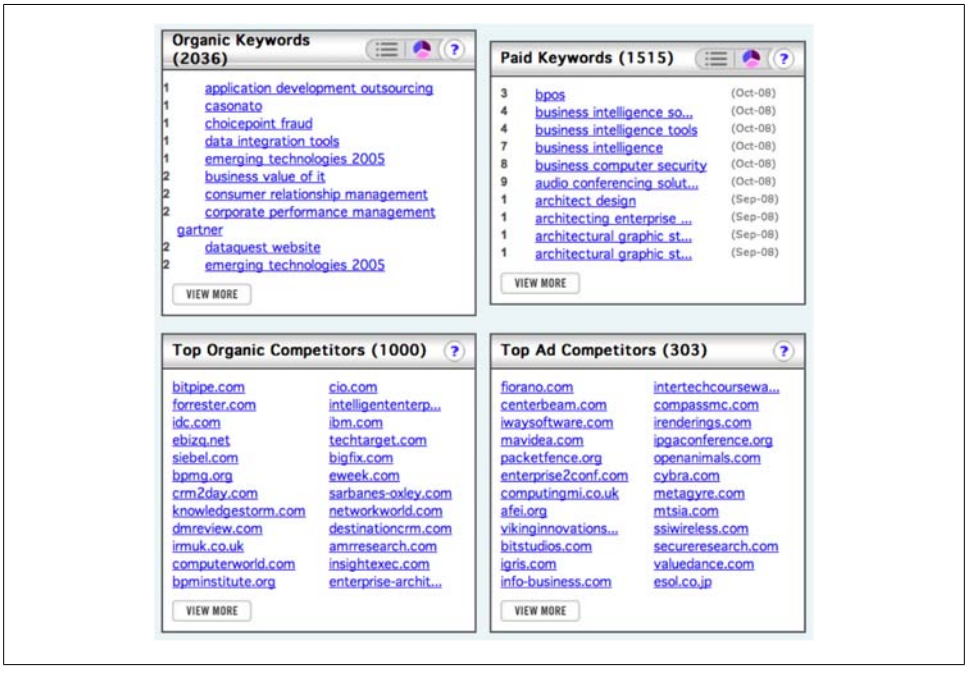


*Figure 13-9. Using a competitor's name, you can find out which terms are driving traffic to its site, then include these terms in the topics you monitor across communities*

You can also look at topics by geography. Google Insights, shown in Figure 13-10, will show you important news events, mentions of terms, and countries in which those topics are heavily discussed. You can use this information to check popular forums and portals in those areas.

Referring sites that are sending you traffic are often communities within which you're being discussed, as shown in Figure 13-11. Depending on your analytics tool and the amount of detail the referring site provides, you may be able to pinpoint the individual submission or comment that generated the traffic. You can then add that community to those you watch.


*Figure 13-10. Geographical distribution of the term "social networking" in Google Insights*

Unfortunately, analytics tools are often limited in what they can tell you about the sites that referred traffic to you. Many community sites won't identify themselves properly with referrers. This can happen for several reasons:

- The community platform may rely heavily on browser *bookmarklets* (StumbleUpon and Twine are two examples), which interfere with referral mechanics.
- The community may want to make money from analytics and as a result may actively hide referrers, replacing them instead with a promotion of its advertising platform (StumbleUpon does this).
- The content may have come from a desktop client—most notably microblogging tools like Tweetdeck or Twhirl—that doesn't leave a referrer.
- The community may have third-party websites that post content on behalf of a user's account. For example, Twitpic, shown in Figure 13-12, lets community members upload an image and then posts it to the account on behalf of the member.

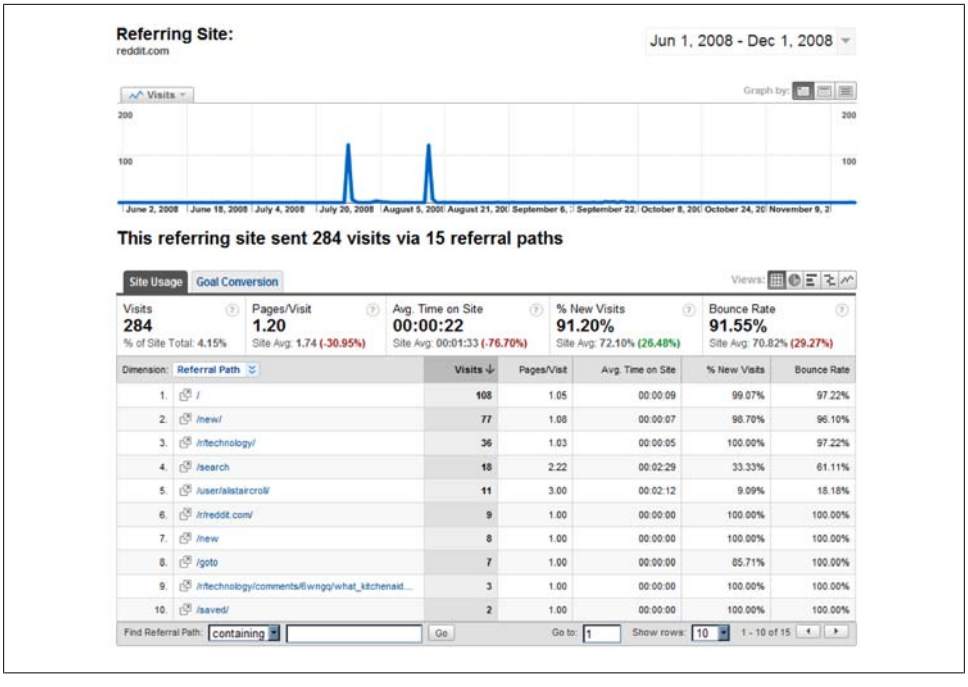


*Figure 13-11. Traffic from reddit, broken down by source in Google Analytics*

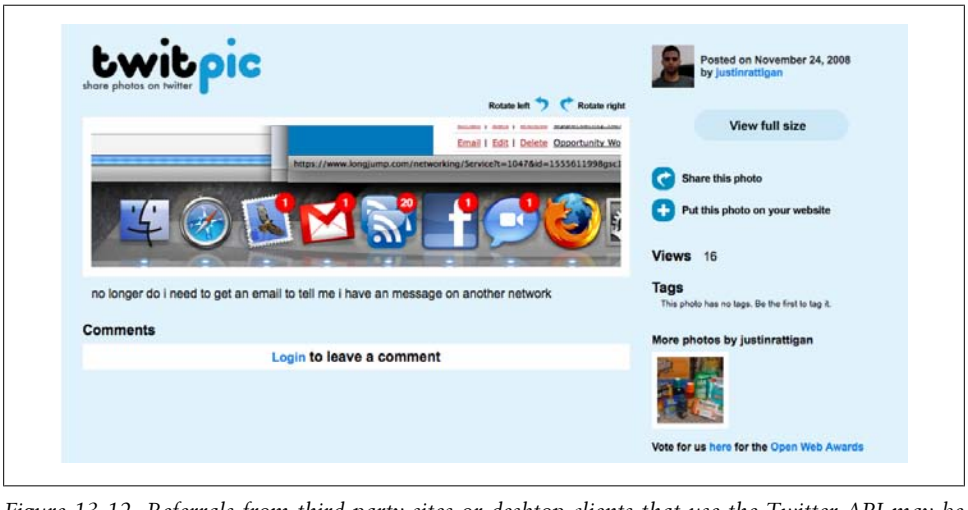


*Figure 13-12. Referrals from third-party sites or desktop clients that use the Twitter API may be obfuscated, appearing to analytics tools as if they were typed in directly rather than linked from Twitter*

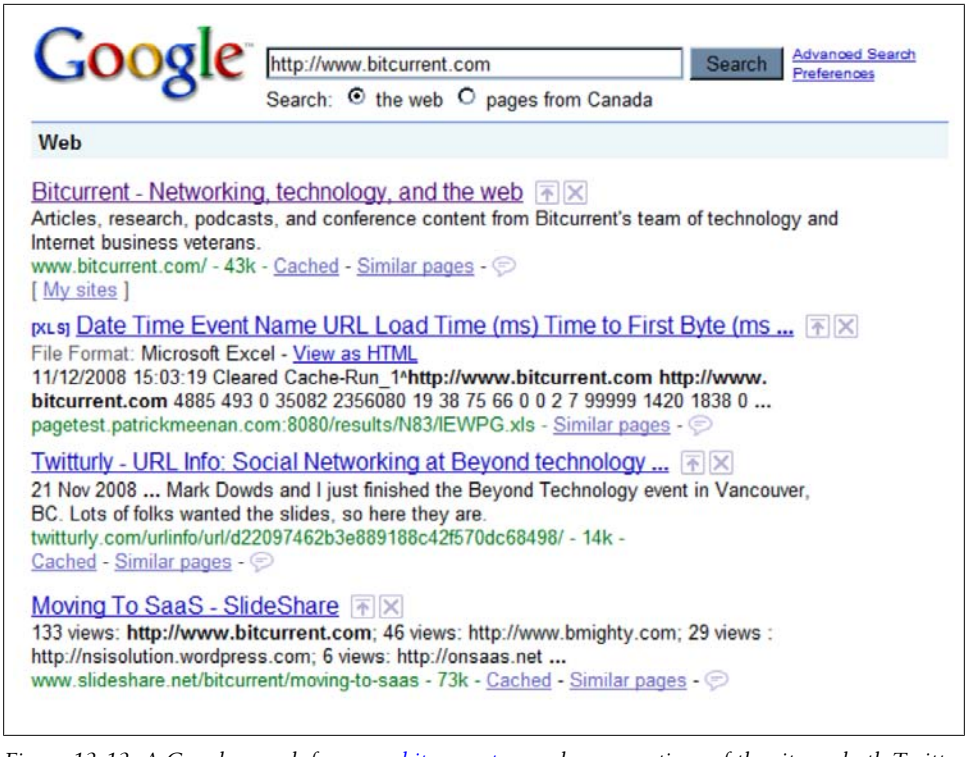


*Figure 13-13. A Google search for www.bitcurrent.com shows mentions of the site on both Twitter and Slideshare*

Because of how community sites often break referrers, if you want to find community mentions, you need to employ other techniques to track traffic back to the communities from which it came. You need to search for your URLs online and see where they're referenced. The search shown in Figure 13-13 shows not only the *www.bitcurrent.com* site, but also mentions of the site on both Twitturly and Slideshare.

URLs in their raw form are becoming less common. URL shorteners first emerged as ways to transmit long URLs, complete with parameters, through email without having them break across multiple lines of an email message, which would have rendered them useless. Shorteners have found new popularity in Twitter and other microblogging platforms where URL length is precious. Today, dozens of URL shortening services reduce long URLs into a simple, short string automatically using the process shown in Figure 13-14. This means that for many social networks and communities, your URL will be obscured when people mention you.

So how do you find places in which your URL has been mentioned online when it's been shortened? Here's one way.

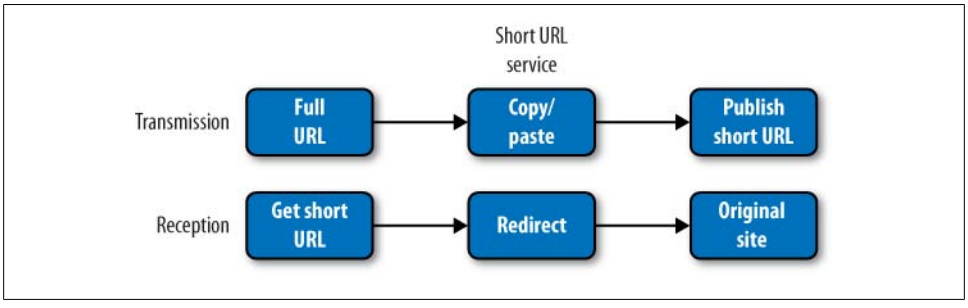


*Figure 13-14. URL shortening services such as bit.ly and tinyurl abbreviate your URL, then redirect those who click on the shortened version to the original site content*

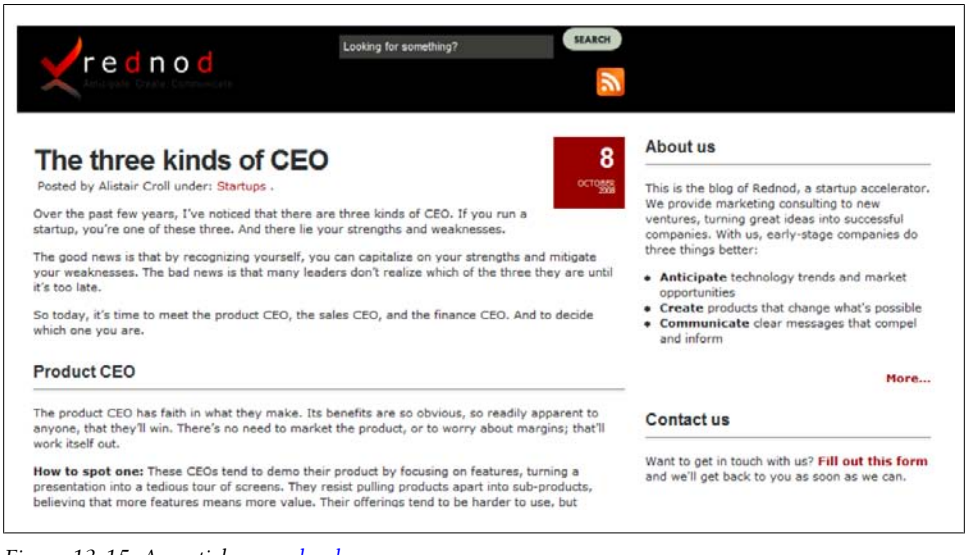


*Figure 13-15. An article on rednod.com*

Start with the online content whose incoming traffic you want to track back to a community, such as a blog post (Figure 13-15).

Enter the long URL of the content (in this case, *www.rednod.com/index.php/2008/10/08/the-three-kinds-of-ceo/*) into a URL shortening service (such as is.gd). The service will return a shortened URL, as shown in Figure 13-16. If someone has used the service to generate a short URL recently, you'll often get the same one that individual did (this depends somewhat on the service).

Enter the URL, in quotes, into a search engine (Figure 13-17).

The results will show you where your shortened URL has appeared on the Web. You can often find out more about the poster or the discussion thread in question,

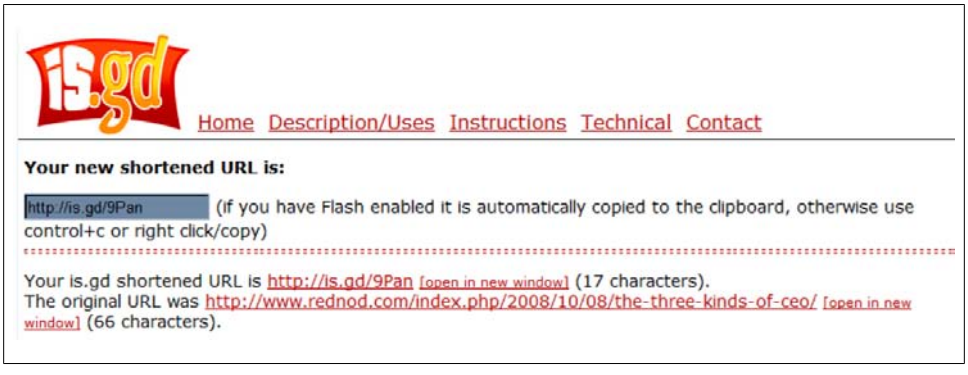


*Figure 13-16. Generating a shortened URL using the is.gd shortening service*

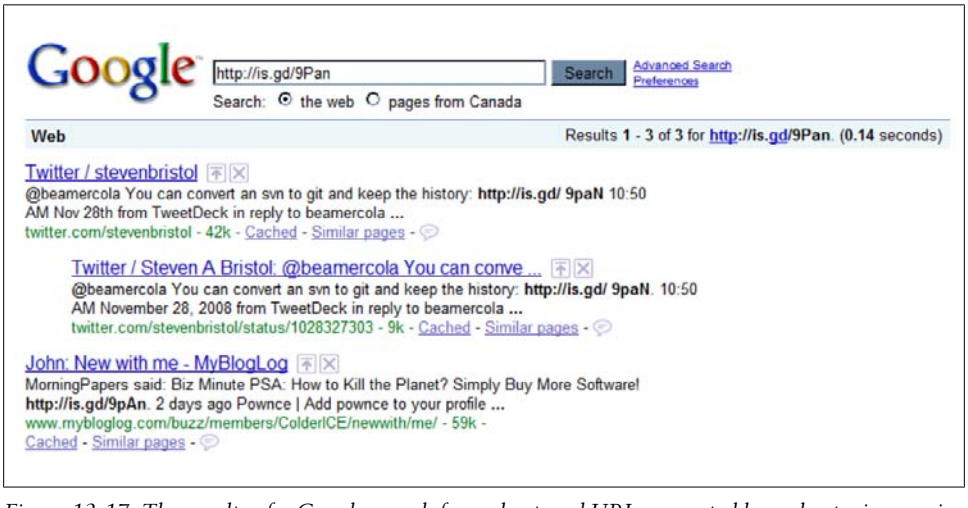


*Figure 13-17. The results of a Google search for a shortened URL generated by a shortening service*

particularly if that person is a member of a public community with a profile of some kind, as shown in Figure 13-18.*

Unfortunately, there are so many URL shortening services—each of which generates a unique shortened URL—that it's time-consuming to search for every mention you encounter online. Instead, use this approach judiciously to investigate particularly popular or contentious postings.

* We'd like to apologize to John for stalking him, featuring him in a book, and generally being creepy. As this example shows, you need to be careful when reaching out to your community. And we owe John some beers if he ever reads this.

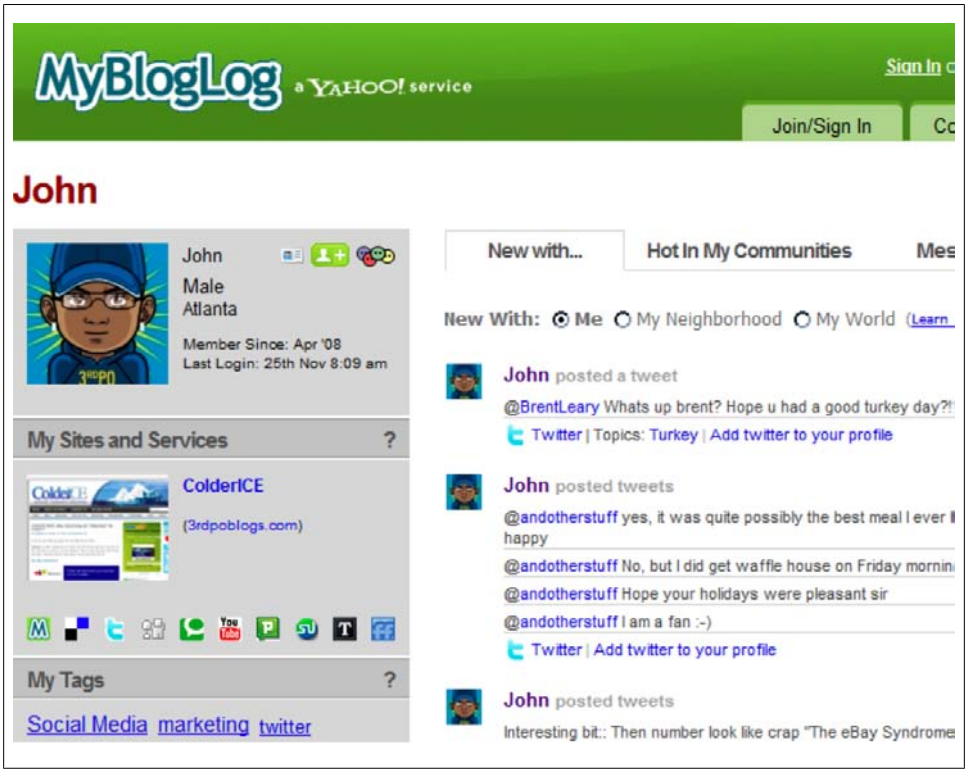


*Figure 13-18. Finding out more about the people who are amplifying your community message, what they talk about, and where they hang out*

Shortened URLs do more than save space. They provide us with a method of embedding additional context into short messages. You can generate a unique URI for each recipient or embed parameters in the URL that help to track its spread with web analytics. We provide a complete example of how to do this in Chapter 14, in the section "The Mechanics of Tracking the Long Funnel" on page 520. For now, remember that short URLs will be an important part of microblogging analytics; the URL has become the new cookie. We recommend using cli.gs, bit.ly, awe.sm, or some other service that provides its own analytics capabilities.

## The Places: Where Are They Talking?

We've looked at the people involved in a conversation, and at the topics you should be investigating. Where are those conversations happening?

There are hundreds of thriving communities, thousands of forums, and millions of comment threads to cover, but there are a small number of communities that will drive much of your traffic. Fortunately for us, popular communities also follow power laws.

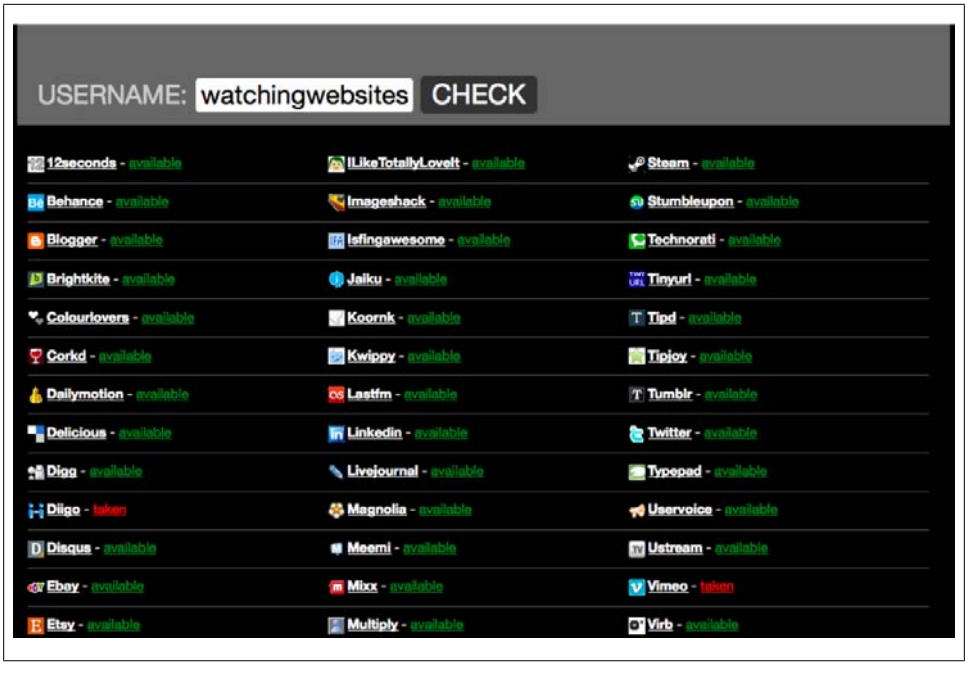


*Figure 13-19. The now-defunct usernamecheck.com allowed you to verify whether your name was taken on popular Web 2.0 sites*

We also have tools such as search to help us, and many of the communities make their sites searchable through APIs and aggregators.

Your first instinct might be to register yourself on as many sites as possible, some of which are listed in Figure 13-19, and search for conversations involving your brand.

That's a good step, but while it might reserve your seat at the communal table, it won't let you sit down and interact. Before we delve into the ways we can track online activity, let's first cover the types of community on the Web today.

Also check out *www.claimid.com* and similar offerings that manage all of your online profiles with OpenID.

## Different Community Models

We can divide up models of online interaction according to two important dimensions: the complexity of their messages, and how openly those messages are shared.

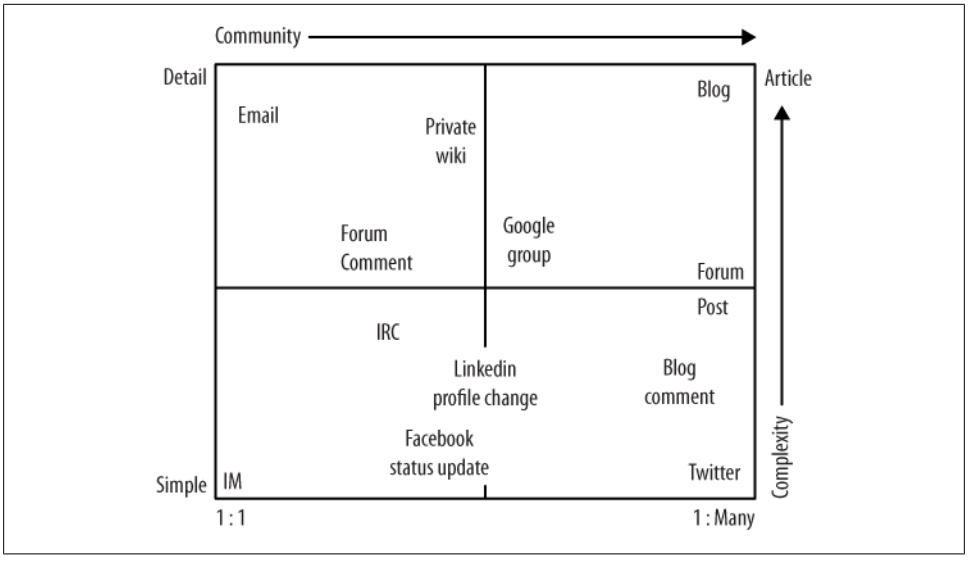


*Figure 13-20. Classifying community platforms by message detail and degree of openness*

*Message complexity*
: Some communications methods, such as IM and Twitter, are suited to simple thoughts and quick references. Others, such as email messages and blog posts, are better for detailed discussions and structured arguments.

*Openness of transmission*
: Some communications methods are one-to-one, such as IM, while others, like Twitter or blogs, are open to the whole world. Those that are more open encourage message spread, but make it harder to track individual users from a community message to a web analytics outcome the way we can with a one-to-one medium like email because we can't identify a unique conversation with them in order to start tracking their behavior.

Every community falls somewhere in this spectrum. A mailing list, for example, is a detailed message for a group interested in the same topic. A Facebook update is a relatively simple message for a group of friends. Twitter is a very simple message aimed at anyone who wants to hear from you. And so on.

Figure 13-20 shows how various community models fit within this classification.

All that's needed for a community is interaction and a group of people. The community can form in many places: groups and mailing lists, forums, real-time chat systems, social networks, blogs, wikis, micromessaging tools, and social news aggregators. Each has unique characteristics that are worth looking at in detail.

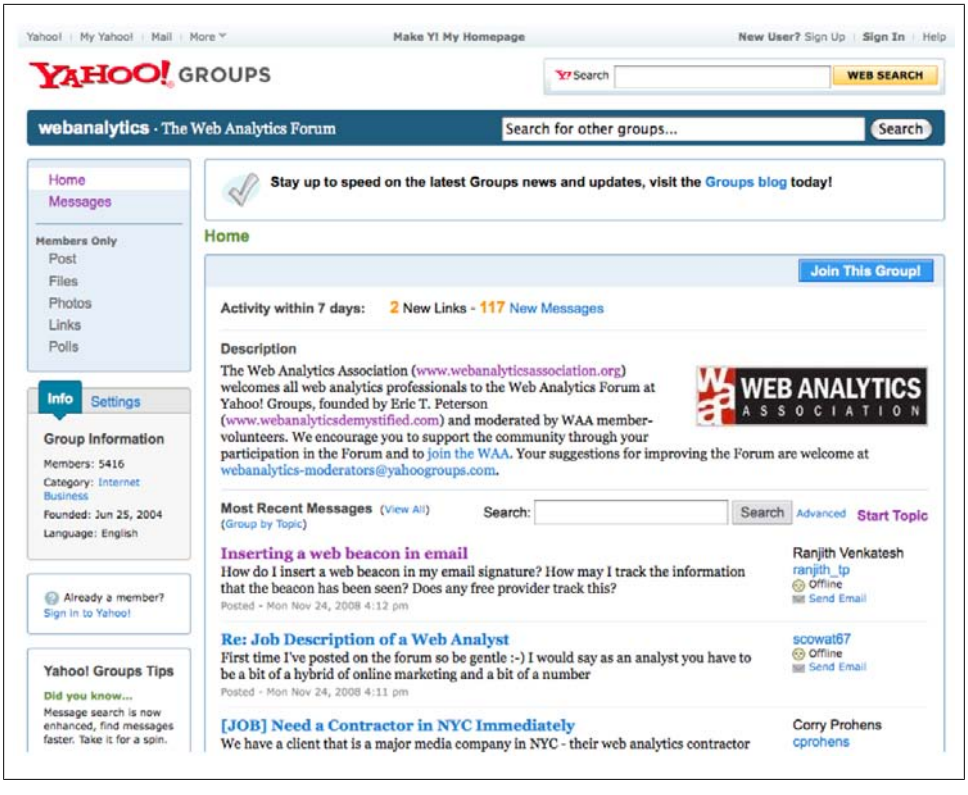


*Figure 13-21. The Web Analytics Forum found on Yahoo! Groups reaches the inbox of many web analytics professionals*

## User Groups, Newsgroups, and Mailing Lists

At the start of the previous chapter, we saw several early forms of community, including mailing list servers (listservs). These systems are still alive and well, and have been reskinned for the Web by the major portals, such as Yahoo! Groups (shown in Figure 13-21). Many of these portals have comprehensive frontends for community organizers to manage members, change policies, and send out bulk messages.

Mailing lists are unsurpassed in their ability to distribute messages to a relatively large audience that's interested in the same topic.

Most groups have four main functions, as shown in Figure 13-22. They include:

- *Maintenance* of the group and its policies. The administrator creates the group and decides who can access it, how messages are distributed, and the degree of control he exerts over the community (for example, he needs to individually approve new participants).

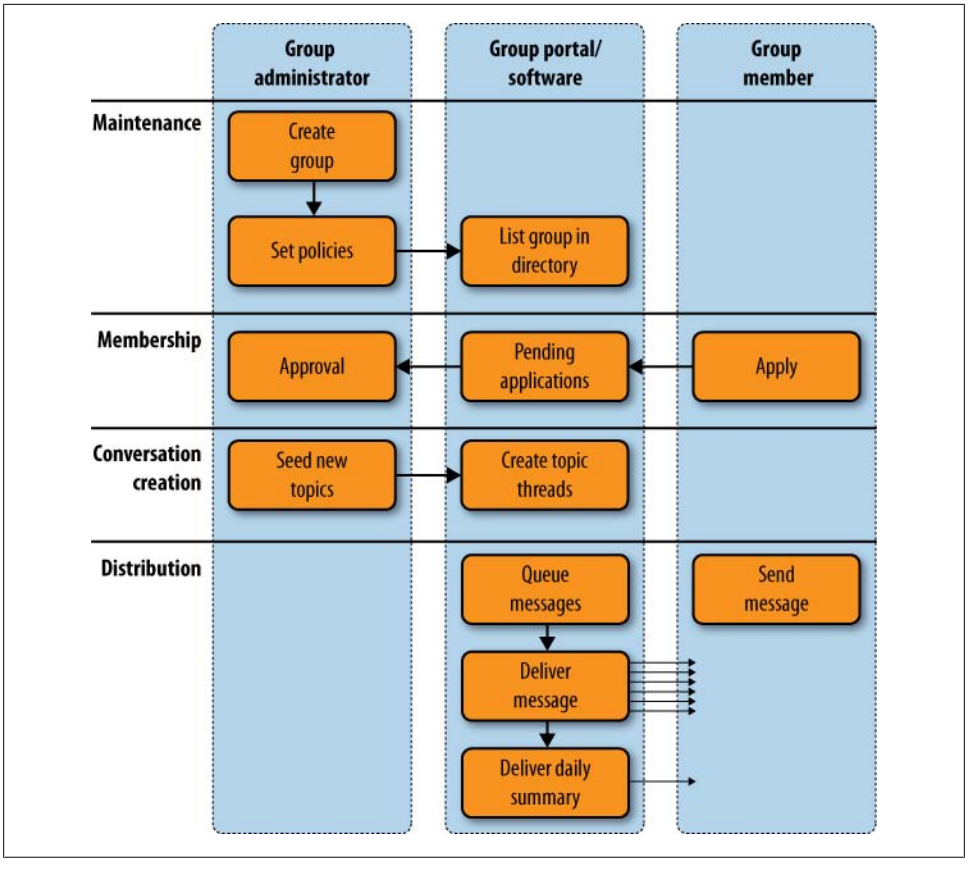


*Figure 13-22. The core functions of a mailing list management system*

- Managing *membership* to the group if applicable, including the approval of new applicants and blocking misbehaving members.
- *Creating conversation* by seeding new topics to spark conversation, and responding to messages from the community.
- *Distribution* of the day's conversations, either individually when they happen or as a daily summary, according to recipient preferences. This may happen via email or on the group's website.

Groups haven't changed much since the early days of the Internet. They're open and easy to join, and there are many of them out there. Some, such as the ones based on Usenet that are now available on Google Groups (Figure 13-23), are open and relatively unregulated. Others, such as mailing lists you run, are easier to oversee.

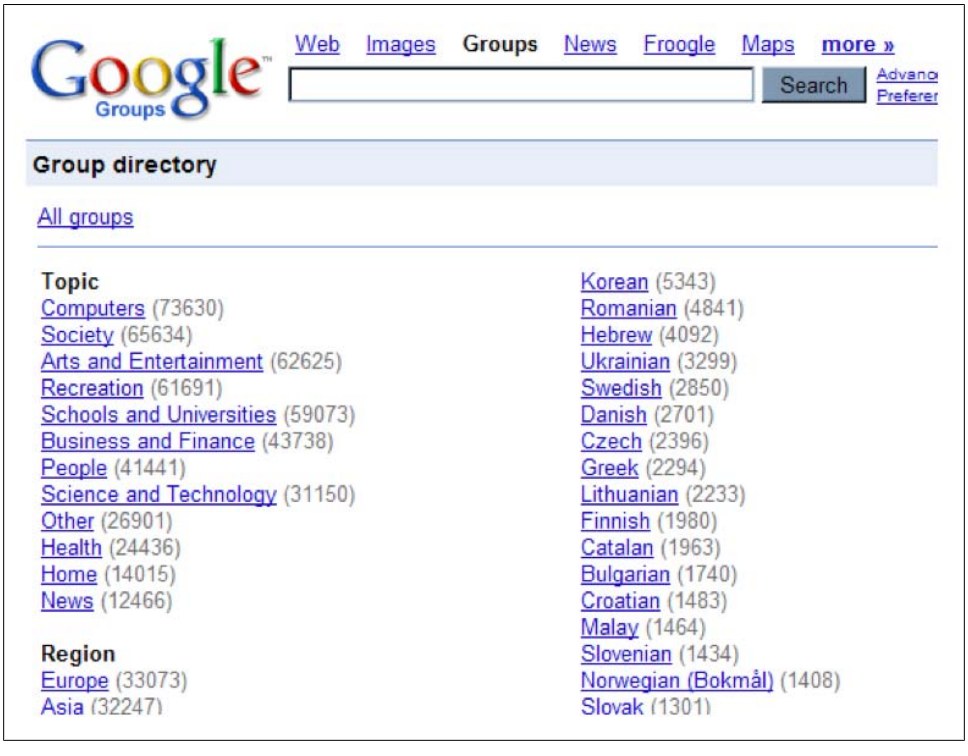


*Figure 13-23. There are hundreds of thousands of newsgroups available through Google Groups*

The most important changes to groups have been in their usability. At one time, the only way to search a group was by parsing its archives or digging through your inbox for a specific message. Now, however, these lists and discussions have been indexed by powerful search engines and are accessed through familiar web-based interfaces. As Figure 13-24 shows, search tools make it possible to find decades-old content quickly and easily across all groups.

Conversations on mailing lists and groups are comparatively slow-paced. Usually, email messages consist of longer, more well-thought-out messages. Many members elect to receive their lists only once a day when they subscribe to a list, as shown in Figure 13-25, meaning a conversation happens over several days.

The slower pace of mailing lists and groups makes them easier to administer, since you can review queued messages before they're sent out to remove content sent by spammers and trolls. If you hand-moderate posts, mailing lists often have higher-quality content than real-time models like comment threads. They're also easier to measure because you know how many subscribers you have; in some paid mailing list tools, you'll also have access to metrics such as open rate, which let you see how many subscribers actually read the content.

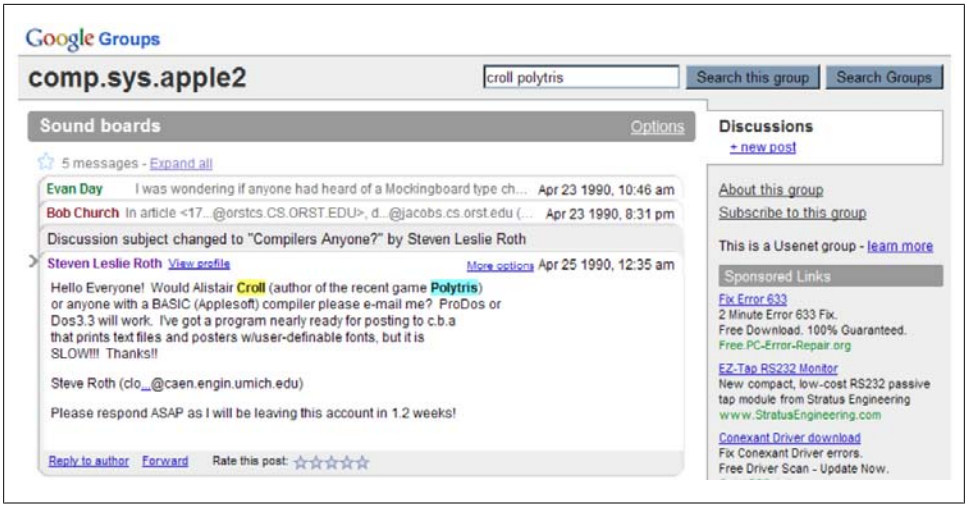


*Figure 13-24. Searching Google's index of Usenet posts for a specific string over decades of content*

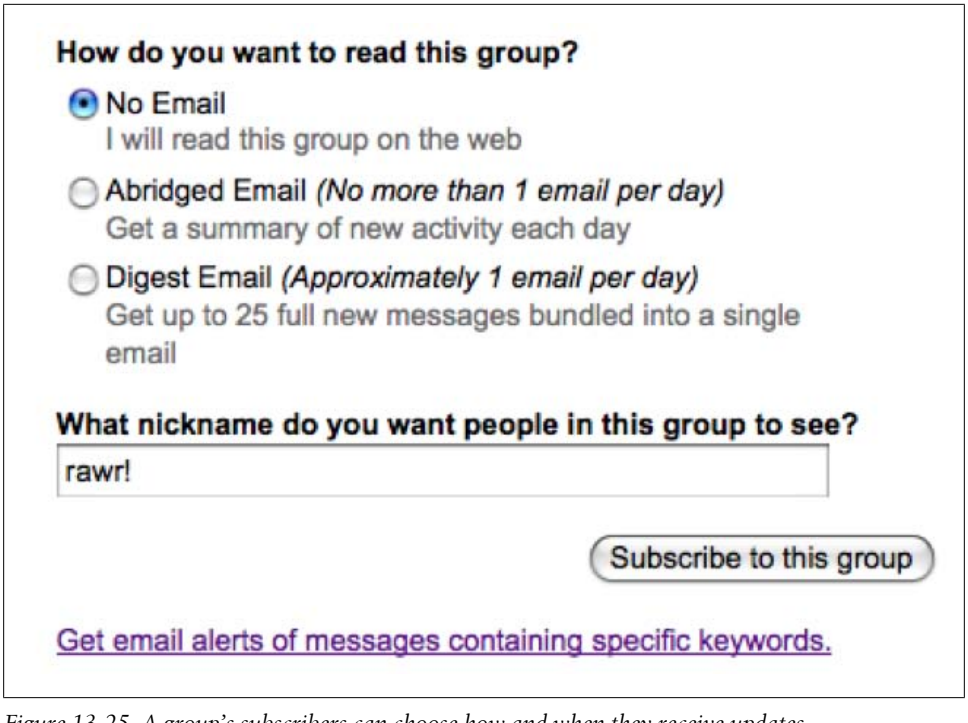


*Figure 13-25. A group's subscribers can choose how and when they receive updates*

These are the main roles you'll need to fill in a mailing list or group:

- *Administrators* maintain the mailing list or group. They may have created the group in the first place, and they define whether subscribers can join without moderator approval, whether the archives of the group are searchable, and how much editing and review happens before updates or comments are transmitted.
- *Moderators* deal with the content the group generates. They're responsible for enforcing the terms of service, weeding out abusive or illegal postings, and letting the administrator know about problems. They may be self-appointed or nominated subscribers who want to help, and have been given control by the administrator.
- *Subscribers* are the participants in the mailing list, and they may include the frequent participants, occasional contributors, lurkers, and disengaged members we saw earlier. If you're joining a mailing list or group that relates to your organization, you'll be doing so as a subscriber.

## Forums

Close cousins of mailing lists, forums are online discussions that emerged from the BBS world. They consist of shorter responses that are viewed immediately within a web page, and are organized chronologically within topics. An example of a forum is shown in Figure 13-26.

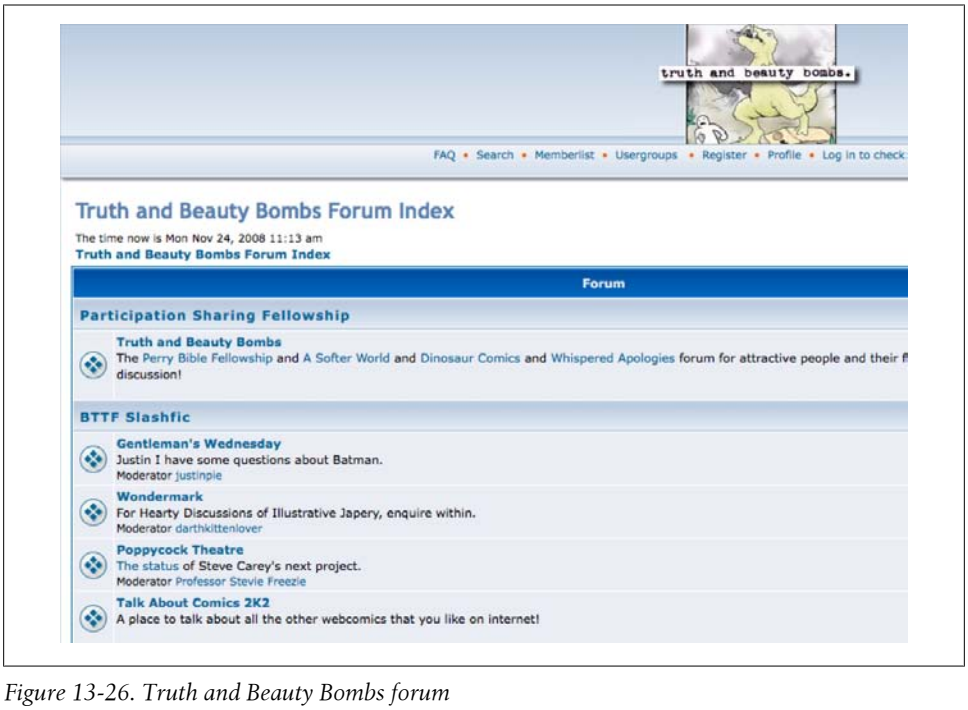


*Figure 13-26. Truth and Beauty Bombs forum*

Forums have low barriers to entry, often letting visitors read their content when they aren't enrolled, and asking for only basic information, such as a nickname and email address, in order to reply to an existing post or create a new one.

Some forums organize conversations into requests and responses in the form of threads, letting participants respond to responses and nesting the various threads hierarchically to make them easier to follow. Other forums simply organize all posts on a topic sequentially.

Forums are monitored by administrators and moderators:

- *Administrators* maintain the forum itself. They add, modify, and delete sections, patch the system, and collect and aggregate metrics for reporting. Their main involvement in content is when they must act as final arbiters when things get out of hand.
- *Moderators*, on the other hand, form the front line against spam and abusive contributors. Because forums are relatively open, the medium is untrusted. Moderators may also start new topic areas in a forum.
- *Members* respond to topics and to other people's posts. In a forum, members often rank their peers based on contributions. This ranking can be used when filtering content. In other words, a member might suppress comments from members with a low ranking. As a result, you need to add value and earn the respect of other forum members if you want what you say to have a wide audience.

Forums offer a variety of tools to keep you up-to-date with what's happening, so even if you're not administering the forum, you can still receive alerts when someone responds to your content or subscribe to an RSS feed on a particular topic. Most forums are also searchable. Though they provide a more interactive back-and-forth than mailing lists and groups, forums still aren't real-time. That's the domain of chat.

## Real-Time Communication Tools

Instant messaging is a one-to-one medium, but real-time chat among groups is a common community model that can involve hundreds of participants across tens of thousands of topics. The biggest of these groups is *Internet Relay Chat*, or IRC.

IRC has existed since 1988, predating the Web. It has its own protocols that make it distributed and highly resilient. Entire books, such as *IRC Hacks* by Paul Mutton (O'Reilly), are devoted to IRC, and as you're reading this, hundreds of thousands of people are currently connected to an IRC network.

IRC is organized into topic "channels" hosted across servers that relay conversations to participants, as shown in Figure 13-27. IRC requires a greater degree of technical acumen than simple web-based chat, despite the efforts of sites like Mibbit to try to make using IRC easier. While this means IRC isn't a popular tool for interacting with

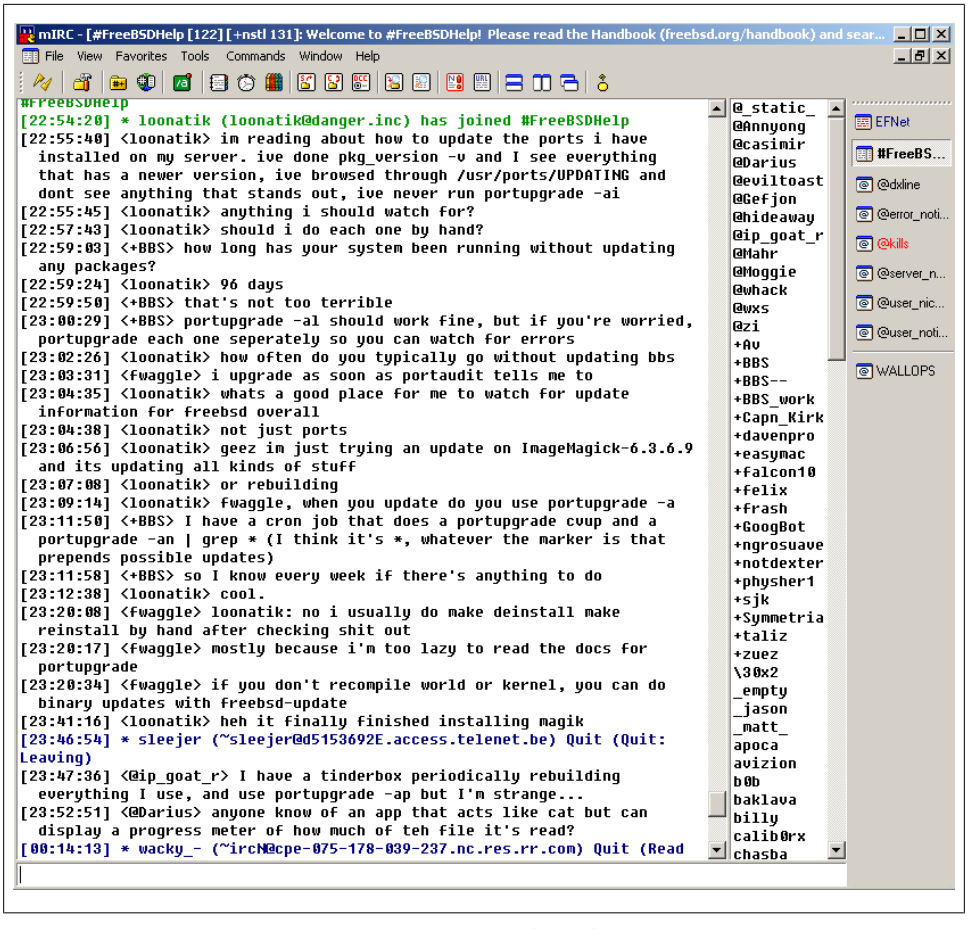


*Figure 13-27. A FreeBSD support community on IRC (EFnet)*

a mass-market audience, it can be an elegant and cost-effective community platform for technical groups.

Because IRC is its own protocol and platform, and can be completely anonymous, it's subject to exploitation:

- *Hackers* use IRC to control armies of infected PCs (botnets). The much-maligned IRC community has tried to address these problems itself, but its reputation has been tarnished and some carriers have tried to block IRC traffic in an effort to reduce the spread of botnets.
- *File sharing groups* can contact one another anonymously to share files that may violate copyrights.
- *Spammers* use scripts to join channels, send URLs, or advertising messages, then leave—a process known as *drive-by spam*.

- *Trolls* post inflammatory content in the hope of provoking a strong backlash, such as ruining the end of a movie ("Bruce Willis is dead") and then hoping for inflammatory responses that take the community discussion far off topic.

If used well, IRC can encourage customer self-support and provide easy interaction between your organization's technical team and knowledgeable end users, particularly in IT fields. Many employees in your organization can casually monitor a chat room, answering user questions as they come in throughout the day. IRC is also an excellent source of breaking news, given the real-time, fast-paced nature of the platform, so it can quickly communicate zero-day exploits or other technical concerns.

Recall, however, that IRC is an untrusted network. There are a high number of one-time visitors to an IRC community channel, and only a few recognized regulars. It takes time and patience to build a reputation, and channel operators are usually distrustful of newcomers. Problems happen far less often in topical, uncrowded channels, where it's harder for a botnet manager or file swapper to linger anonymously, and spam reaches a smaller audience.

Unlike forums and mailing lists, IRC doesn't have built-in monitoring tools or APIs. Administrators and operators rely on *channel bots*—scripts that eavesdrop on conversations and aggregate statistics on popular keywords, top talkers, and so on—to keep track of what's happening.

Existing communities, such as Facebook and phpBB, may also have simple real-time chat functions (as shown in Figure 13-28) to make them more interactive and current. This kind of chat may encourage members to stay engaged, but it's seldom the primary vehicle for communication within the community.

Simple chat platforms that are part of a larger social network inherit the trust model of their overall platforms. In other words, you've already got a profile on Facebook and a network of contacts within the broader community, so you're less likely to abuse Facebook's chat. Members will likely behave within the standards of the broader community (such as the phpBB forum) when using real-time tools such as its chat function. As a result, real-time interactions that are an extension of other community platforms require less policing and are less prone to abuse.

Real-time chatrooms have a three-tiered hierarchy of moderators and coordinators:

- *Administrators* focus on the technology. They often own the machine, manage traffic flows and routing, and maintain the latest versions of software. They understand peering and are often experts in routing protocols such as BGP (Border Gateway Protocol), because they must stitch together thousands of user sessions across hundreds of machines. They may also work with carriers to resolve security issues and IRC abuse.
- An *operator* ("oper") has the ability to disconnect users from the network (known as "killing" them), but in return for this power must arbitrate channel disputes,


*Figure 13-28. Extended chat functionality integrated into the phpbb forum software*

work on spam prevention, and must try to kill botnets as quickly as possible by blocking bot traffic.

- Individual *channel operators* (known as "ops") control individual channels, changing permissions and banning users from their immediate channel as they see fit.
- For *participants*, joining IRC is easy; getting some reputation and the respect of your peers takes work.

## Social Networks

Social networks are communities built around the interpersonal relationships of their members. By keeping track of who each participant knows—and how they know one another—members each create a private community through a whitelist of known friends.

While two members' friends lists might overlap considerably, very few people have identical social graphs. The community grows as members explore their friends' contacts, befriending them and growing their own social graphs. Requests from strangers are discouraged, which improves the quality of interactions and reduces traditional spam and trolling.

Facebook, in particular, differs from other social networks because it's a place for connecting with real-world friends. Facebook social graphs look more like their real-world counterparts.

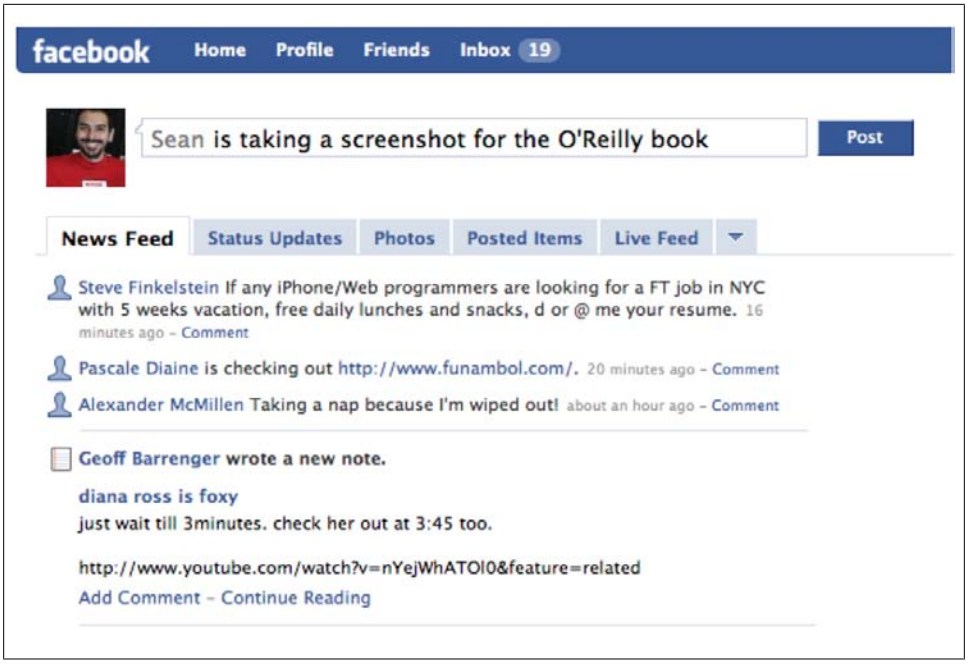


*Figure 13-29. A Facebook lifestream*

Social networks like Facebook and Friendster began as hyperlinked online contact lists, but have increasingly shifted toward sharing small updates among friends, as shown in Figure 13-29. This interaction, commonly referred to as a *lifestream*, contributes to a feeling of ambient awareness within the community, much as microblogging does.

Social networks often include many other community features, such as image sharing, group games, discussion groups, and real-time chat, all built around the central idea of the friend list and the social graph. Messages may be open to the community (as is the case with Facebook's "write on my wall" and "My News" features) or limited to only the sender and receiver.

Behind the few very large social networks are many smaller, more focused ones. *Compete.com* estimates that only five social networks have over five million visitors a month, but there are many smaller ones, as shown in Figure 13-30.

Outside these top five, niche communities thrive. Community Connect operates several smaller sites (BlackPlanet, Glee.com, MiGente, and AsianAvenue) that target specific groups based on ethnicity and sexual orientation. And Disaboom provides a social network for people dealing with or caring for those with disabilities.

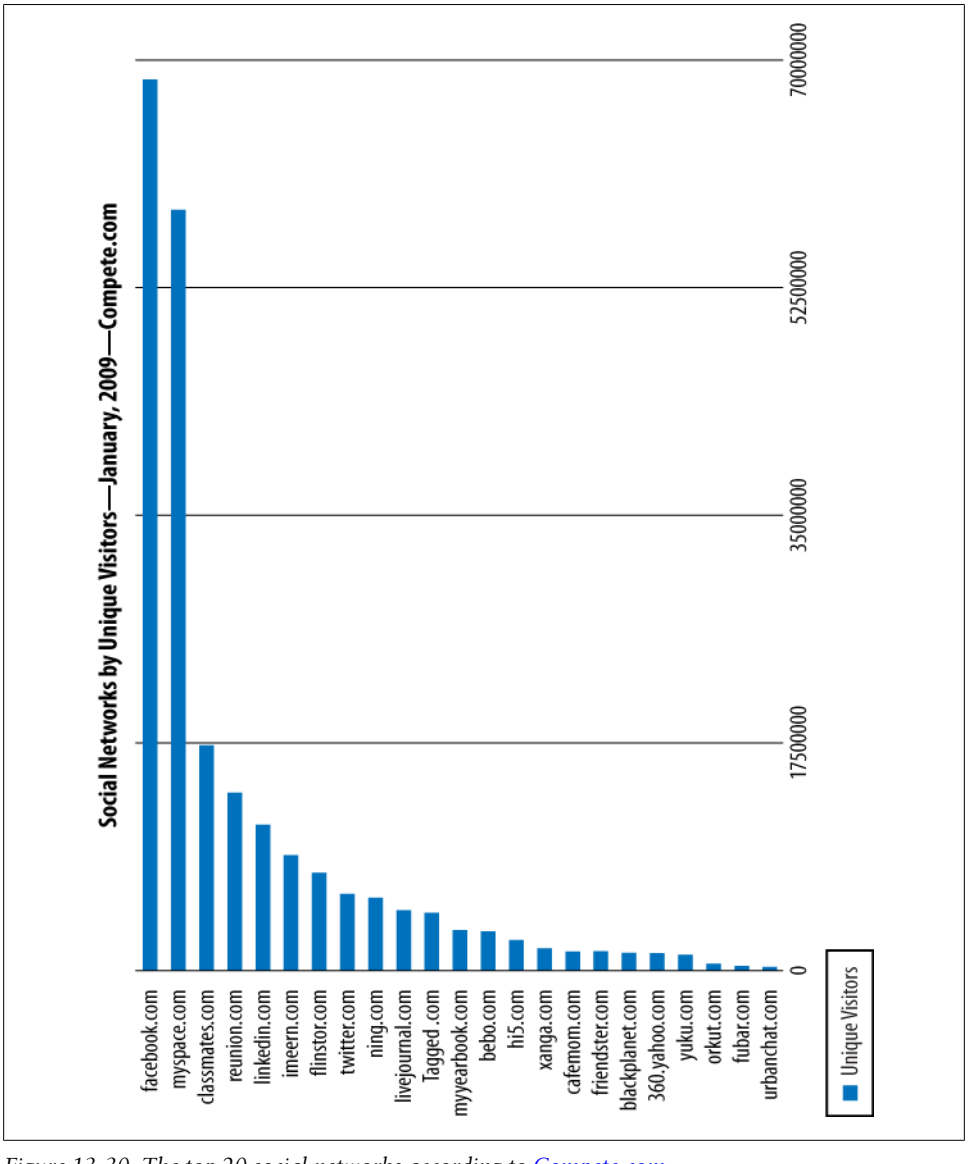


*Figure 13-30. The top 20 social networks according to Compete.com*

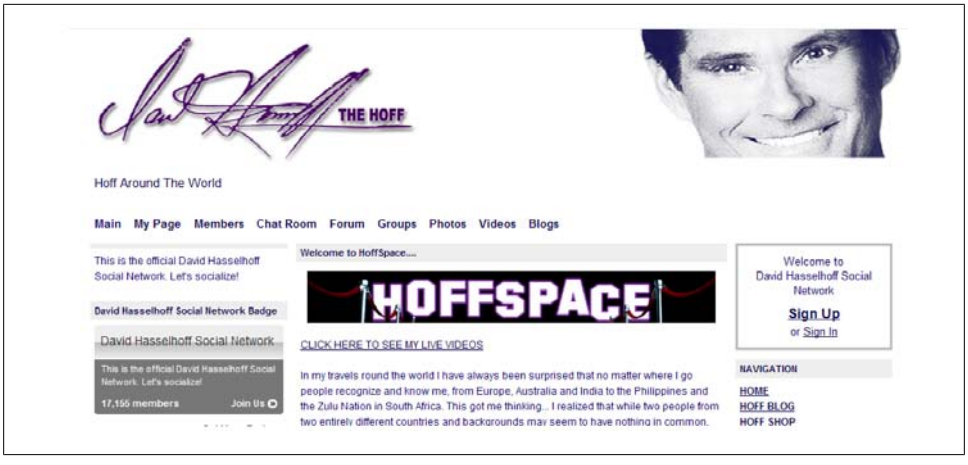


*Figure 13-31. A social network built around an individual*

Many of these social networks contain groups in which a member may occasionally be the topic of conversation. In some cases, the platform may arise around a particular person, brand, or company, as shown in Figure 13-31.

At the other end of the social network spectrum, companies like Ning, shown in Figure 13-32, allow users to create small social networks with features similar to those found on Facebook. This can be an easy alternative to building your own platform—you won't have to build and deploy your own software, but you'll be limited in the data you can collect about your members and their discussions, as well as your ability to moderate and control what's said.

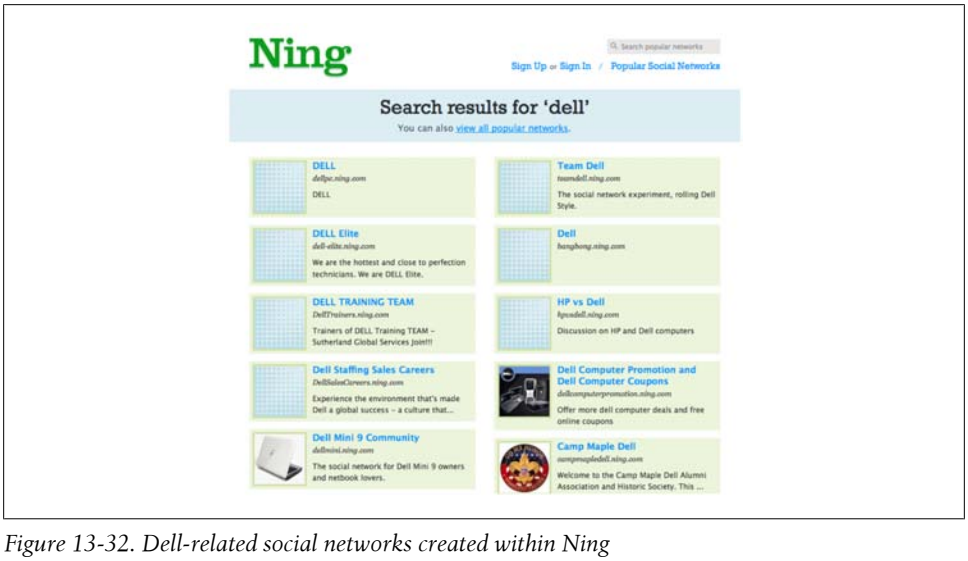


*Figure 13-32. Dell-related social networks created within Ning*

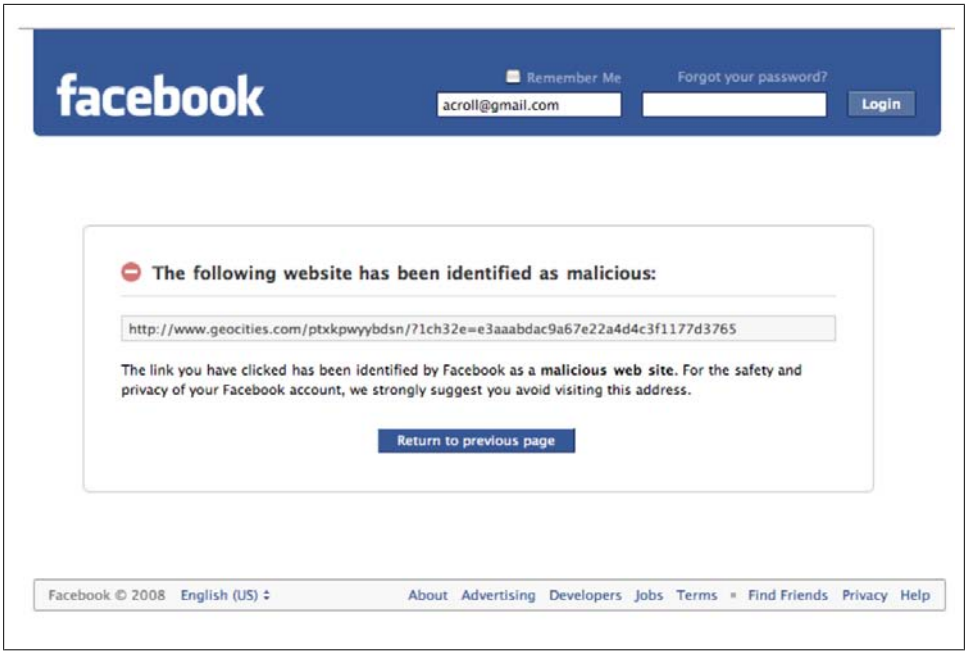


*Figure 13-33. A Facebook warning about a malicious site invitation sent from an infected friend's computer*

Platforms built around a social graph strike a balance between openness and trust, which helps to prevent abuse. But spammers still work on social networks, either by sending friend requests that encourage users to check the spammer's profile or by infecting users' PCs and then using their accounts to send spam and malware links to their social groups, as shown in Figure 13-33.

There's a conflict here for community managers. The more open your social platform, the easier it is for casual users to contribute and join. However, the more closed it is—and the more of it you control—the better the content, the greater your analytical visibility, and the less likely the network is to be abused by spammers and trolls. If there's a closed network whose subject matter you care about, you have to join the conversation in earnest.

Here are the roles in a social network:

- The *platform operator* (the person running the servers) has complete control over the tools, functions, and messages within the system. Unfortunately, if you run your own platform, you'll need to convince others to join you there—and getting a community's attention is often the hardest part.
- On a site like Ning, the *social network administrator* creates a new social network and defines which features are available to whom.

- Social networks often include groups and events that form around a particular topic. *Group administrators* have a certain degree of control around membership and policies, as well as the ability to communicate with the group and to create events associated with the topic.
- *Members* can participate in a social network, but they need to build their network of friends by inviting others and connecting with people already using the network.

## Blogs

Blogs are content management systems that simplify the task of publishing on the Web. They often focus on a particular topic or community of interest. Blogs handle the distribution of content via RSS feeds, automatically update aggregation tools like Technorati, and structure posts chronologically in an archive. They also have comment threads and several other features to promote interaction between the blog's authors and its audience.

A blog is often the public face of a company to its community. Unlike other social networks seen here, in which all participants have a roughly equal voice, in a blog, you have clear authority to create content and to decide which comments to post. It's up to you to pen insightful posts that will encourage your audience to comment, tell others, and return. For an excellent list of tips on writing and maintaining blog content, see Chris Brogan's article at *www.chrisbrogan.com/40-ways-to-deliver-killer-blog-content/*.

If you're writing your own blog, know that it's not a magazine. The goal is to get in, make a point—often provocatively—and get out fast. Put another way, on a blog, a mystery story would begin with the words, "The butler did it. And he was right to. Here's why." An effective blog posting is the starting point for a conversation, offering a viewpoint with some backup material and inviting the community to respond within the comment threads.

Blogs have their own challenges, primarily spam and trolling, but there are many automated tools that can reduce spam and bulk moderate unwanted comments.

Here are the main roles you'll need to fill when running a blog:

- The *operator/administrator* implements analytics and measures uptime for the system, handles upgrades, adds and removes others from the system, and maintains the stylesheets and blog layout.
- The *editor/publisher* gets the final say over what goes live, and it's her responsibility to validate what you're writing as well as to ensure it's not libelous. Most blogging tools support an editorial cycle in which only editors can publish content.
- *Contributors* generate content, which is saved as drafts for review by the editor. Contributors probably have several posts in the queue pending publication. They should also read and respond to comment threads for posts they've written and, in conjunction with the editor, correct or retract information.

- Anyone can be a *commenter* on a blog. If the blog uses a consolidated comment system like Disqus or Mybloglog, the commenter may have a "persona" that ties together comments across multiple sites or over time.

#### A note on megablogs

Power laws tell us that while there are millions of blogs on the Internet, the majority of traffic goes to only a small number of them. These are the megablogs.

In recent years, the Huffington Post and Engadget have dominated Top 100 blog lists. In every sector, megabloggers have risen to prominence. For technology, it's Mike Arrington, Om Malik, Tim O'Reilly, and Robert Scoble. They have the ability to create a community within the comment threads of a single post, and have tens of thousands of micromessaging followers. When these sites and bloggers talk, others listen.

Does that mean you should focus community efforts only on the "big talkers"? Should your community strategy focus on trying to convince Kara Swisher or Walt Mossberg to blog about you?

The short answer is no, for three reasons:

- The megablogs are increasingly like mainstream media. This is a simple fact of being big. They're businesses, with bills to pay and an editorial beat to cover. As their audiences grow, it's harder for them to find a niche (though notably, GigaOm has worked hard at creating a family of smaller, topical blogs within its network). They're inundated with messages from traditional PR firms, who see them as just another outlet. It's impossible for them to read as much as they write.
- The big bloggers are wise to the games PR and promoters play to get their attention. They'll use their own networks of informants and experts—the less vocal "Interesting Middle" of the long tail—to find topics and validate information. So if you don't focus your attention on this secondary group, what you're saying to the megabloggers won't ring true when they check their sources.
- On megablogs, comment threads are where much of the community interaction happens. Every writer reads her own comment threads (those who tell you they don't are lying), and contributing useful information such as statistics or links to useful additional content is a great way to join the conversation without "pitching."

What you're after are the newcomers who are quickly gaining influence. For that, you need to engage with them on topics they care about—giving content to gain attention.

## Wikis

A wiki is a collection of pages written by users that can be edited by other users. While other social platforms have a hierarchy (topics for groups, threads for forums, time for blogs, and a social graph for social networks), wikis have no rigid hierarchy; rather, they are a collection of documents linked to one another by topics.

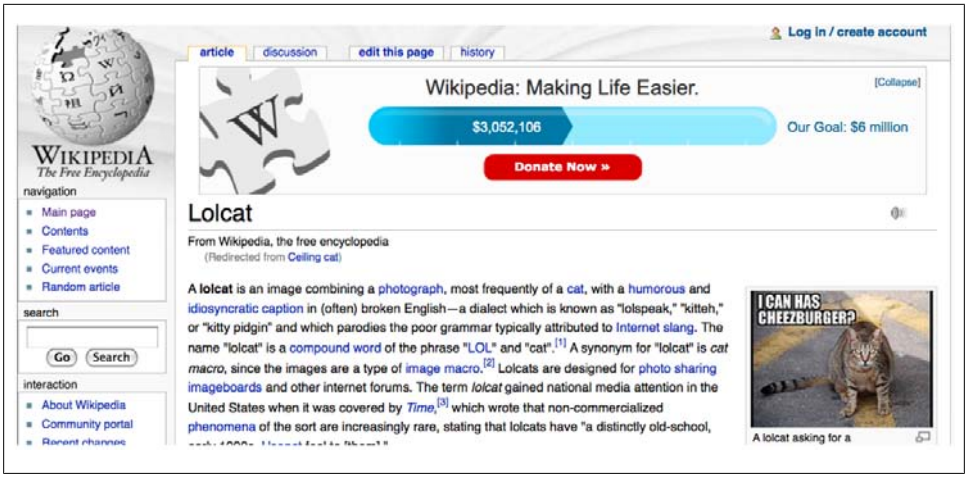


*Figure 13-34. User-generated content on Wikipedia*

The main focus of a wiki is collaboration, and a single page may have hundreds of edits, all preserved chronologically. Wikis make it easy to link to a new topic from within an entry and then return to complete the new topic later. The most popular example, shown in Figure 13-34, Wikipedia, is the world's largest and most comprehensive encyclopedia, but a wiki can be used for many kinds of content.

Wikis often contain reference material, but their openness makes them inherently less trusted than more formal reference sources. Wikipedia deals with this problem constantly; in some cases, the Wikipedia team has had to lock certain entries due to warring factions of editors or social pranks. For example, on July 31, 2006, Stephen Colbert urged his viewers to modify the Wikipedia entry about elephants, declaring them no longer endangered, and Wikipedia had to halt edits to the page when thousands of edits ensued, as shown in Figure 13-35. Only weeks later could the page be safely unlocked.

Despite this, and thanks in no small part to the efforts of hundreds of dedicated content auditors who check references and flag questionable content, Wikipedia has become one of the world's largest online information resources. But it's allergic to subjective or promotional content, so any content you contribute to it must be independently verifiable, contain reliable references, and be considered nonpartisan.

Wikipedia, which runs atop the MediaWiki platform, is a good model for any wiki you create or moderate. There are many other wiki platforms out there that you can use or run yourself, including TikiWiki, Wetpaint, and Socialtext.

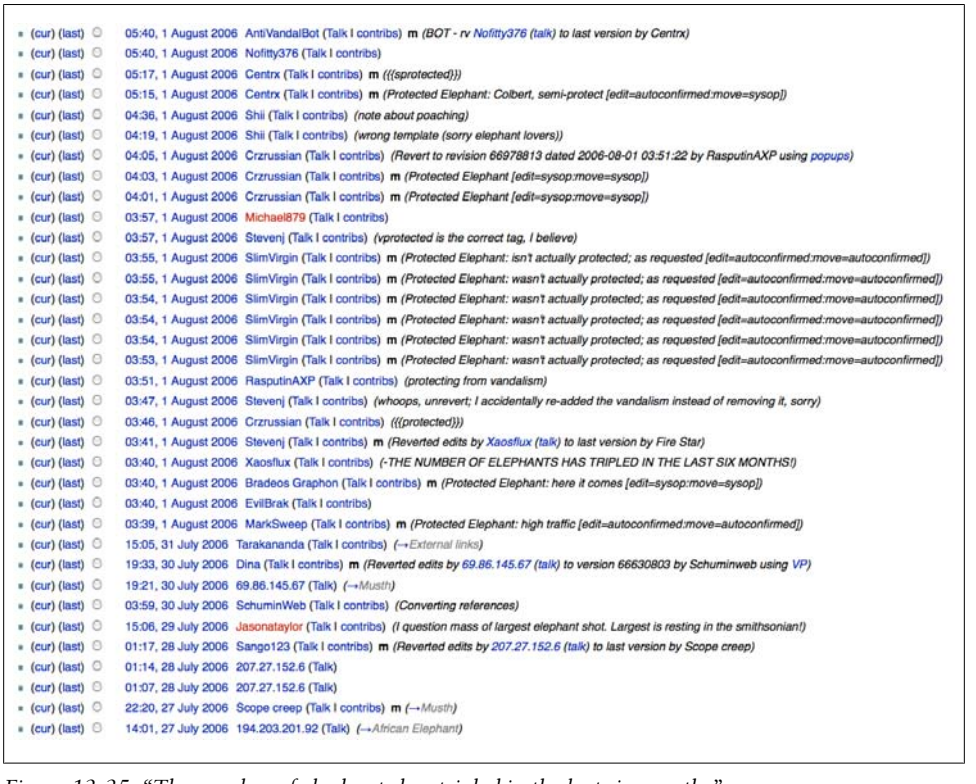


*Figure 13-35. "The number of elephants has tripled in the last six months"*

Here are the roles in a wiki:

- The *operator* runs and maintains the wiki platform and servers, and manages the "blunt instruments" of content management, such as locking down topics to prevent change. The operator also decides on layout and functionality. In Wikipedia's case, the operator is Jimmy Wales' Wikimedia foundation.
- *Authors* can be anyone with a claim to expertise; their authority should be supported by third-party material and references, or content they generate they will be quickly questioned. If you're running a wiki as a reference source for your customers, you may limit authors to employees or accredited subject matter experts. For example, genetic analysis site 23andMe has a reference wiki with articles authored by doctors built into the service.
- *Editors* are the content police of a wiki. They correct and update content, and flag problems such as poor syntax or missing references so that authors can fix them. They also have the ability to revert contributions to a previous version and try to identify bad behavior that an operator needs to address.

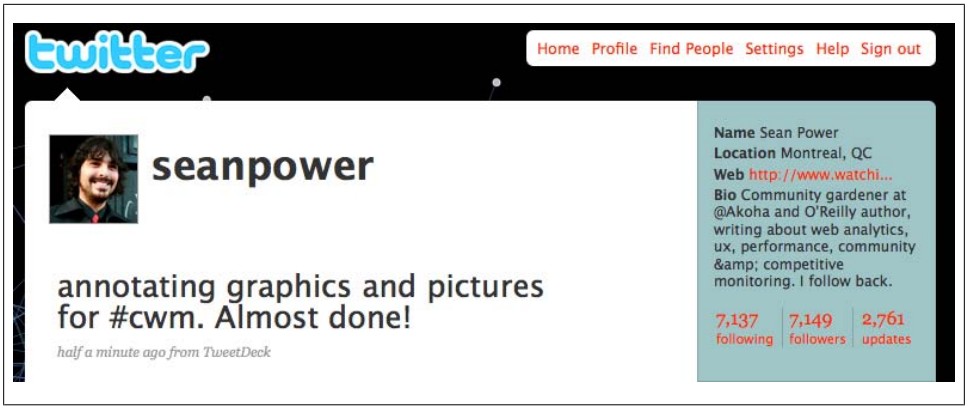


*Figure 13-36. A Twitter feed, showing the number of followers and updates*


*Figure 13-37. Facebook's "status" is a micromessaging tool*

## Micromessaging

Micromessaging, or microblogging, applications are streams of small messages shared with a group. There are dozens of platforms, but Twitter, shown in Figure 13-36, is by far the largest independent microblogging tool in terms of adoption and message volume.

Micromessaging does exist within larger social networks such as Facebook in the form of status updates like the one shown in Figure 13-37. In these models, however, interactions generally occur among a preexisting social graph rather than between strangers.

Micromessaging is a relatively new communications model, with unique characteristics that affect how we use it and what's appropriate. It's an RSS feed for people, a way to direct the attention of audiences, and a means of reaching the famous without burdening them with an obligation to respond. In short, *Twitter is an API for human attention*. While the rules of micromessaging are still being written, it has some unique features worth noting.

### Asymmetric following

On Twitter, following is politics. Well-known Twitter members have huge followings, and often use this network to promote blog postings or amplify messages they find interesting. We often judge members according to their ratio of followers to people they're following (see *http://twitter.grader.com/* for one ranking system created by

marketing consultancy Hubspot). To build a following, you need to strike a balance between creating regular, interesting updates and replying to those you follow in the hope that they will return the favor.

Twitter is neither one-to-one (unicast) nor one-to-many (broadcast). Call it "sometimescast". This produces strange behaviors. People who are highly followed selectively amplify messages, and people who are less followed send nuggets of wisdom to the more followed in the hopes of getting their attention.

## When Twitter Followers Become Botnets

British author, actor, and technology aficionado Stephen Fry (@stephenfry) has hundreds of thousands of followers. We asked him how being such a Twitter heavyweight affects his online behavior.

*CWM: Is there a burden of responsibility that comes with being a widely followed community member?*

**Stephen Fry:** I have learned to be very careful about revealing frustrations with any individuals, companies, or institutions. I once—almost casually—mentioned an individual who had expressed bizarre antiscience opinions that I found absurd and some followers hounded this poor person with really unkind and abusive emails and posts. That's a worry. It's a tiny minority, but a small percentage of a very large number is a force to be reckoned with.

The same can happen with people who tweet me rudely: if I, as it were, raise my eyebrows at their posts in a public reply their Twitter stream can be looked at by my followers and they can be harried for their eccentricities, politics, or general manner. If, therefore, someone is rude to me, they have to be very rude or unreasonable for me to draw attention to it.

*CWM: What's the most havoc you've (unintentionally) wreaked on a site, now that you're a human denial-of-service attack and Twitter is your bot network?*

**Stephen:** Perhaps justly the most havoc I've wrought has been on a site of my own. In a vain effort to get a revenue stream from the increasing costs of hosting my website I tweeted the arrival of a t-shirt store which carried some "I tweeted @stephenfry and all I got was this luxurious t-shirt" shirtings—within seconds the site was down. I didn't have the nerve to reTweet when it was finally up and running—looked too much like huckstering. "What a waste of time!"—as Bill Murray expostulates in *Groundhog Day* on the subject of French poetry.

*CWM: Counting followers seems awfully "high school yearbook" as a way of measuring online success. What metrics do you care about beyond followers?*

**Stephen:** Yes, the pecker test really won't do. In the early days it was all one had and that was before institutions like TV networks, musicians, and PR machines clambered onto the tweetwagon. Nowadays, I suppose one measures by the levels of genuine interactivity. I think my followers know that it is always me tweeting, that I don't use Twitter to sell anything (with the exception of aborted t-shirts from time to time!) and

that such views, observations, banalities, and quotidian mundanities as issue forth are always my own.

*CWM: What tools can you possibly use to monitor a conversation with hundreds of thousands of followers? Are you forced to resort to hashtags? How do you collect the zeitgeist of a following?*

**Stephen:** I do use hashtags for what one might consider important purposes. I have a hashtag bin (I'm using bin in the film editor's way, not in the sense of trashcan!) for those wanting me to retweet a message or highlight a charitable website. I have another hashtag bin for those who are anxious for me to follow them, and yet another which explains the general ethos of Twitter as I see it.

As to how one can monitor, well, with tens of thousands of new followers arriving weekly I do hope they are able to understand that the chances of me seeing and reading a tweet are very low. It's a simple question of probability.

Imagine a gusty autumn day: I'm standing in a forest as thousands and thousands of leaves stream past me borne on the wind. I snatch at one or two to look at them as they pass. That's my position. There's the option of direct messaging, which increases the chance of me getting a glimpse (which is why people want me to follow them, I guess, as that means they can DM me), but they're coming in pretty thick and fast, too, now.

So long as people understand that it's a kind of lottery as to whether I respond, that's fine. But Twitter is not email, it is not about access and chat. It is, as it tells anyone, about what the tweeter is tweeting. It is "microblogging." I like reading others' tweets and get a little frustrated if all I get are questions (99% of which don't need to be asked, but are clearly a pretext for initiating a conversation—which is cool, but hard to act upon)—but it's not a frustration that matters, and can be set down as the grumpiness of a busy old man.

### Fluid relationships

Twitter's explosive growth is due in part to the fact that you can follow anyone and they can follow you. *It's just as easy to "unfollow" someone*, and we'll soon see more sites that track who's being followed or unfollowed the most. This keeps people honest; there's always a lingering fear of being abandoned by your flock.

The notion of whether you care if you're followed is an elephant in the Twitter room. People want to seem smart. They want the affirmation of reTweeting. They want to be noticed, to be perceived as having authority. Like it or not, Twitter's fluid social graph unearths a yearbook psychology buried deep in our high school psyches, and has more of an impact on our behavior than we think.

### Follower count

Users who follow many people but have few followers themselves are often seen as sources of spam. And those who have many followers but aren't following others are assumed to be simply broadcasting messages rather than interacting. During the 2008


*Figure 13-38. A following: followers' comparison of Hillary Clinton and Barack Obama during the 2008 democratic elections*

Democratic party campaigns, Hillary Clinton's team was criticized for not reciprocating Twitter follows, as shown in a comparison of Clinton and Obama's Twitter profiles in Figure 13-38.

It isn't really possible to read what 146,000 people say (Obama's numbers in Figure 13-38) unless you're using tools for automation. Many Twitter users rely on clients like Tweetdeck to filter traffic from the people they're following, sometimes going so far as to create their own groups of "real" friends they actually follow. This is disingenuous, because it makes people appear to be listening to more people than they really are, but in a world of reciprocal following, it's standard practice. Consequently, *seeing who's mentioning others is a far better indicator of actual social engagement than follower count is*.

#### Limited context

Twitter doesn't thread messages, and the messages themselves are brief. As a result, there's less context for a conversation, since third-party observers don't easily see who's talking to whom. While back-and-forth chats happen, they're usually brief.

Twitter messages tend to remain a mix of instant messaging, URL referrals, bumper-sticker humor, and status updates. *The lack of conversational context stops Twitter from devolving into a party line*.

### Extensible syntax

At the outset, Twitter didn't have much functionality built in (prefacing a member's name with the letter "d" for direct messaging was its only "feature"). This meant that *the community created and selected its own conventions:* reTweeting (prefixing a message with RT), @naming, and #hashtags started out independent of Twitter, and only later became part of the system.

One convention—flagging topics with a hashtag ("#")—is similar to the IRC channel naming convention. When an IRC user talks about a channel called "#channel," he's referring to a chatroom with that name. Similarly, if a Twitter user includes the tag "#memes," she's tagging the message with a particular topic. Users can then converge around that tag for something approximating real-time chat on sites such as Twemes, or form a spontaneous community on thread.io. While hashtags aren't formally part of Twitter, some clients, such as Tweetdeck, will persist hashtags across replies to create a sort of message threading.

Hashtags can also serve as backchannels or a means of circumventing Twitter problems, since they're another way for groups to organize and communicate without being one anothers' followers, as shown in Figure 13-39.

### The constraints of brevity

Twitter's founders attribute much of their success to self-imposed constraints like message size (*www.inc.com/magazine/20080301/anything-could-happen_pagen_2.html*). Developing for a constraint like 140 characters has an important side effect: short messages make traditional brand marketing hard. Spamming a slogan on Twitter is useless, and bad behavior (such as letting companies like Magpie send messages on your behalf) gets you quickly unfollowed.

This has kept the medium relatively spam-free despite the fact that it's so open. Other, richer social networks have to contend with far more spam; *it's hard to hide a virus in only 140 characters* (or is it 255 characters? Or 920? Check out *www.radian6.com/blog/101/learnings-from-twittersecret/*). This limitation also makes users prune and clarify their thoughts, which means your audience can process more ideas with less effort.

### An open API

Twitter has made it very easy to extend its functionality, in part because it hasn't focused on monetizing the system (yet). This means there are quite literally new Twitter applications going live every week. There are also several popular desktop clients and web plug-ins.


*Figure 13-39. Using hashtags to circumvent Twitter's follow mechanism and identify members of a group*

Some of these applications are hashtag trackers, some are tag clouds, and some are ranking tools. There's a vibrant development ecosystem around micromessaging, and it's creating new ways to extract meaning from the thronged masses. These days, even if individual conversations won't get your attention, a groundswell of objections will force you to get involved.

Micromessaging nay-sayers should remember that HTTP was a nascent protocol once, too. It had a few very simple verbs (GET and POST, mostly) and a lot of possibility. Developers and innovators came up with cookies, embedded content, URI parameters, and so on. They were optional, which allowed browsers and web servers to evolve independently rather than having to be in lockstep, resulting in much faster adoption and innovation.

As a communication model, we're still defining micromessaging. If early adoption is any indication, we'll build rich functionality atop it just as we have with HTTP.

The following roles all exist in micromessaging:

- While Twitter is a privately operated service, micromessaging platforms like Laconi.ca are available that can be run independently. The *operators* of these platforms have considerable control (and in Laconi.ca's case, the ability to modify the

platform itself, since it's open source). On the other hand, operators have to convince a community of microbloggers to join them.

- While micromessaging platforms have very little hierarchy, one of the ways that people can rally others around a topic is to create a hashtag that others can use to flag their messages with a particular topic. The person who creates this group is a *group/hashtag creator*. People don't get any special privileges as the creator of the hashtag, but they can track other posts that include it. Some sites also invite hashtag creators to claim a particular hashtag and define it, so there's an advantage to being the first to claim a hashtag for your community.
- For *participants*, the bulk of micromessaging interaction is simply participation—following others and interacting with them. Participants may also follow certain automated services so that those services can track and analyze their content over time.

## Social News Aggregators

With hundreds of millions of people posting their thoughts to the Web, we rely on the wisdom of crowds to decide what's valuable. Social news aggregators let members nominate content they find interesting, allowing others to upvote (promote) it if it's worthy. These sites also encourage commenting on the submitted content, and the comment threads themselves are subject to voting so that the most popular topics rise to the top.

Digg, reddit, and Slashdot are three of the biggest social news aggregation platforms. Each has a community of fanatics responsible for many of the submissions. Other news aggregators, such as StumbleUpon, Twine, and ShareThis, rely more on browser plug-ins that make it easy to bookmark and nominate sites, but are less about member comments on submissions. As URL redirection becomes more commonplace, still more link aggregators will emerge to show which sites web users are sharing the most within social networks and micromessaging platforms.

If your website gets noticed by a social news aggregator, you'll get a huge number of visitors in a very short period of time. This means a heavy burden on your servers. To make matters worse, few of the new visitors attracted by your newfound popularity will have been to your the site before, so they will all be retrieving every object from your servers (since none of it was cached on their browsers). The effect will often break a website entirely, as shown in Figure 13-40.

News aggregators have different reputation systems. On some systems, such as reddit, submitters receive "karma" for popular submissions, but that karma has no effect on a person's votes; it's simply a mark of community respect. On other sites, past performance gives more power: a well-regarded Digg user's votes count for more.

This is a controversial matter, because reddit's one-voice-one-vote means the site is less prone to having a few very influential users—something that has plagued Digg—but

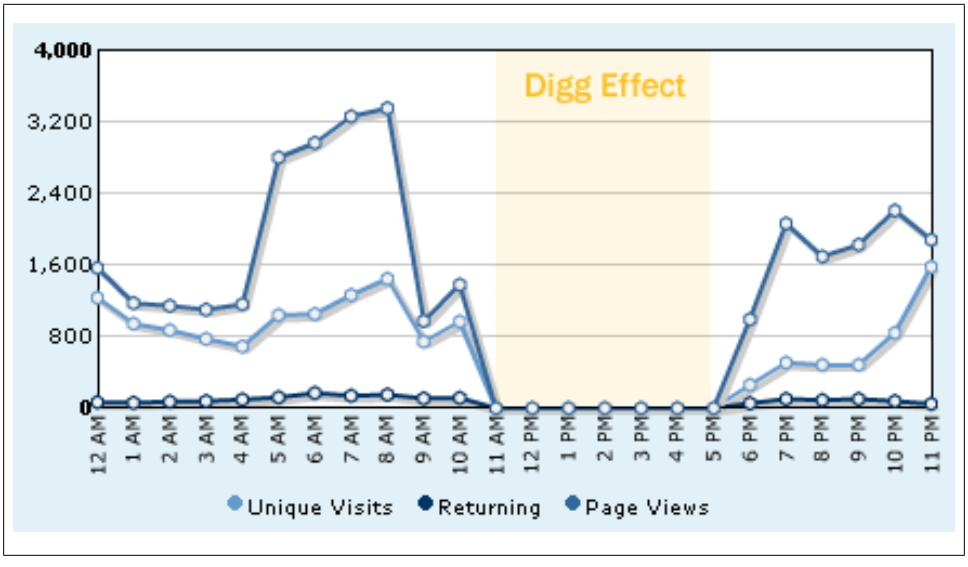


*Figure 13-40. The "Digg Effect" on Nick La's N.Design Studios portfolio site—after a sudden burst of interest, this site became inoperable*

spammy content is harder to filter out automatically on a site like reddit where every submission, whether from a one-time spammer or a veteran, has the same initial merit.

These are the roles in social news aggregation:

- *Operators* may run platforms of their own on software such as Pligg. Because the power of aggregators comes from the collective wisdom of a large community, however, these will probably be part of an intranet.
- A *group creator* creates a community of interest within an aggregator (on reddit, this is called a subreddit). Group creators don't get additional visibility into the communities they create, but they will have a place to unite like-minded contributors.
- When *submitters* submit links, those submissions become part of their personal histories, and upvotes for something they've submitted contribute to their reputations. Submitters are expected to explain why the submission is interesting and summarize it for others. If submitters want the subject to get attention, they need to make the summary compelling, controversial, and brief.
- *Commenters* comment on submissions, and the comments themselves are subject to voting by the community. Commenters may receive some kind of ranking for content that rises to the top.
- *Voters* alter the ranking of content and comments.


*Figure 13-41. A community site built entirely in Drupal that includes many social network features*

## Combined Platforms

Major social networks have incorporated many of the communications models we've seen above. Facebook, for example, has real-time chat, open message threads (walls), one-to-one messaging, groups, mailing lists (messages to groups), and so on.

On the private side, servers like Drupal combine blogging, wikis, and other components in a single platform that many developers have extended with plug-ins and new functionality, as shown in Figure 13-41.

## Why Be Everywhere?

Clearly, there's a daunting variety of community platforms on the Internet. Can you just pick a few?

It depends.

If you have a specific audience with a strong preference for one type of community, one platform may be enough. An extremely technical, always-connected set of server administrators, for example, may be content with just an IRC channel. But if you're trying to engage a broader market, particularly for consumer products or public causes, you need to be on all of the dominant platforms. Here's why:

*Reach*
: By being active on many platforms, you increase the chances that your messages will be picked up by the community and amplified.

*Early days*
: It's still hard to know which social platforms will dominate the Internet, and you need to hedge your bets. Will Twitter stand on its own? Will it fracture into several federated micromessaging platforms? Or will micromessaging simply become a feature of social network sites? Because it's too soon to tell, you need to engage on all of them.

*Different platforms for different messages*
: A last-minute announcement about a party goes best on Twitter, while a detailed list goes in a blog posting, and a question looking for responses goes to a mailing list.

*Awareness*
: You never know when you're going to become the subject of discussion, so you should be monitoring as broad a range of sites as possible in case conversations about you or your competitors emerge.

*Multiple audiences*
: You engage with multiple people in an organization. The accounting department may not use the same social networks as the executive team, who may work with different tools from the folks in support. Different audiences gravitate toward different platforms.

The media you choose strongly influences the messages that are discussed, the people who engage, and the detail of their responses. Real-time platforms require quick thinking, wit, and a thick skin. Micromessaging tools demand bumper-sticker cleverness, pithy quips, and an understanding of shorthand conventions and abbreviations. Blogs have to engage visitors with a single thought early on to capture attention. Wikis need deeply linked, comprehensive, defensible content.

## Monitoring Communities

Regardless of the community platform, several fundamental building blocks keep emerging. These include the following:

*User reputation*
: This is a score of how well respected you are within the community.

*Threading*
: Threading is a way of organizing replies and responses between community members, often with comments.

*Voting by the community*
: This helps point out what's good, in the form of reTweeting, upvoting a submission, or flagging an inappropriate addition.

*Notifications*
: These let members know about what's happening within the community and keep community members engaged while fighting attrition.

*Lifestream*
: A lifestream provides status updates on community members for those who choose to follow them.

*Social graph*
: A social graph describes the relationship of members to one another, either in a whitelist (Facebook-like) or blacklist (Twitter-like) model.

*Searching*
: You can search across historical interactions, often using tags and other filters.

*Topics, subgroups, and events*
: These are groupings on which a community can converge.

*Befriending and unfriending*
: This can change the structure of the social graph.

These building blocks can all be measured and analyzed. They represent the KPIs and metrics of community monitoring. In the next chapter, we'll turn our attention to how to monitor them.