

---

# Tracking and Responding

Engaging a community is one thing; understanding the results of those engagements is another entirely. If we're going to link community activity to business outcomes, we need to track conversations and respond when appropriate.

Our ability to track what's said online varies widely depending on where that conversation is taking place. Some community platforms, like Twitter, are wide open, myriad tools for analysis and search. Others, such as Facebook, are firmly closed—third parties are forbidden from extracting information and publishing it elsewhere, and while the company is putting some carefully designed gates in its walled garden, Facebook can't easily be crawled, curtailing the number of community management tools that can work with it.

It's not just Facebook. StumbleUpon, shown in [Figure 14-1](#), shows a promotion for marketers when you follow a referring link from its site. There are ways around this, primarily by installing and using the StumbleUpon toolbar, as described by Tony Adam at <http://tonyadam.com/blog/checking-stumbleupon-referrals-reviews-urls>. For community operators, there's a constant tension between the desire to be open (with the community development, rapid growth, and innovation this entails) and the desire to monetize the community somehow by controlling who has access to information on its activity.

The way in which you engage a community will depend on how much visibility you want into that community's activity. In the coming pages, we'll look at four distinct levels of engagement you can have with a community.

- You can *search* the Web using tools like Google, or search within a specific community using the platform's search functions. In this case, you're simply a passive observer, and this is a hands-off approach. It's also the one that yields the least information, because you can only search what you're allowed to see as an outsider.
- You can *join* communities, seeing what's happening from within. To do this, you must get an account and reciprocate interactions (follow users who follow you, comment on others' posts, or join mailing lists they create).

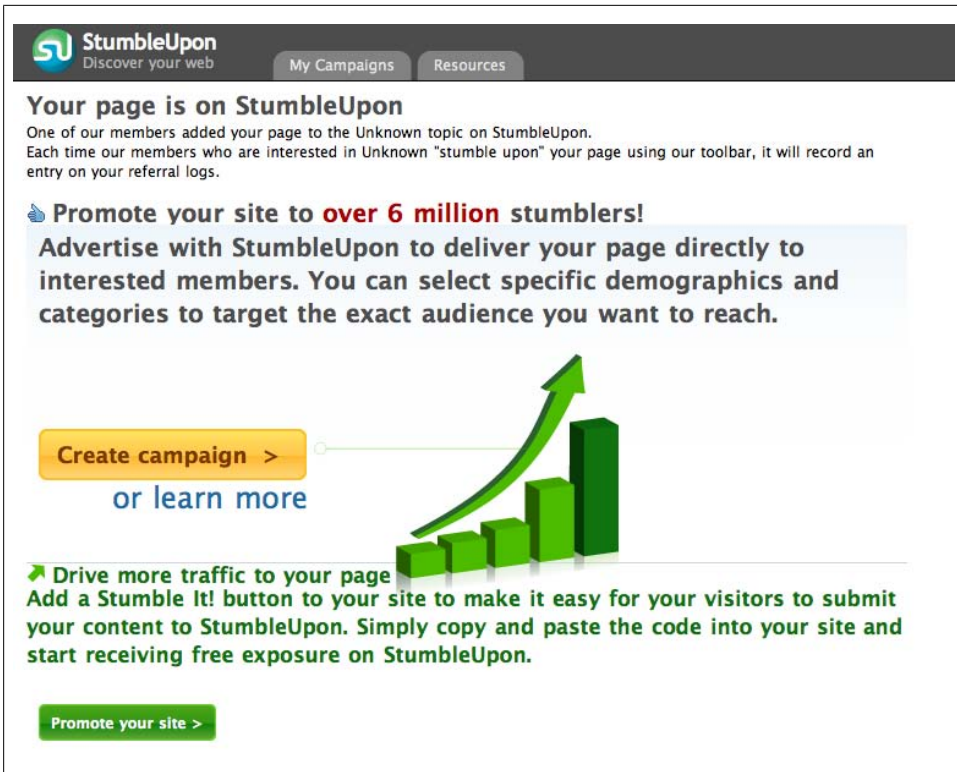


Figure 14-1. StumbleUpon's referral landing page doesn't show you which content drove traffic to your site

- You can set up and *moderate* a community within an existing platform, inviting users to join you while remaining responsible for the operation of the group. This may be a Facebook group, a Yahoo! mailing list, a subreddit, or even a hashtag you define on sites like thread.io.
- You can *run* your own community platform, giving you complete control, but requiring that you convince others to join you. This works best when others have a reason to join you, for example, when they want product support or access to people within your organization.

Figure 14-2 shows these four levels of engagement, along with the degree of commitment required, the level of analytical detail you can hope to extract, and the number of existing users you'll find with the community when it launches.

In the preceding chapter, we looked at eight distinct community models: groups and mailing lists, forums, real-time chat systems, social networks, blogs, wikis, micromessaging tools, and social news aggregators. By combining these eight models with the four levels of engagement, we can put together a community monitoring strategy.

	Search	Join	Moderate	Run
Commitment	Low	Low	Medium	High
Analytical detail	Low	Low	Medium	Complete
Existing users	All of them	Many	Few	None

Figure 14-2. What you put in and get out of various levels of community engagement

	Search	Join	Moderate	Run
Group/mailling list	Mailing list search tools	Subscribe to Yahoo Group	Set up Yahoo Group	Operate a Listserv
Forum	Forum search tools	Post in online forums	Start a thread in a forum	PHPBB
Real-time chat	N/A	Post on an IRC channel	Create an IRC channel	Operate an IRC node
Social network	Search Facebook	Subscribe to Facebook	Create a Facebook group	Drupal
Blog	Blog search tools	Comment on a blog	Contribute to a blog	Wordpress.com installation
Wiki	Wiki search tools	Contribute/edit Wikipedia	Create a Wetpaint wiki	Mediawiki
Micromessaging	Twitter search sites	Get a Twitter account	Create a group or hashtag	Indenti.ca installation
Social news aggregators	Reddit, Digg search	Vote & post content	N/A	Pligg

Figure 14-3. Examples of engagement across each of the eight community models

Figure 14-3 provides some examples of how we might search, join, moderate, or run each of the eight types of community.

## Searching a Community

Web crawlers are constantly crawling the public Web, indexing what they find into dozens of online services, such as search engines, plagiarism detectors, alerting tools, or social sentiment collectors. Search is your main way to find conversations that concern you, especially when you don't actually want to join or run the communities yourself. There are two kinds of search: traditional manual searches, known as

The screenshot shows the Google Alerts (BETA) page. At the top right, there is a user email 'spower@gmail.com' and links for 'Settings', 'FAQ', and 'Sign out'. The main heading is 'Welcome to Google Alerts'. Below it, a paragraph explains that Google Alerts are email updates of the latest relevant Google results. A list of 'Some handy uses of Google Alerts include:' follows, with bullet points: monitoring a developing news story, keeping current on a competitor or industry, getting the latest on a celebrity or event, and keeping tabs on favorite sports teams. To the right is a 'Create a Google Alert' form. The form has a title bar, a text input for 'Search terms' (containing 'social network'), a 'Type' dropdown (set to 'Comprehensive'), a 'How often' dropdown (set to 'as-it-happens'), and a 'Deliver to' dropdown (set to 'Feed' with 'spower@gmail.com' listed below it). A small disclaimer at the bottom of the form states 'Google will not sell or share your email address.' Below the form, a link says 'You can also click here to manage your alerts'. At the very bottom, a footer contains copyright information and links to 'Google Home', 'Google Alerts Help', 'Terms of Use', and 'Privacy Policy'.

Figure 14-4. Setting up a Google Alert for a prospective search

*retrospective searches* because they look for data that’s already been published, and alerts, known as *prospective searches* because they let you know when new data appears.

You need to search across all sites that matter to you—industry publications, analysts, and readers. Once you find a search that’s useful, store it so you can run it again. You can store all of the searches as bookmarks, so if you want to share them with your team, consider a social bookmarking site like Delicious.

Stalking yourself daily with search engines might be fun and flattering, but it gets old fast. Fortunately, you can get alerts to notify you when content that you care about appears, using prospective search. To do this, set up a search term in advance using tools like Google Alerts (shown in [Figure 14-4](#)). In this case, Google will let you know whenever that term is seen.

You can also receive updates via RSS feeds. While these will only pull in content from postings and feeds that are published, anyone tasked with monitoring a community needs to run an RSS feed reader to keep track of influential posters or relevant posts.

Finally, you can also use alert services that check web pages for you and let you know when their content changes, such as WatchThatPage.com (shown in [Figure 14-5](#)) or [ChangeDetection.com](#). Then you can dig deeper to see if the changes matter.

## Searching Groups and Mailing Lists

Mailing lists that are archived online may be searchable without your participation if they’re public. Yahoo! Groups, shown in [Figure 14-6](#), lets you view archives and search

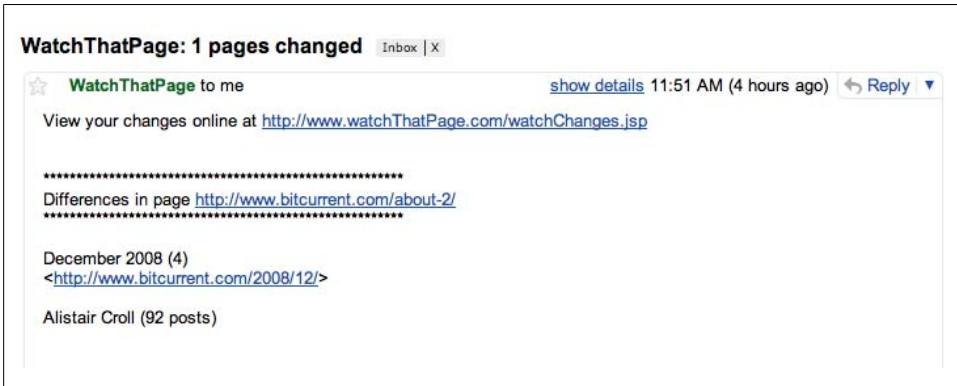


Figure 14-5. A report from [WatchThatPage.com](http://www.watchThatPage.com) showing changes to a monitored website

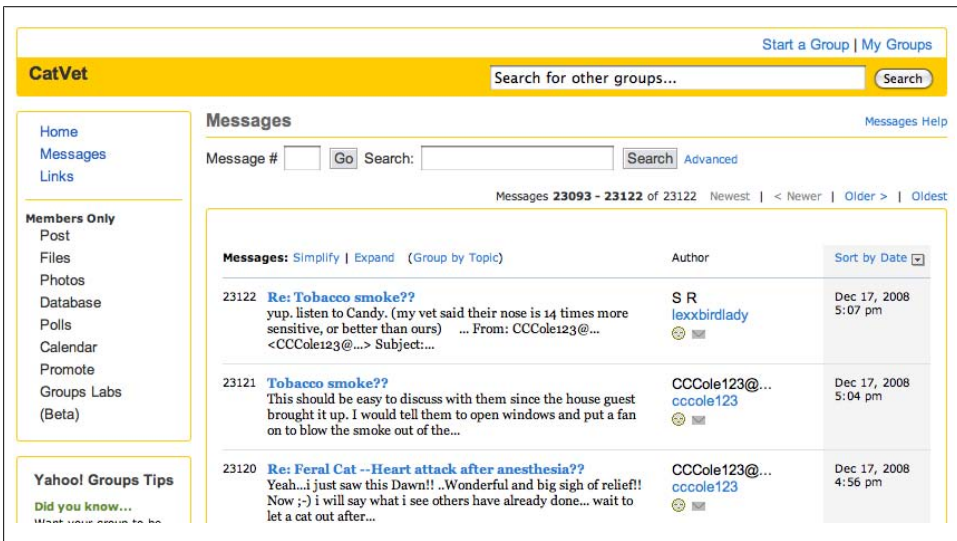


Figure 14-6. Searching the archives of a Yahoo! group

for strings in many mailing lists for which the administrator has made the archives available.

You can also use a prospective search portal like Rollyo to search for specific topics within the constraints of specific sites that contain mailing list archives, as shown in [Figure 14-7](#).

Some mail tracking tools index mailing lists. Markmail ([Figure 14-8](#)), for example, indexes roughly 500 Google Groups lists (about 3.8 million email messages) and thousands of other mailing lists, and makes them searchable through a web interface

that also provides some keyword analytics showing which mailing lists and authors mention a topic most often.

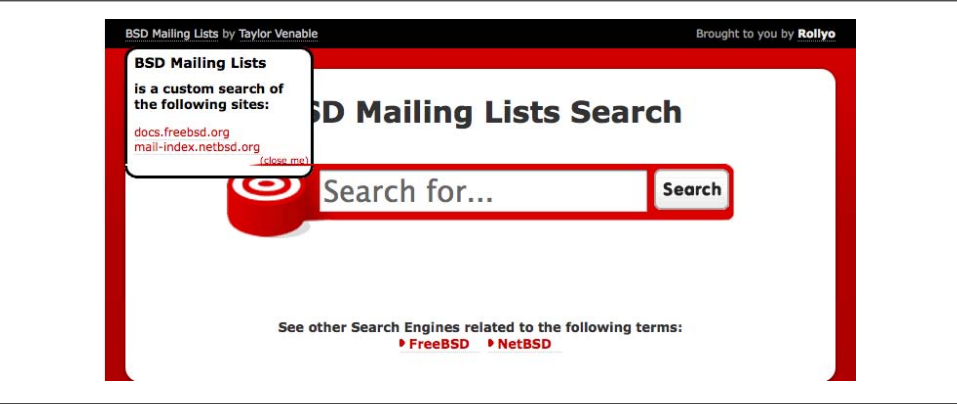


Figure 14-7. A prospective search query on Rollyo for the NetBSD and FreeBSD mailing lists

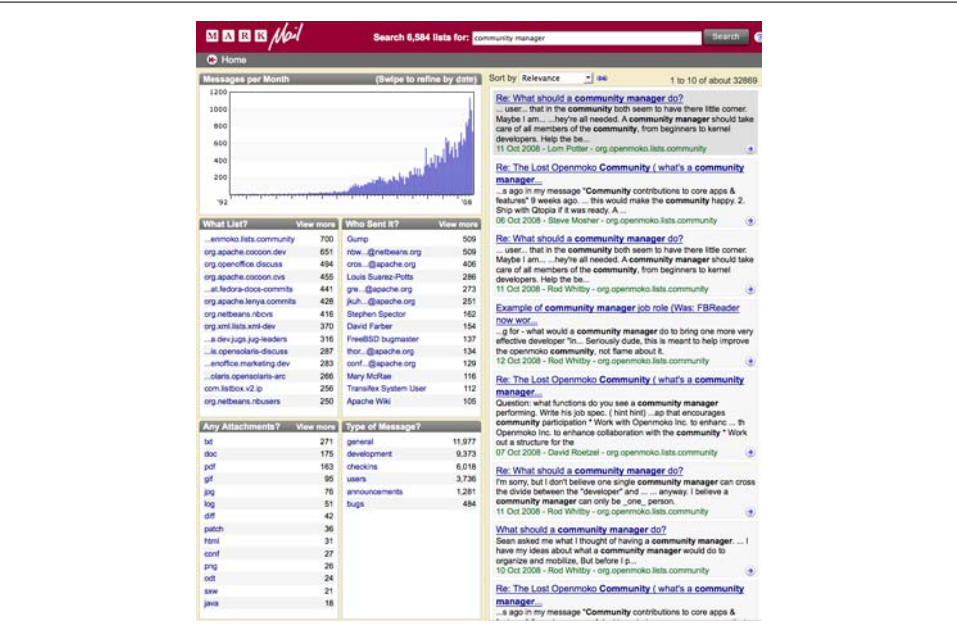


Figure 14-8. Searching mailing lists with Markmail

In addition to finding topics and content, Markmail reports on how popular specific topics or keywords are becoming by graphing the percentage of messages that contain a particular term (Figure 14-9).

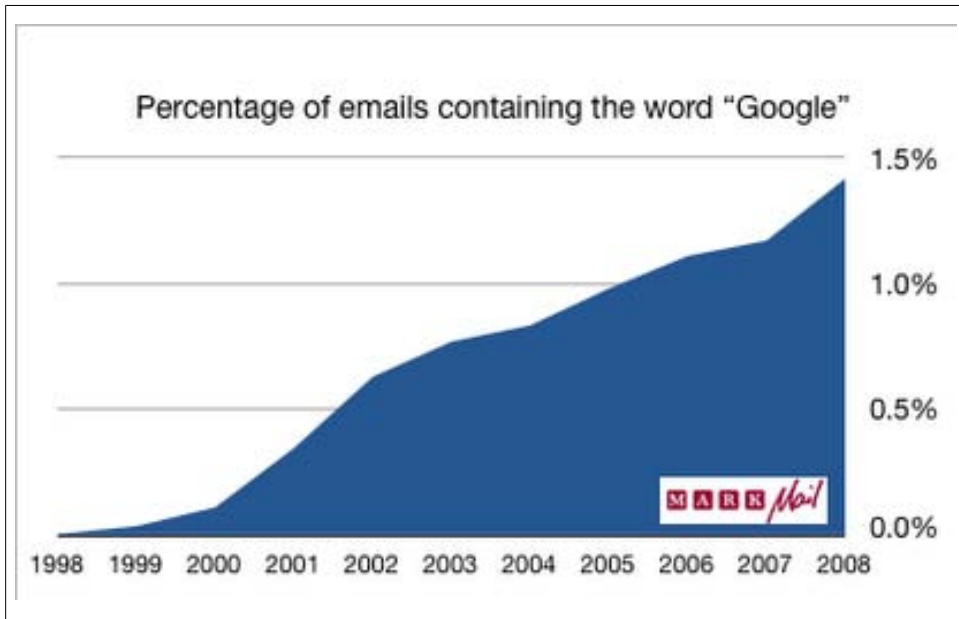


Figure 14-9. Percentage of email messages containing the word “Google” across 3.6 million email messages in 10 years

## Searching Forums

Public forums are often indexed by web crawlers, so you may be able to search them by narrowing the scope of a Google search to just the forum in question. Using Google’s advanced search, shown in Figure 14-10, you can constrain a search to a single site (in this case [forums.worldofwarcraft.com](http://forums.worldofwarcraft.com)) and find specific keywords.

Forum ranking sites, like Big Boards and Boardtracker, index many forums on the Web and may also help you to find the forums you should be tracking (Figure 14-11).

Once you’re on a forum, you can use its internal search tools to look for specific content. But many forums require you to be a member, so you may be forced to join the group, forum, or mailing list in order to access its search features or to view archives.



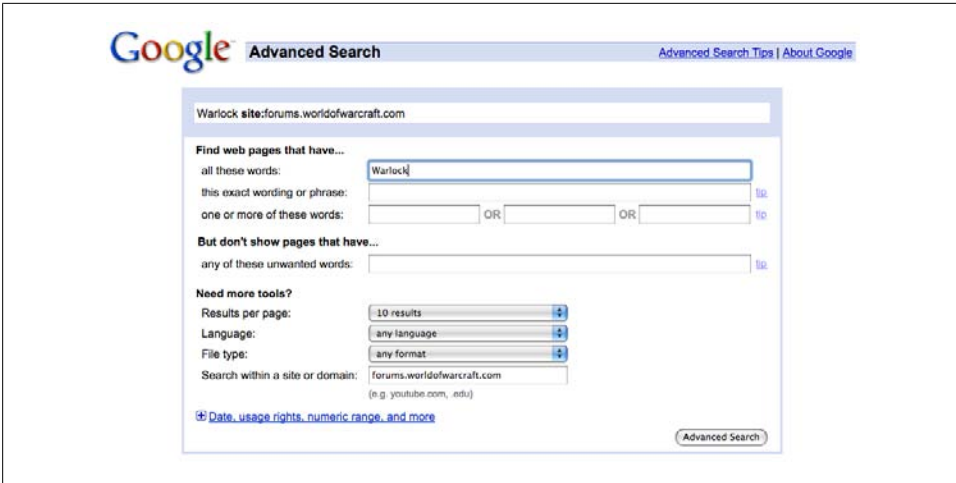


Figure 14-10. Using Google's advanced search to find content on a specific forum

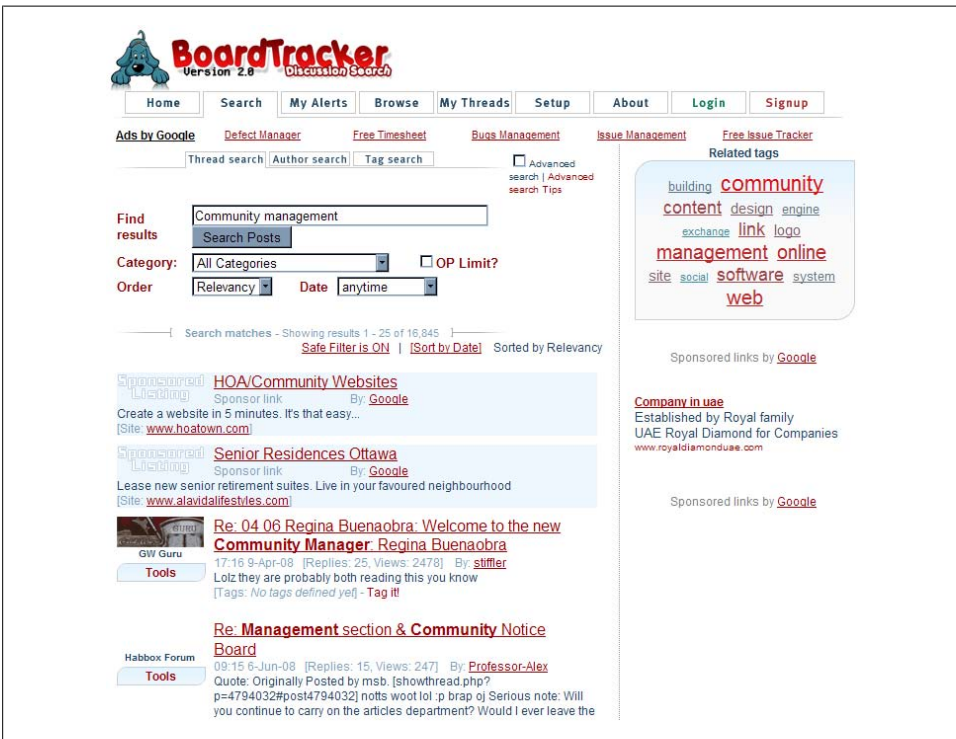


Figure 14-11. BoardTracker searches online forums for specific content



The screenshot shows the SearchIRC website interface. At the top, the logo "SearchIRC" is displayed with a magnifying glass icon. Below the logo, a navigation bar includes links for "Home", "Networks", "Community", and "Need Help?". The main search area shows a search for "Web performance" with 370,762 channels found. The results are displayed in a list format, showing the first 15 results (of 1,686) matching the search term. Each result includes the channel name, network, review status, and a freshness factor. The results are as follows:

Channel Name	Network	Review	Freshness Factor
irc:// #ads Established	GameSurge	There are no reviews of this channel	99.5
irc:// #advertise Established	GameSurge	There are no reviews of this channel	99.5
irc:// #performance	GIMPnet	There are no reviews of this channel	97.8
irc:// #Web	wildirc.net	There are no reviews of this channel	99.4

On the left side of the page, there is a sidebar with a "Please visit" section containing links to "Alabama Crimson Tide", "Pool Cue Case", and "Current events discussion".

Figure 14-12. SearchIRC indexes and searches content on selected IRC channels

## Searching Real-Time Chat Systems

Real-time chat isn't archived in the same way that groups and forums are. To "crawl" an IRC channel, an indexing service requires a bot that listens to traffic on the channel and indexes it. If you don't feel like joining a channel yourself and indexing content, you can use sites like SearchIRC, shown in Figure 14-12, to look for specific terms.

IRC isn't the only chat platform on the Internet, of course. Many instant messenger platforms today (such as Skype and MSN Messenger) give users the ability to create ad hoc group chats, but if you're not party to the conversation, you can't track it.

There are some public conversations that resemble chats. Twemes, for example, strings together micromessages around a specific topic using a hashtag (Figure 14-13 shows the Twemes analysis of the #food hashtag). You can stream the content of a topical discussion to an RSS feed and examine it later, or use Twemes itself to search a topic's conversation.

## Searching Social Networks

Social networks come in two flavors: open models and walled gardens. Simply put, you can't browse the walled ones from outside their walls.



Figure 14-13. Twemes looks like IRC or near-real-time chat, but is based on Twitter

A social network is only as valuable as its social graph. So some sites, such as Facebook, LinkedIn, and MySpace, have guarded them jealously (though they've opened up considerably in the face of rapid growth from competing communities like Twitter). These sites don't let you view profiles or navigate relationships unless you're logged in. And even if you're logged in, it's a violation of their terms of service to try to map the social graph with outside mechanisms. In other words, you need to be a Facebook member to access your social network, and even then you can only do so *within* Facebook's website.

Facebook does provide some data on popular topics, using Lexicon analytics, shown in Figure 14-14. Lexicon measures the popularity of terms within the public portions of Facebook, such as wall posts and status updates.

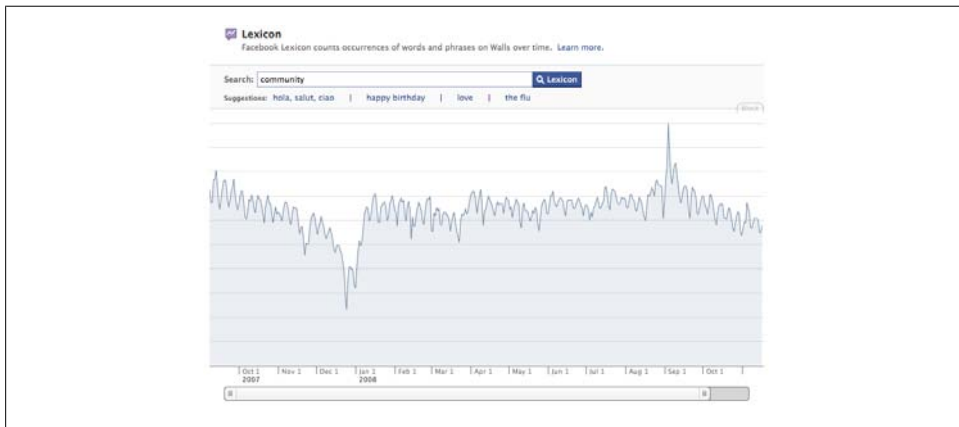


Figure 14-14. Facebook Lexicon shows occurrences of particular words over time within the Facebook Walls application

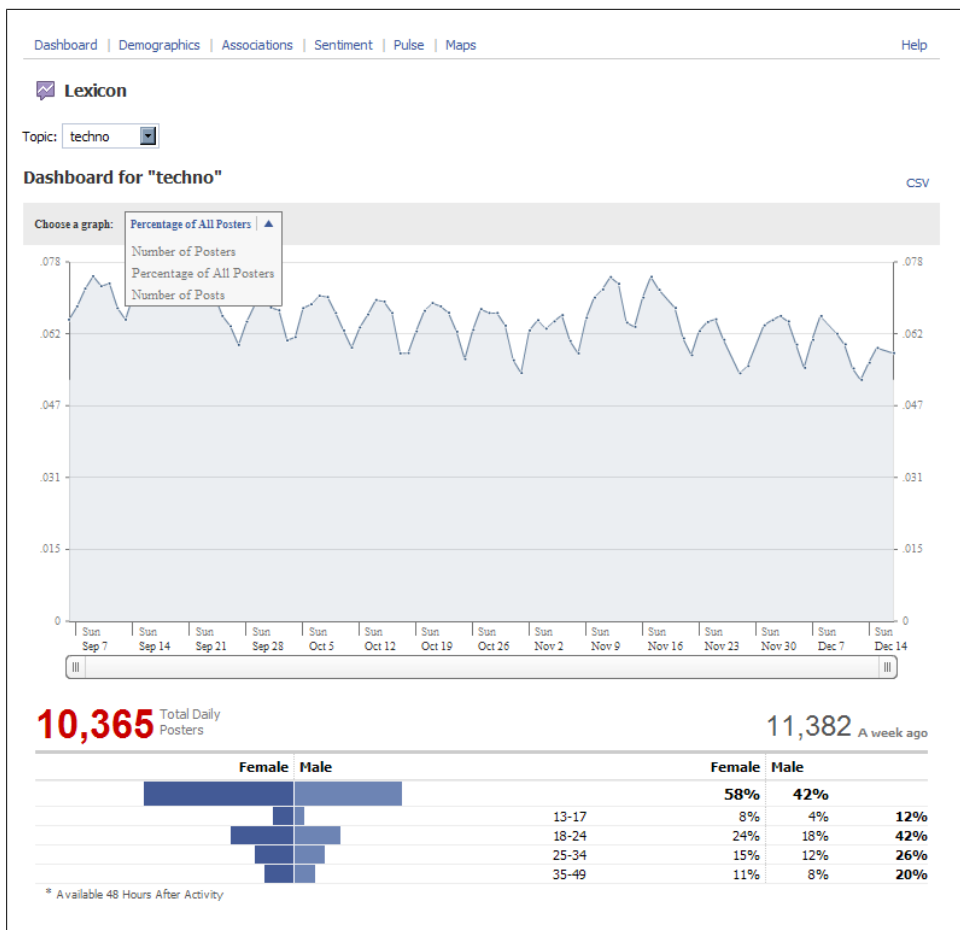


Figure 14-15. The public view of Facebook's Lexicon only indexes a few terms and doesn't offer the capability to analyze terms you specify

However, Facebook doesn't provide much detail within this model. You can't see who's talking or what they're saying. And they provide deeper analysis only for specific terms they dictate, as shown in Figure 14-15. Facebook's advertising platform offers additional visibility for paying customers, as the company is sitting on a motherlode of public sentiment data.

By contrast, more open social networks (most notably Google's OpenSocial model and those based on OpenID) make it easier for anyone to browse a social graph through a set of well-defined APIs.

There's one exception to this walled garden limitation, however. Business-oriented social networks are only partly walled. Sites like LinkedIn and Spoke have to be somewhat open, since one of their main functions is to help their users promote their profiles

and extend their business networks. This means you can usually search and view profile information on a site like LinkedIn without being logged in.

## Searching Blogs

Blog content is fairly persistent (unlike chat), seldom password protected (unlike groups), and open to everyone (unlike social networks). Consequently, most blogs are indexed by search engines and you can track them using the prospective and retrospective search methods outlined earlier. You can also collect and archive blog content via RSS feeds or mailing list subscriptions, and many blogs have their own search functions that let you find content by keyword or author.

Google's Blog Search service (<http://blogsearch.google.com>) is a good place to start. Sites such as Technorati and IceRocket also rank blogs and categorize them by topics, providing a good way to identify websites that cover topics in which you're interested.

## Searching Wikis

Wiki platforms are searchable, as shown in [Figure 14-16](#), and provide tools for tracking content you care about, letting you know by email or RSS feed when something has been edited.



Figure 14-16. Searching on Wikipedia

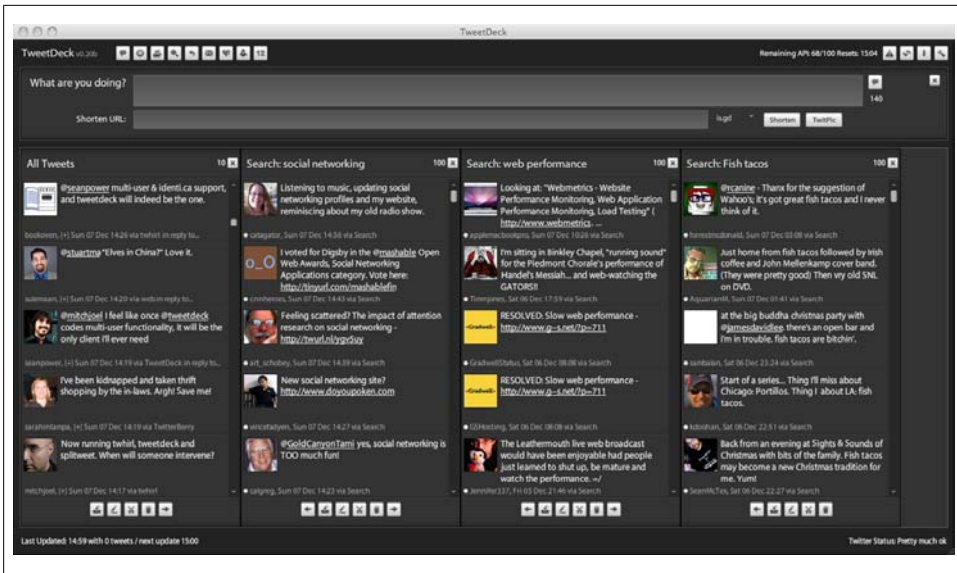


Figure 14-17. Tracking several live search results with Tweetdeck

Wikis are also indexed by search engines, but if a particular topic interests you, it's a good idea to subscribe to changes to that topic, because a heated debate on a particular subject may mean that the topic changes more frequently than the search engines can index it.

## Searching Micromessaging Tools

Twitter is notable for its open APIs, which make searching and monitoring specific topics relatively easy. Many desktop Twitter clients let you search Twitter for content. Tweetdeck, shown in Figure 14-17, can maintain multiple search windows that are refreshed periodically.

In addition to desktop clients, Twitter's open APIs have also given rise to a huge number of other sites for searching and analyzing micromessaging content (for a list, see <http://mashable.com/2008/05/24/14-more-twitter-tools/>), so you can use these to see what's going on, even without an account. Sites like Twitscoop, shown in Figure 14-18, provide visual information on popular Twitter topics.

Twitscoop can also show the popularity of specific topics in recent days, as shown in Figure 14-19.

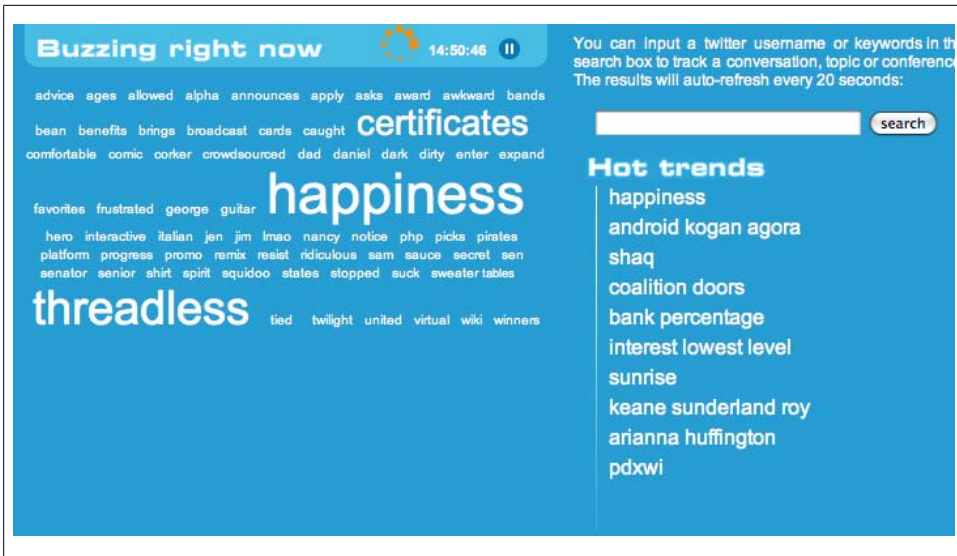


Figure 14-18. Twitscoop provides tag clouds of specific topics within Twitter, segmented by username or keyword

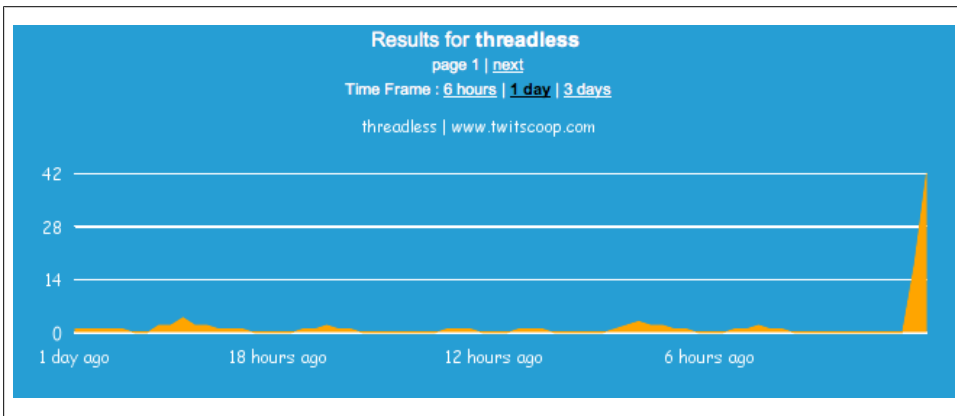


Figure 14-19. Twitscoop's view of topic mentions over time

Tweetstats, shown in Figure 14-20, can show how prolific a user is, who he talks with, and what systems he uses to send messages.

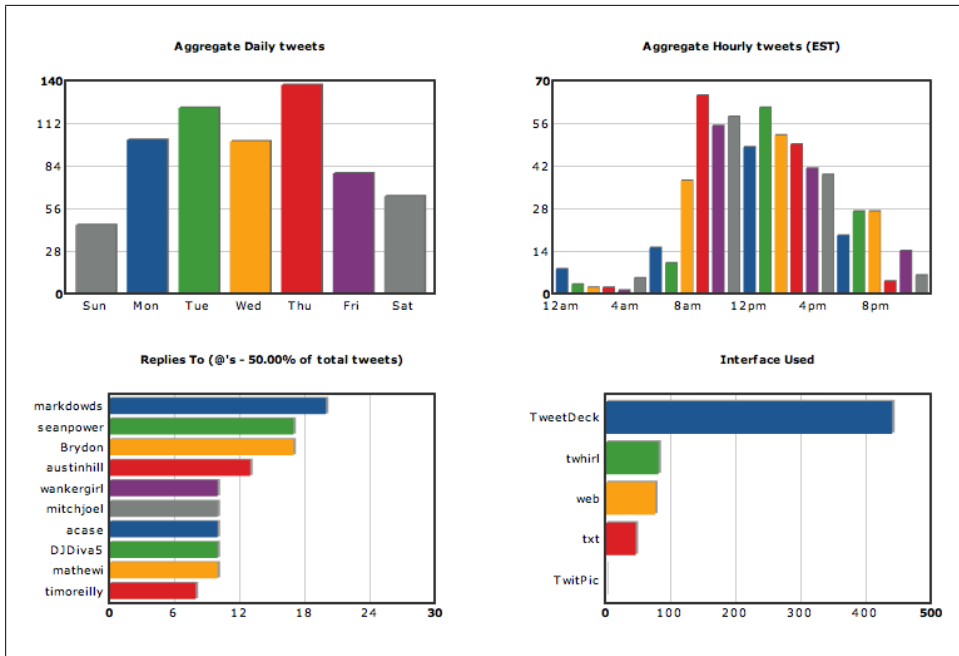


Figure 14-20. Tweetstats report on a Twitter user's volume of messages

Tweetstats also shows a profile of when a user is most active, shown in [Figure 14-21](#). If you're trying to get noticed, this kind of analysis can give you an idea of when your target audience is most likely to be on Twitter and therefore to see your message.

If you want to find topics and posters, you can search for them within Twitter or on directories such as Twellow, shown in [Figure 14-22](#).

Unlike social networking sites, micromessaging platforms are so open that there's little you can't analyze, even when you don't have an account. If you join a social network, you get much more visibility, but if you join a micromessaging platform that has open APIs, the only additional visibility you'll gain is detail about your personal profile.

## Searching Social News Aggregators

Social news sites that rely on upvoting are important sources of sudden attention, and their content isn't easily crawled by search engines, because it changes rapidly. You'll need to use built-in search interfaces (shown in [Figure 14-23](#)) or follow referrals to find content that relates to you. The good thing about this is that it will show you not only stories, but also comments and other content you should be tracking.






Figure 14-21. Tweetstats report on patterns and timing of Twitter use



Figure 14-22. Twellow lets you browse a directory of topics, then see who's discussing the topic most actively

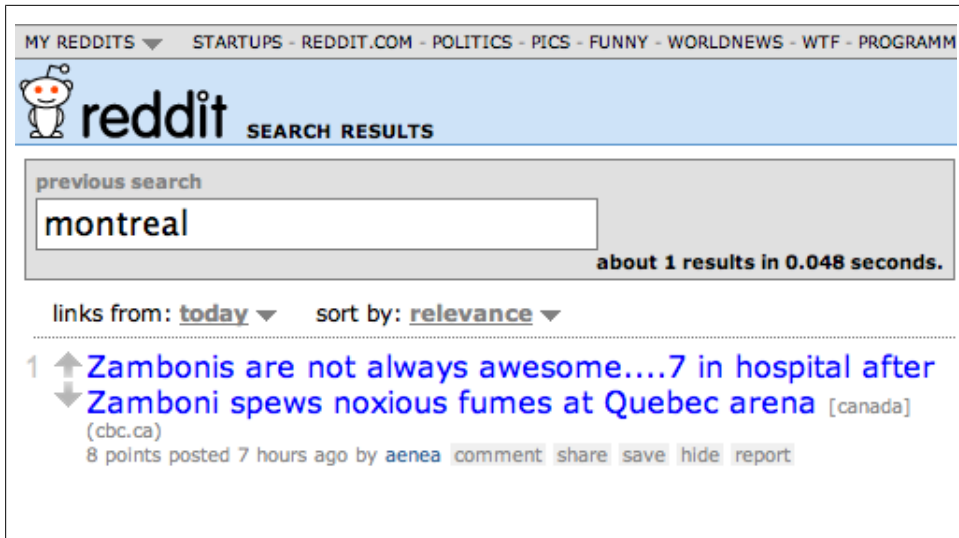


Figure 14-23. Searching social news aggregator reddit using built-in search tools

## Cross-Platform Searching

With so many communities generating content, sites like BuzzFeed assemble a “river of news” that shows messages across many communities. You can search this feed. Unlike Google, which focuses on slower, more comprehensive search, sites like BuzzFeed and Serph update themselves more frequently, but don’t try to index the entire Internet.

You can also use keyword popularity tools like SiteVolume, shown in [Figure 14-24](#), to find and compare topics against one another across multiple platforms. This will show you how well you rank against competing terms, companies, or individuals.

By keeping a close eye on sites that refer to you, you can add new places to search. Maintaining a comprehensive list of social sites that should be searched on a regular basis and keeping prospective search alerts up-to-date is a key job for community managers.

Despite all of these tools, you need to join a community. Not just because it will give you greater visibility—though that’s reason enough—but because the whole purpose of communities is to build interaction with your audience.

## Joining a Community

To really analyze and track communities, you need to marinate in them. This means getting an account, connecting with others, and analyzing the results. Often, you’ll be rewarded with greater visibility as a result of signing up.

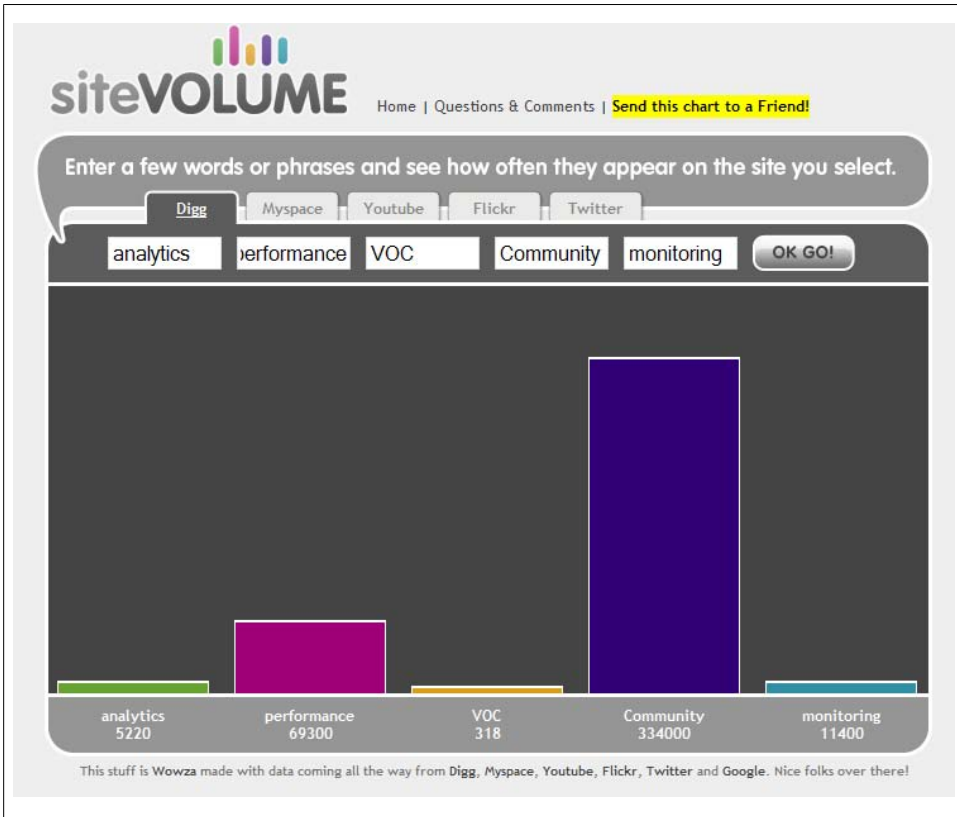


Figure 14-24. SiteVolume compares the popularity of mentions across several social platforms

When you join a community, you'll have little control over the conversations that take place and you'll be at the mercy of the tools that the community makes available. Once you're a part of the network, however, you can mine your own social graph and use internal tools such as groups to better understand your audience. Many of these tools are still rudimentary and are often constrained by how much data the platform owners are willing to share with users.

Unlike searching, which leaves no trace, becoming a part of the community has consequences. When you sign up, others will know you're there. You need to decide whether you want to identify yourself as a member of your organization (we strongly recommend that you do—these things have a way of leaking out) and understand what legal ramifications your participation will have. For example, if you work for a car manufacturer, you may be subject to lemon laws that require you to respond to complaints—so tread carefully.

**Choose search criteria** Specify the criteria you'd like to use for determining what to do with a message as it arrives. Use "Test Search" to see which messages would have been filtered using these criteria. Messages in Spam and Trash will not be searched.

From:

To:

Subject:

Has the words:  Doesn't have:

☐ Has attachment

[Show current filters](#)

---

**Choose action** - Now, select the action you'd like to take on messages that match the criteria you specified. When a message arrives that matches the search: list(<webanalytics.yahoo.com>), do the following:

- ☐ Skip the Inbox (Archive it)
- ☐ Mark as read
- ☐ Star it
- ☒ Apply the label:
- ☐ Forward it to:
- ☐ Delete it
- ☐ Never send it to Spam

[Show current filters](#)    ☐ Also apply filter to 3000 conversations below.

Figure 14-25. Defining custom filters in Gmail to label inbound messages

## Joining Groups and Mailing Lists

For groups and mailing lists that are members-only, you can join through a mailing list management portal, after which you'll be able to read messages. Depending on how the moderator has configured the mailing list, you may have access to historical archives as well.

If you join many mailing lists and groups that are relevant to your community, you can use webmail as a form of prospective search. Using Gmail, you can append topic labels to a username (in the form *username+topic@gmail.com*) and have messages automatically tagged when they arrive, making it easier to search them later. To do this, first send all of the mailing list content to a dedicated Gmail inbox and create custom filters that label messages according to their groups (Figure 14-25).

Next, create a second set of filters that watches for specific topics and either flags them for review or forwards them to your primary mail account, as Figure 14-26 illustrates.

With a bit of work, you can create your own searchable archive of relevant mailing lists and define your own alert system to handle thousands of messages a day, letting you search many mailing lists with a single query and archiving list content indefinitely.

## Joining Forums

While some forums are indexed by web search engines and let anonymous visitors search for content, some require that you be a member of the forum in order to see content.

Once you're logged in, you can interact with other participants and engage the community that's already in place. Web forums are relatively mature and have good

**Choose search criteria** Specify the criteria you'd like to use for determining what to do with a message as it arrives. Use "Test Search" to see which messages would have been filtered using these criteria. Messages in Spam and Trash will not be searched.

From:   
 To:   
 Subject:

Has the words:   
 Doesn't have:   
☐ Has attachment

[Show current filters](#)

---

**Choose action** - Now, select the action you'd like to take on messages that match the criteria you specified. When a message arrives that matches the search: **[webanalytics] "social media"**, do the following:

☐ Skip the Inbox (Archive it)  
☐ Mark as read  
☐ Star it  
☐ Apply the label:   
☒ Forward it to:   
☐ Delete it  
☐ Never send it to Spam

[Show current filters](#)    ☐ Also apply filter to 32 conversations below.

Note: old mail will not be forwarded

Figure 14-26. Defining actions for messages in Gmail for archiving and management

notification features, letting you know when specific topics are discussed, as well as allowing you to subscribe to content via email updates or RSS feeds.

A more mainstream kind of forum is a review site. These sites are a cross between forums and blogs, and often revolve around ratings systems. Yelp and TripAdvisor are examples of these. When you join review sites, you face a dilemma: if you identify yourself, others may be able to see your actions and accuse you of falsely inflating scores. On the other hand, if you choose to use a pseudonym and are discovered, you may face accusations of dishonesty. Most of these sites allow you to view their content anonymously, so you don't need to join them to see content and ratings.

## Joining Real-Time Chat Systems

With IRC, you simply subscribe to a channel in order to analyze it. Some channels may be available only to members, and if you want to join them, you'll need permission from the moderators. Once you're in, you use the same bot tools to eavesdrop on what's said and analyze the conversation.

## Joining Social Networks

As we've seen, social networks are hard to measure from the outside because of their walled garden design and their restrictive terms of service. Even once you're inside the walls, the tools aren't very flexible and are usually limited to your immediate social network. Some Facebook applications, such as Socialistics (shown in [Figure 14-27](#)), provide you with rudimentary information about your social graph.

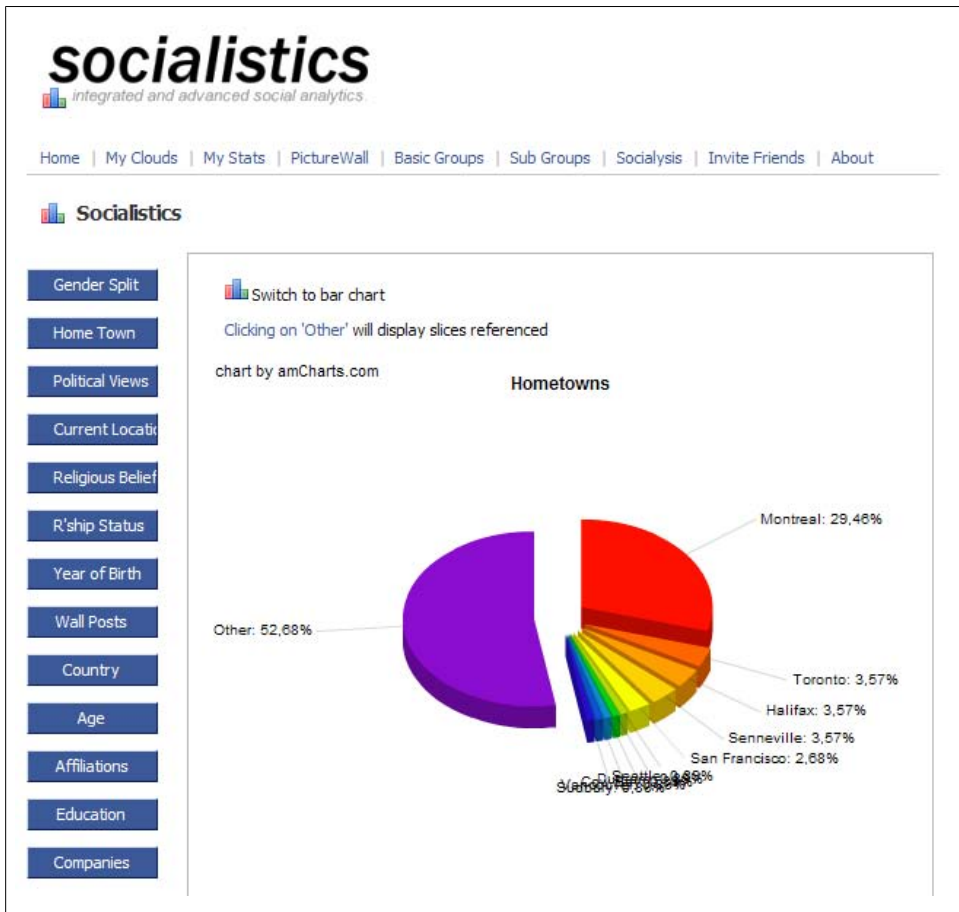


Figure 14-27. Socialistics is a Facebook application that provides basic segmentation of your social graph

The only way to find out what's going on within a social network is to befriend people, and with growing concerns about spam and malware on these sites, users may be reluctant to accept friend requests from strangers. Facebook is starting to offer marketers data on members as a part of its advertising system, but these are paid services.

## Joining Blogs

Blogs are public, so you don't "join" them the way you do a forum. By posting to a comment thread as part of a blog, you're joining a transient conversation on a particular topic.

Comment communities such as MyBlogLog, BackType, and Disqus consolidate an individual's comments and make them searchable, and if you use them you'll have a

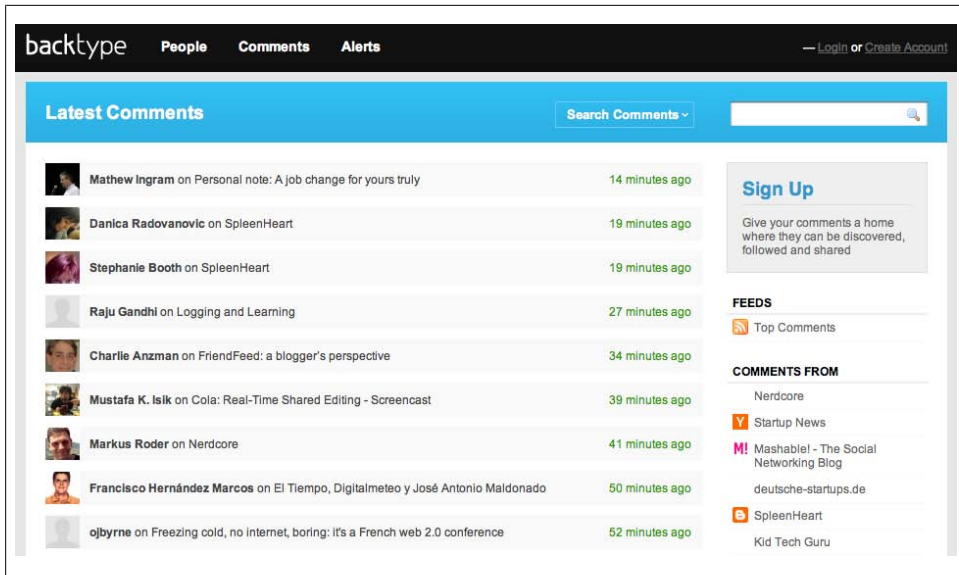


Figure 14-28. BackType monitors your comments and those of others

profile others can see. These services offer limited aggregate analysis, but they can be useful for looking into individual commenters and seeing where else they're posting. In other words, if you run a comment community on your blog, when someone comments on it, you'll be able to see what he has said elsewhere.

There may be situations in which you simply want to chime in on a comment thread, or even become a contributor in a blog network. The comment threads for an active post are the best places to join the conversation. When you comment on a blog post, you may get the option to be notified via email if someone responds to something you've posted. BackType, shown in [Figure 14-28](#), will let you track your comments and those of other community members.

## Joining Wikis

When you contribute content to a wiki, you're effectively joining it. You'll have an account—or at least an IP address—that can be tracked back to you. You're also implicitly agreeing that your content is subject to the modifications of others.

You don't get additional visibility into wiki activity when you start creating content, since transparency is one of the cornerstones of a collaborative reference. However, you do have the opportunity to correct issues or add information, provided it's non-partisan. If you believe content is false, you may want to enlist the assistance of editors or flag it as problematic. If you react directly within the entry by posting a rebuttal or deleting someone's work, you risk an escalating battle.



If you become a contributor to a wiki, you should make sure your user page is complete and discloses any interests or affiliations you have.

## Joining Micromessaging Tools

Twitter's APIs allow you to get detailed information about activity without joining. The main reason for being a part of a micromessaging platform is to interact with others and to build a social community.

Once you've got an account, you can track your own statistics: how many posts you've sent, how many friends you have, and how quickly that group of friends is growing. You can turn the many micromessaging analytics tools on yourself and analyze data, such as direct messages, that aren't publicly available.

## Joining Social News Aggregators

Participating in a news aggregator means submitting stories, commenting, and upvoting. As with a wiki, you're subjecting yourself to the whims of others and need to be transparent about who you are and what you represent. You can use the comment threads associated with an entry to provide additional details or differing viewpoints.

The people who run social news aggregators are constantly adjusting the algorithms by which votes are tallied in order to fight paid upvoting and spammers.

Social news aggregators have many clues they can use to figure out whether a submission's popularity is legitimate, and simply submitting good content to the right group with an appealing title is your best policy. By all means, join in and vote on stories you think are worthwhile, and leave comments on submissions when appropriate. Just don't be disingenuous, or the community will find out.

## Moderating a Community

Search is hands-off community monitoring, and joining communities only grants you more visibility if you're entering a walled garden you couldn't otherwise access. If you *really* want to see what's going on within a community, you need to help it form. As moderator, you'll have access to more details about the community than you would have as an outsider or even a member.

Many community platforms let you set up subgroups within the community and manage membership. The only problem with running or moderating a community, rather than joining an existing one, is convincing others to join it.

People will join your community if it provides value or gives them information they can't get elsewhere. For example, Dell runs a mailing list to which they send special offers. Subscribers get discounts, and Dell can drive traffic to its site and dispose of overstock.

**Step 1 of 3. Customize thisisatestgroup1123**

Do you want the group to be listed in the Yahoo! Groups directory? [Learn More](#)

☐ Yes, list the group. Default

☒ No, make the group unlisted.

---

How can people join your group?

☐ Anyone can join immediately. Default

☒ People can join only with my approval.

---

Who can post messages to your group?

☒ Only group members. (Private discussion group) Default

☐ Anyone. (Open discussion group)

☐ Only group owner. (Newsletter)

< Back   Next >   Finish   Cancel

Figure 14-29. Configuring a Yahoo! group's privacy settings

Similarly, people will visit a forum, Facebook group, or IRC chatroom if they know they'll get access to subject matter experts or celebrities.

## Moderating Groups and Mailing Lists

Google Groups, MSN Groups, and Yahoo! Groups all offer hosted tools to create and moderate a mailing list (Figure 14-29). Unless you have specific privacy or security needs that require you to manage your own mailing systems, consider these services. Not only will they let you control access to the mailing lists, they'll also remove the burden of things like spam blocking and unsubscribe requests. Additionally, the group directories that each site maintains will automatically index them.

As a moderator, you'll have considerable control over which features and functions members can use, as well as how easily they can search through content (Figure 14-30).

These services are lacking in analytics, however. If you want to tie mailing list activity to website outcomes, such as purchases or enrollments, you may need to run the group yourself. At the very least, you'll need to use custom URLs built with a tool like Google Analytics' URL Builder for the links you share with the mailing list so that you can track them through to conversion on your website.

## Moderating Forums

Forums are made up of topic areas, and you can create and moderate a topic without running a forum yourself. If you're a vendor, you may want to start a topic thread within a larger forum. In this way, you'll already be close to your audience. As topic moderator, you'll have additional control over what people can say and whether others can search the forum.




Figure 14-30. Administrative controls in a Yahoo! group

## Moderating Real-Time Chat Systems

By creating an IRC channel, you launch a discussion around which others can rally. If IRC users know that they'll have access to expertise on that channel, they're more likely to seek it out.

Creating channels on IRC is easy; you simply need to join a channel and hope that no one is already there. If you are able to join an unused channel, you will usually be automatically promoted to channel operator and you can begin to set up the channel topics and bots before you start funneling traffic to it. IRC channels are examples of an *emergent community* within a social platform; hashtags on Twitter have a similar function.

As a moderator (“@” or channel operator), you'll have the ability to enforce any rules that you see fit. You can do this by either moderating the channel when conversations veer away from what you deem acceptable or by kicking and banning a user from the channel for a designated period of time. These actions are generally reserved for times when things truly get out of hand, or when spammers begin to join the channel.

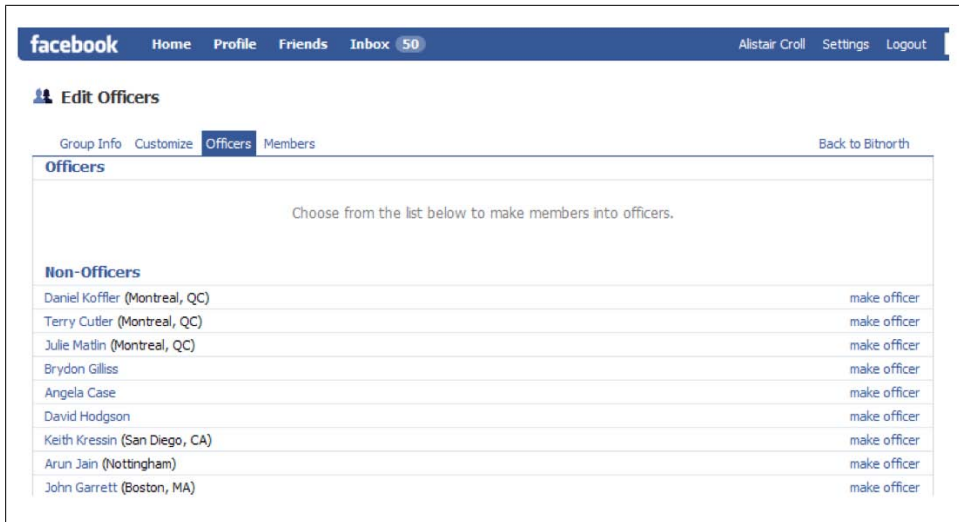


Figure 14-31. Managing members in a Facebook group

## Moderating Social Networks

You can create groups within social networks like Facebook (an example of such a group is shown in [Figure 14-31](#)). Groups can have events associated with them, and you can use the system to send messages to members, manage membership, and so on.

However, in keeping with the walled garden model of social networks, you have limited access to analytical details, such as whether people are reading your messages or accessing the group.

## Moderating Blogs

If you want to do more than comment on a blog, but don't want to run one outright, you may be able to contribute to an existing blog on occasion. For example, members of both U.S. political parties were given access to post to the site [BlogHer.com](#) during the elections, in order to foster an ad hoc community. Michelle Obama posted, as did a Republican chairwoman.

The blog operator may not give you access to information on your posts, such as how many people read the message or subscribed to an RSS feed, but you'll create a transient community in the comment thread that follows the post itself.

You can also encourage your community to let you know when they comment on you. If you have a blog, you'll get linkbacks whenever people reference you in their own blog postings. If someone mentions you on a community platform, acknowledge it and link to it when possible—this will encourage that person to engage more with you.



Figure 14-32. Popularity of specific terms on Hashtags.org

## Moderating Wikis

While Wikipedia is an open system, there are other topically focused wikis that may welcome your assistance in moderating and editing content. This is particularly true if you have subject matter expertise. As a moderator you'll have additional privileges, such as the ability to block contributors and undo changes; whether you'll get additional visibility into the activity of the community (e.g., how many people have read an entry) depends on the platform itself.

## Moderating Micromessaging Tools

You don't have to run your own microblogging site to create a community. Something as simple as a hashtag can be a rallying point for a community. For example, when the organizers of the Mesh conference wanted to launch a holiday party in Toronto, they used the hashtag #hohoto. Hashtags.org, a site that tracks hashtag use, immediately started showing the popularity of the term as shown in Figure 14-32.

Just remember that micromessaging is so open, you don't have any control over who joins or leaves the group (though you can claim a particular hashtag on sites like Thread.io).

## Moderating Social News Aggregators

There's not a lot you can do to moderate news aggregators. While the systems for submitting, commenting, and voting on content are relatively open, most of these sites handle the backend moderation themselves.

MY REDDITS ▾ REDDIT.COM · POLITICS · PICS · WTF · FUNNY · PROGRAMMING · WORLDNEWS · SCIENCE · TECHNOLOGY · ATHEISM · BUSINESS · COMICS · ENTERTAIN! MORE »

**reddit** WATCHINGWEBSITES

alistaircroll (30) | preferences | logout

manage your reddit

**name** Watchingwebsites

**title** Analytics, performance, usability, and online comm

**description** Subreddit focused on monitoring web applications, including web analytics, end user experience, performance and availability, usability, voice-of-the-customer feedback, and communities

**language** English [en]

**type** ☒ public anyone can view and submit  
☐ restricted anyone can view, but only contributors can submit links  
☐ private only contributors can view and submit

**age** ☐ viewers must be over eighteen years old

**media** ☐ show thumbnail images of content

**domain** redditwatchingwebsites.com [what's this?](#)

**look and feel** [edit the stylesheet](#) (leaves this page)  
 upload header image

**Watchingwebsites**  
 unsubscribe (1 subscriber)  
 a community for 369 milliseconds  
 you are a moderator of this reddit. [\(leave\)](#)  
[moderators](#) [edit](#) [ban users](#) [spam](#)

**Submit a link**  
 to anything interesting: news article, blog entry, video, picture...

**Create your own reddit**  
 ...because you love freedom.  
 ...for your favorite team.

Figure 14-33. Creating and administering a subreddit

You can, however, launch a subgroup. In reddit, these are known as subreddits, as shown in Figure 14-33.

These subgroups act as communities of interest in which members submit and vote on content. You can add moderators, block users, and see the size of the community, but you don't have access to analytics such as how often links were clicked.

## Running a Community

There are times when you'll need to operate the community yourself. Much of what we've already seen regarding watching your own websites applies to communities, and as a result, you have access to everything—server logs, analytical details, and more. Because you'll control design and layout, you can often embed JavaScript for analytics. You can also more tightly control the look and feel of the community and enforce terms of use that suit your business. In some cases, you may be forced to run a community yourself for legal or privacy reasons.

Note that we're talking here about *communities that are still intended for the public*. You may have internally facing communities intended for collaboration and project work, which we'll look at in the next chapter.

In addition to the headaches of setting up and running a community platform, you'll also need to convince others to join you. If you can manage to get your customers and fans to use a community that you run, you'll have an important asset for your

organization. You'll also have unbeatable insight into the community's behavior. Best of all, you'll be able to tie community activity back to business outcomes, such as viral spread, contribution, purchases, and so on.

## Running Groups and Mailing Lists

There are a number of open source and commercial mailing list servers. *majordomo* and GNU *mailman* are both mature listserv implementations, and while you can configure Microsoft's Exchange Server to support and run mailing lists, most administrators prefer to manage them through Microsoft's Live Communication Server.

If you don't want to run the servers, but still need control over mailing lists and the resulting analytics, consider a hosted mailing list service company such as Constant Contact or a private-label service like CakeMail, which can provide greater customization and analytical detail.

## Running Forums

There are two major kinds of forum software: general-purpose and commercial platforms.

The first kind of forum grew out of the open source world. These are general-purpose forums that aren't targeted at a particular business problem or industry. They're relatively mature and have ecosystems of developers who have released various plug-ins and extensions to add functionality such as analytics, topic voting, and spam filtering. Some of the most popular include phpBB and vBulletin, although there are many alternatives (see [http://en.wikipedia.org/wiki/Comparison\\_of\\_Internet\\_forum\\_software](http://en.wikipedia.org/wiki/Comparison_of_Internet_forum_software)).

The second kind of forum platform consists of commercial software that you can use for internal or external communities. Vendors such as Jive Software, Passenger (shown in Figure 14-34), Telligent (Figure 14-35), LiveWorld, and Mzinga fall in this category.

Because you're running the forum, you get far more visibility into what users are up to, including keyword popularity, sources of visitors, and more.

## Running Real-Time Chat Systems

If you want to run a private chat system, you have several options. You can use the IRC stack, which is a mature and reliable platform for real-time messaging but requires considerable know-how and demands that your community use client software to connect. These systems scale to hundreds of thousands of concurrent users, but they require dedicated clients such as *irssi*, *epic*, and *mIRC*, and aren't particularly user-friendly.





Figure 14-34. Running a forum yourself gives you far more insight into what's happening and who's participating




Figure 14-35. Contributions per day in Telligent's Harvest Reporting Server

```

z blockheap auth_heap elements used: 0 elements free: 256 memory in use: 0 total memory: 11264
z blockheap monitor_heap elements used: 13 elements free: 1011 memory in use: 416 total memory: 32768
z blockheap member_heap elements used: 226217 elements free: 35927 memory in use: 12668152 total memory: 14680064
z blockheap topic_heap elements used: 11716 elements free: 572 memory in use: 1171600 total memory: 1228800
z blockheap ban_heap elements used: 58329 elements free: 3111 memory in use: 1633212 total memory: 1720320
z blockheap channel_heap elements used: 25419 elements free: 7349 memory in use: 3863688 total memory: 4980736
z blockheap away_heap elements used: 5895 elements free: 761 memory in use: 565928 total memory: 638976
z blockheap user_heap elements used: 53043 elements free: 12493 memory in use: 1697376 total memory: 2097152
z blockheap lclient_heap elements used: 536 elements free: 488 memory in use: 237984 total memory: 454656
z blockheap client_heap elements used: 53100 elements free: 12436 memory in use: 14868000 total memory: 18350000
z blockheap nd_heap elements used: 0 elements free: 512 memory in use: 0 total memory: 24576
z blockheap confitem_heap elements used: 48298 elements free: 9046 memory in use: 2511496 total memory: 2981888
z blockheap librb_linebuf_heap elements used: 12 elements free: 2036 memory in use: 6336 total memory: 1081344
z blockheap librb_dnode_heap elements used: 202112 elements free: 84608 memory in use: 3233792 total memory: 4507520
z blockheap librb_fd_heap elements used: 568 elements free: 456 memory in use: 40896 total memory: 73728
z blockheap Total Allocated: 52943872 Total Used: 42500484
z Users 53042(1485176) Invites 4(48)
z User channels 226217(2714604) Aways 5895(209895)
z Attached confs 536(6432)
z Conflines 0(0)
z Classes 19(1064)
z Channels 25419(3991989)
z Bans 39259(1413324)
z Exceptions 4773(171828)
z Inxev 14297(514692)
z Channel members 226217(2714604) invite 4(48)
z Whowas users 30000(840000)
z Whowas array 30000(6840000)
z Hash: client 131072(1572864) chan 65536(786432)
z linebuf 16(8448)
z scache 61(1779)
z hostname hash 131072(1572864)
z Total: whowas 7680000 channel 8119965 conf 0
z Local client Memory in use: 536(388064)
z Remote client Memory in use: 52564(14717920)
z TOTAL: 30908792 Available: Current max RSS: 0
End of /STATS report

```

Figure 14-36. A small sample of available IRC server-side stats

Should you go this route, you'll have access to a wealth of information on community activity, as shown in [Figure 14-36](#). You can also run clients or automated scripts to monitor traffic and alert you to content or suspicious behavior.

For most self-run real-time chat communities, however, IRC platforms are overkill. Several community platforms include chat capabilities, and you can get chat plug-ins for blogs. Third-party services like Firefly, shown in [Figure 14-37](#), let you layer chat atop your website, but these haven't seen broad adoption.

Any real-time chat you run may become the target of drive-by spam and will require dedicated moderation to identify and block abusive users.

## Running Social Networks

Given the lack of analytics within public social networks, you may want to run your own. You can use hosted services like Ning, Flux, KickApps, or Me.com to create private social networks for your community, but you're still dependent on these platforms to provide you with the means to implement your own page instrumentation or to supply you with reporting and analytics.

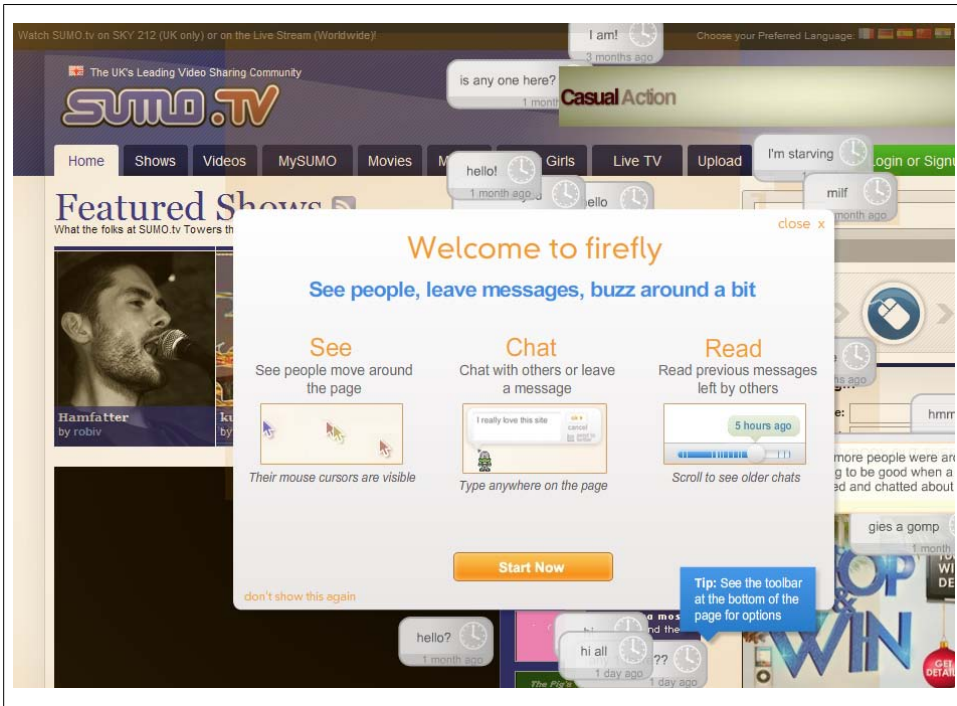


Figure 14-37. Firefly puts a chat page atop your website, where you can see what visitors are saying

## Running Blogs

With thousands of new blogs appearing on the Internet each day, starting one is trivial. If you use a hosted service for blogging (such as Six Apart, WordPress, or Blogger), you won't have to worry about operating it. Because you can control the stylesheets and content of the blog, you'll also be able to embed many monitoring tools—JavaScript-based analytics and real user monitoring, voice of the customer surveying, and web interaction analytics tools—directly into the pages to capture a great deal of information on your blog's activity. Blogs also support plug-ins such as StatZen, shown in [Figure 14-38](#), which will help you to identify popular content.

## Running Wikis

You can run your own wiki using one of several platforms, including the MediaWiki software on which Wikipedia runs. You can populate your wiki with your own content. This is one way to build a community around a specific topic; if you can corral subject matter experts and get them to start populating a wiki, you may be able to convince your fans to continue the work.

Title	Inferred	Web	Feed	Attention
With Customers Comes Bugs, Results in Progress	0	0	53	
Idea Valuation Revisited	7	7	35	
Cut to the Chase	13	0	30	
Goodbye Palm Treo, Hello Blackberry Curve	0	26	0	
Home Page	0	13	0	
HowTo: Extend the Apple iWork 08 Trial Period	0	11	0	
Back To School Pictures	0	9	0	
Goodbye Sprint, Hello AT&T	0	8	0	
Looking Forward Looking Backward	0	7	0	
More Info on Syncing a SprintPCS Samsung a900 with Mac OSX	0	6	0	

Figure 14-38. The StatZen reporting tool gives feed and post information

## Incipient links

One of the most important metrics on a wiki is the *incipient link*. When members create wiki entries, they provide links to pages they anticipate writing in the future. These incipient links identify places where additional content is needed. Incipient link analysis provides several key KPIs for a wiki:

- A page with no links leading away from it, incipient or otherwise, is an orphan that's not connecting to the rest of the wiki.
- A page with many incipient links needs love and attention because it's incomplete.
- The average time that an incipient link exists before its destination page is complete is a measure of turnover and the performance of a content generation team.
- The most frequently clicked-on incipient links are those that should probably be completed first, since they show topics of interest that your wiki hasn't yet completed to its audience's satisfaction.

## What operating a wiki can show you

When you run a wiki, you'll have access to much more information on which topics are being edited and where people are coming from. You will also be able to understand wiki-specific factors such as:

- *Top contributors*: you'll quickly see which members are generating the most content, as well as whose content is most disputed.




Figure 14-39. Wikipedia: new edits per day

- *Interconnectedness*: a wiki works best when content links to other content. Otherwise, it's simply a repository of articles. Analyzing how interconnected the wiki is overall is a measure of health and utility.
- *Edit patterns*: seeing repeated changes (such as a high number of edits by the same person in a small time period) can be an indicator of abuse or vandalism.



For a good analysis of metrics and analytics surrounding wikis, see [http://www.socialtext.net/open/index.cgi?wiki\\_analytics](http://www.socialtext.net/open/index.cgi?wiki_analytics) and [http://www.socialtext.net/open/index.cgi?wiki\\_metrics](http://www.socialtext.net/open/index.cgi?wiki_metrics).

Many of the projects and tools built for Wikipedia are available as open source code that you can repurpose for your own wiki. Figure 14-39 shows a Wikipedia tool that calculates the number of new edits per day on different language versions of the site.

If you don't want to run your own wiki, consider a hosted service you can embed in a website such as Wetpaint or Wikispaces. One advantage of running a wiki on your own site is that it shows search engines that you're a good source of authoritative information, because others will link to wiki content you create. If you create that content on Wikipedia, it will get a broader audience, but your domain won't get credit for it from search engines.

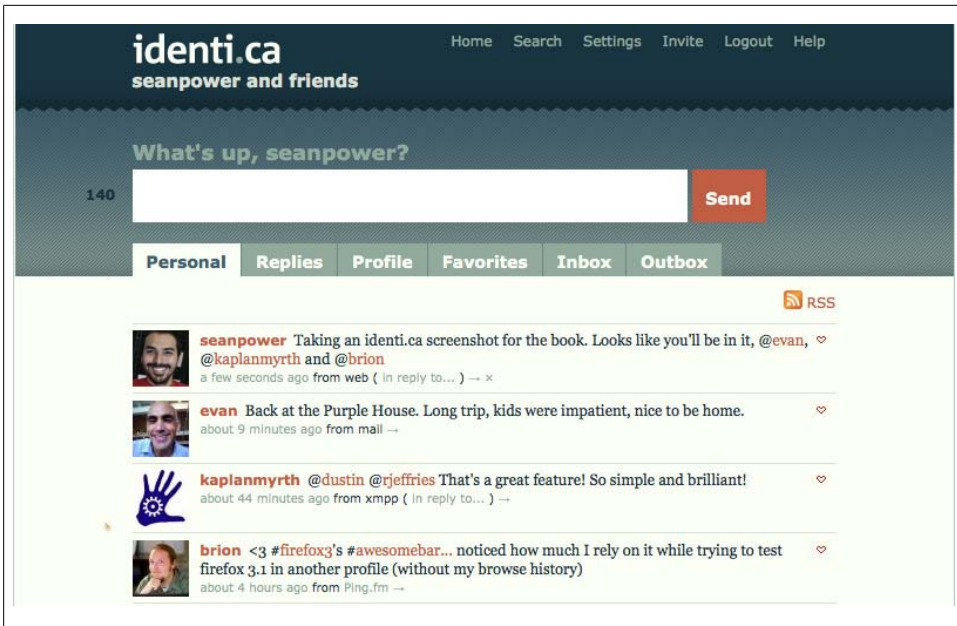


Figure 14-40. Identi.ca is an instance of the Laconi.ca open source microblogging platform

## Running Micromessaging Tools

If you want to operate your own version of Twitter, you can use Laconi.ca, which is an open source alternative that resembles Twitter. The Laconi.ca team runs an instance of Laconi.ca, called Identi.ca, shown in [Figure 14-40](#).

One of the advantages of Laconi.ca is its federated model, which allows you to join together several microblog instances so they can share messages between them. The micromessaging community is still defining what federation will do to the syntax of these platforms (for example, mentioning @seanpower might be limited to a local user with that name, but @@acroll might reference a user with that name on any federated node).

## Running Social News Aggregators

If you want to run your own news aggregator, and need more visibility than what you'll get from a subreddit, you can run an open source service like Pligg (shown in [Figure 14-41](#)) and encourage your community to submit and vote on content.

Pligg has a variety of commercial plug-ins and templates that you can use to customize how it looks.

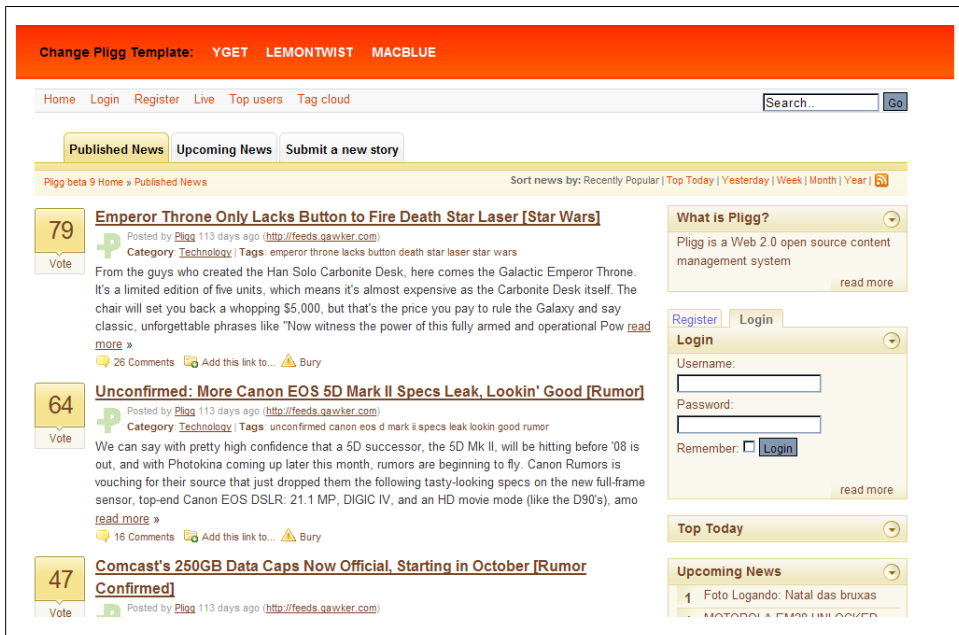


Figure 14-41. Pligg is a social news aggregator you can run on your own

## Putting It All Together

We've seen that conversations are made up of people (the LTOF, the fanatics, and the occasional contributors), the topics they're discussing, and the eight types of community platform.

Your job is to search, join, moderate, or run each of these communities to get the visibility into your online presence that your organization needs.

- *To follow the people who drive your community*, use search tools and interact with community members on social networks. Subscribe via RSS to blogs they write and encourage them to join mailing lists and forums you moderate. Consider interviewing them for blog entries or showcasing things they're doing.
- *To track topics that matter to you*, rely on prospective and retrospective search. Set up a dashboard for topics that matter, featuring sites and search results for those topics. Aggregate data from micromessaging services. Create and use hashtags for the topics you care about, and reference them in your communications and on your site. If mailing lists don't exist, create and promote them.
- *To be sure you're looking at the right sites*, keep track of inbound traffic and the sites that come up in alerts, and aggregate their content using more targeted search and tools (such as Yahoo! Pipes) that can consolidate multiple feeds. Find the sites



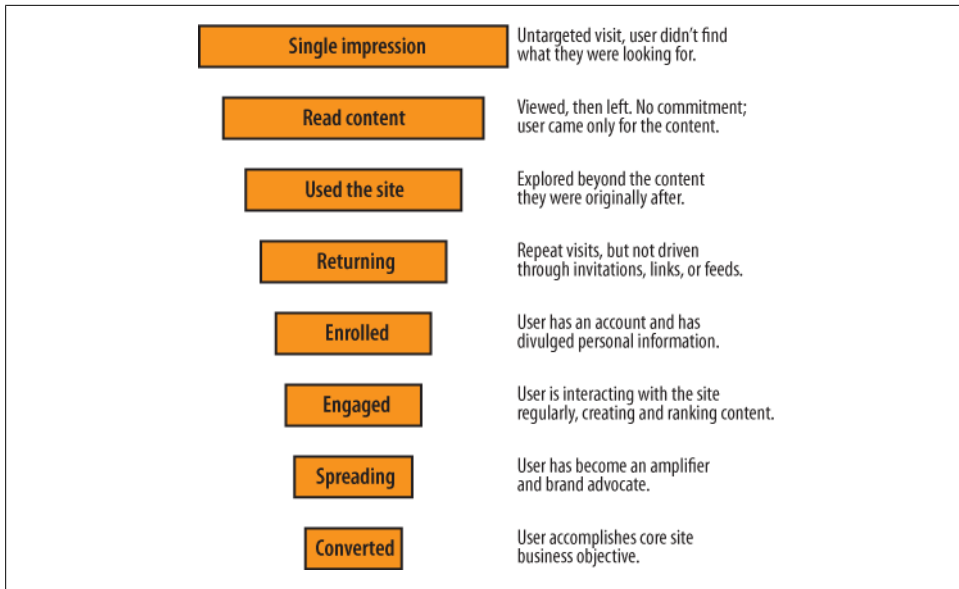


Figure 14-42. *The community funnel*

where you're the subject of conversation, and if those sites don't exist, launch them yourself.

## Measuring Communities and Outcomes

Your ability to track community engagement depends to a large degree on the nature of the community platform. As we've seen, when you run the platform you have far more insight into how your community engages with you. Because you control site design and the insertion of JavaScript, you can use many of the tools we've seen in earlier chapters, particularly web analytics, to understand the degree of engagement you're seeing within a community.

On the other hand, when you're merely a moderator or participant, measuring community engagement can be difficult.

As [Figure 14-42](#) shows, you can think of a community's members as moving through eight levels of engagement, from their first impressions of the community all the way to the moment when they become active participants.

### Single Impression

Any visitor who encounters a community gets an impression. If she doesn't see something of use, she will leave. Such visits typically last less than five seconds, and show up as bounces on your website ([Figure 14-43](#)). Of course, if you're not running the

+ Visitors who resolve to <a href="#">Telus Communications</a>	
👤 Visitors	1
🔗 Actions	1
🔗 Average actions per visit	1.0
🕒 Total time spent	5s
🕒 Average time per visit	5s
🚫 Bounce rate	100.0%

Figure 14-43. Comparing a visitor who failed to be impressed (and had only a single impression lasting five seconds) with the rest of your visitors

platform, you won't have any idea what the bounce rate of your messages is—you simply won't see activity, such as upvoting or comments, around the content.

## Read Content

You've reached a visitor when he discovers your community and reads or interacts with content he sees there. Having found the content he wanted, there's little guarantee he will return. His visit includes only one page, but the average time on the page is greater than that of a bounced visit.

## Used the Site

Visitors who navigated within the community website found a reason to stick around. They either had to navigate for their content or were intrigued by what they saw and wanted to see more. But they weren't interested enough to flag the site for future use (i.e., by bookmarking it) or to opt in for notifications about future content. These show up in traditional analytics tools as first-time visitors who saw more than one page (Figure 14-44).

## Returning

Users who come back to the site over time are more committed than those who use it once. Browser cookies can measure this aspect of the community's visitors; if your blog uses a cross-blog commenting system, you may be able to spot returning visitors, too. Reports on new versus returning visitors show how likely community members are to return.

Visitors who resolve to <a href="#">Qsc Ag Dynamic Ip Addresses</a>	
Visitors	1
Actions	8
Average actions per visit	8.0
Total time spent	2m 18s
Average time per visit	2m 18s
Bounce rate	0.0%

Figure 14-44. Visitors who performed multiple actions and used the site

you'll need to login or register to do that

(don't worry, it only takes a few seconds)

**CREATE A NEW ACCOUNT**  
all it takes is a username and password

username:

email: (optional)

password:

verify password:

type the letters from the image above

☐ remember me

I understand and agree that registration on or use of this site constitutes agreement to its [User Agreement](#) and [Privacy Policy](#).

is it really that easy? only one way to find out...

**LOGIN**  
already have an account and just want to login?

username:

password:

☐ remember me

[forgot your password?](#)

close this window

Figure 14-45. On reddit, users must first enroll before they can vote or contribute content

## Enrolled

Enrollment happens when a visitor provides credentials or tries to establish a relationship with the community. This may mean that she creates an account (as shown in [Figure 14-45](#)), subscribes to an RSS feed, or joins a mailing list.

For social sites, enrollment occurs when a member connects with you. This may involve someone following you on Twitter or joining your group on Facebook. In these cases, you'll see who these users are, and they'll be added to your tally of friends or your social graph.

On the other hand, for sites such as social news aggregators, you won't know users are enrolled until they contribute or comment on content within a subreddit.

## Engaged

Enrollment means establishing a connection; engagement means using it. A friend who sends you a message, a reader who comments on a blog, or a mailing list recipient who posts to the mailing list, is engaged. This may also be the act of upvoting or rating a submission (though this is harder to measure).

Users don't always have to enroll to engage you, however. They may interact with you in other ways—letting you know about inappropriate content, responding to a question via email, or responding to a survey, for example.

## Spreading

Community members who help spread your message are further down the funnel. This may consist of discussing your content on other platforms, submitting it to a social news aggregator, reTweeting it (as shown in [Figure 14-46](#)), or mentioning it in a message or blog post. If your platform offers tools to help members invite others, members may also be willing to tell their social network about you directly by sending invites. This is a higher level of engagement that indicates brand trust, and is a sure sign of fledgling brand advisors.

## Converted

As with all web analytics, community members are most engaged when they reach a business outcome. This is the return on community—it's where you get a financial benefit from your investment. It may be a customer getting help from a fellow community member rather than calling your support department. It may be a journalist contacting you because she heard about you. It may even be a product recommendation or a purchase. Whatever your conversion goals, this is the end of the funnel.

## Reporting the Data

There's no standard way to track community interactions, and there's a wide variety of data to collect across the eight community types we've seen, so you're going to have to stitch together, clean up, and present community metrics yourself. There are paid

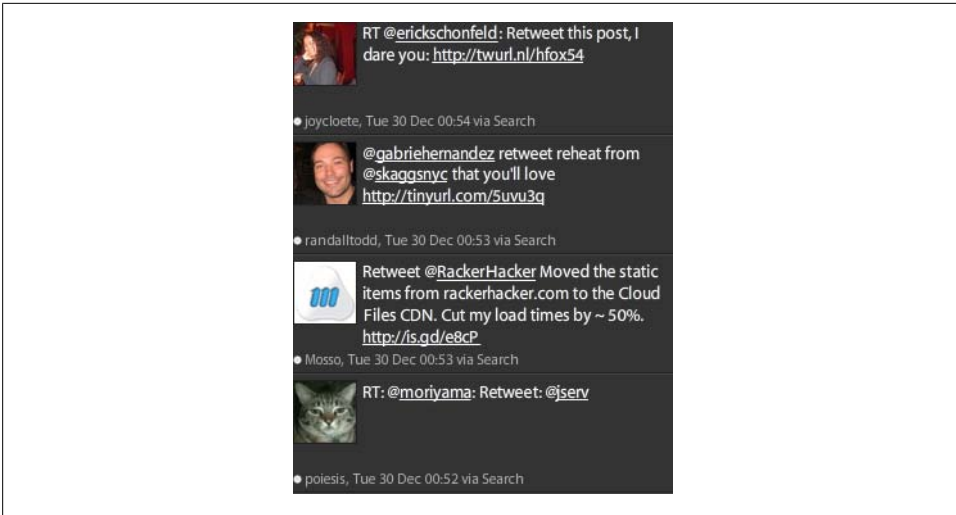


Figure 14-46. mollybenton reTweets an Akoha blog post

tools that can help you collect and analyze community activity, which we'll look at later in this chapter. First, however, let's consider what you need to report.

## What's in a Community Report?

The first rule for reporting is to *remember why you have a community initiative in the first place*. If the goal is to increase brand awareness, you should be reporting on visibility and mentions. On the other hand, if it's to reduce support call volumes, you need to track inbound calls and use of the online support forum.

Remember all the promises you made to your executive sponsor when you launched the community initiative? Tracking your progress toward the goals you agreed upon back then is critical. This is why you needed to know what the purpose of the community was: it's the only way to know what to report. It's easy to get bogged down in statistics and metrics about communities, but you need to distill all of that data down to the few nuggets of insight that help you inch toward your goal. Your executive sponsor won't care how many friends you have—she'll care how many of them told one another to buy something.

Table 14-1 lists some of the data you may want to include in a community report for your organization.

Table 14-1. Suggested list of metrics that you should include in a regularly published community report

Topic	Detail	Statistics
Business outcomes	Conversions	Percent of visitors
		Dollar amount
	Resolutions	Number of queries resolved online
General information	Contributions	User satisfaction with results
		Numbers of posts, posters, and new first-time contributors
	Notes and excerpts	Specific comments worthy of attention
Internal community	Mailing list	New enrollments, bounce rate, reach
	Support forum	Top 10 topics
		Top 10 most frequently asked questions
		Answers per question
		Votes per question
		Question ranking feedback
		Tag cloud of key topics
		Analytics reports
	Chatroom	Volume of chat activity (messages, users)
External community	Facebook group	Analytics reports
		Number of members, number of messages, segmentation if needed
	Micromessaging	Topic activity, hashtag activity, tag cloud of inbound messages, number of followers
		Identification of key reTweets or first mentions
	Blogosphere	Number of post mentions, number of comment mentions, top 10 posts mentioning you ranked by Technorati ranking

Social news aggregators		reddit, Digg, StumbleUpon, and Slashdot scores for any posting

You should always show reports against a historical graph (data is far more meaningful when it’s delivered as a comparison) and you should annotate these reports with events such as mailouts, posts, outages, software releases, or marketing campaigns that might have triggered changes in community behavior.

## The Mechanics of Tracking the Long Funnel

Tying business outcomes to the long funnel of a community is challenging. You can still understand the overall impact of communities on conversion, even if you’re just

overlaying the number of times your product is mentioned alongside website conversions at that time. This won't show you individual conversion rates, but it will show you what impact various conversations had on your business.

But what if you want to segment and analyze traffic more closely?

### **Custom URLs**

An old marketing technique that works well is to create unique links for each social network, which will allow you to track their spread as you distribute them throughout the world.

If someone submits a URL from your website on reddit or broadcasts it through Twitter, that person will grab whatever string is in the address bar. Your site can dynamically modify this (for example, with a timestamp) and use this information to track when the link was first copied. This kind of integration isn't for the faint of heart, since it requires a coordinated effort from engineering and marketing, but if you look carefully at the URLs of some of today's biggest sites, you'll see it happening.

### **The three-step personal invite**

Most of the analytics tools available for microblogging sites don't deal in business outcomes. They're more about popularity and frequency of communications, and less about which people or groups are driving customer behaviors you want. One reason for this is the inherent brevity of the messaging model. As an email marketer, you can send long custom URLs to each recipient and then track sessions by those URLs to assess whether a campaign was successful. However, it's hard to do that in 140 characters.

Traditional direct marketing relied on giving each prospect a unique identifier, allowing businesses to track prospects individually. This only works in a one-to-one model, however; on Twitter, everyone can read your status page and see the same URL.

One way around the constraints of one-to-many messaging is a three-step campaign model outlined in [Figure 14-47](#):

1. Send a promotion to your audience, asking them to reply with a certain string if they're interested.
2. When users reply with that string, reply to them with a direct message that contains a unique URL, preferably one that embeds campaign and follower data.
3. Track referrals from that campaign and the rest of the conversion funnel using traditional analytics.

Today, this kind of tracking is relatively unsophisticated and isn't well integrated into analytics or other tools, but as micromessaging models become more commonplace we expect to see more marketing of this sort. This is a likely source of longer-term revenue for Twitter, too, since the company can dynamically rewrite URLs in transit.

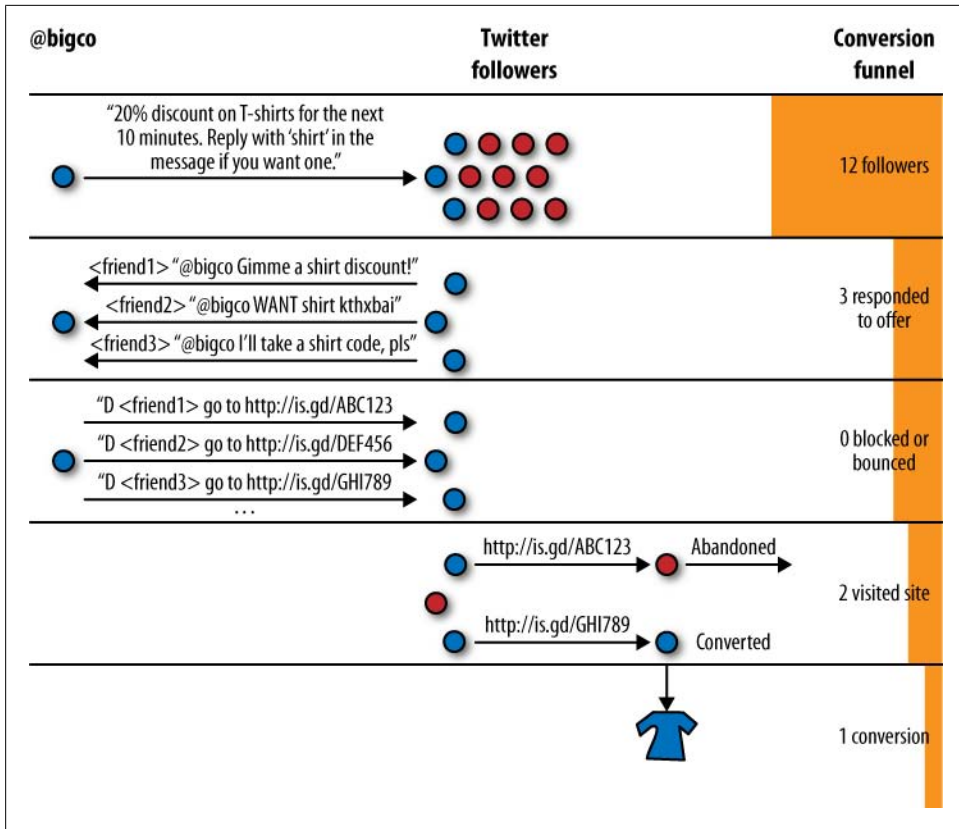


Figure 14-47. A three-step model for giving each follower a tracking URL in a promotion

### Using campaign URLs and URL shortening

You can try out several messages to a community to see which works best by using URL parameters for marketing campaigns. This process is illustrated in Figure 14-48.

1. Create the marketing messages you want to test out.
2. For each message, create a custom URL using a campaign builder (such as Google's URL builder).
3. Using a URL shortening service, generate a shortened URL. This will contain not only the URL of your content, but also the campaign parameters encoded in the URL.
4. Publish the message to the community, along with its associated shortened URL.
5. When members of the community see a message that appeals to them, they'll click the shortened URL.
6. The URL service will redirect members to the full campaign URL.



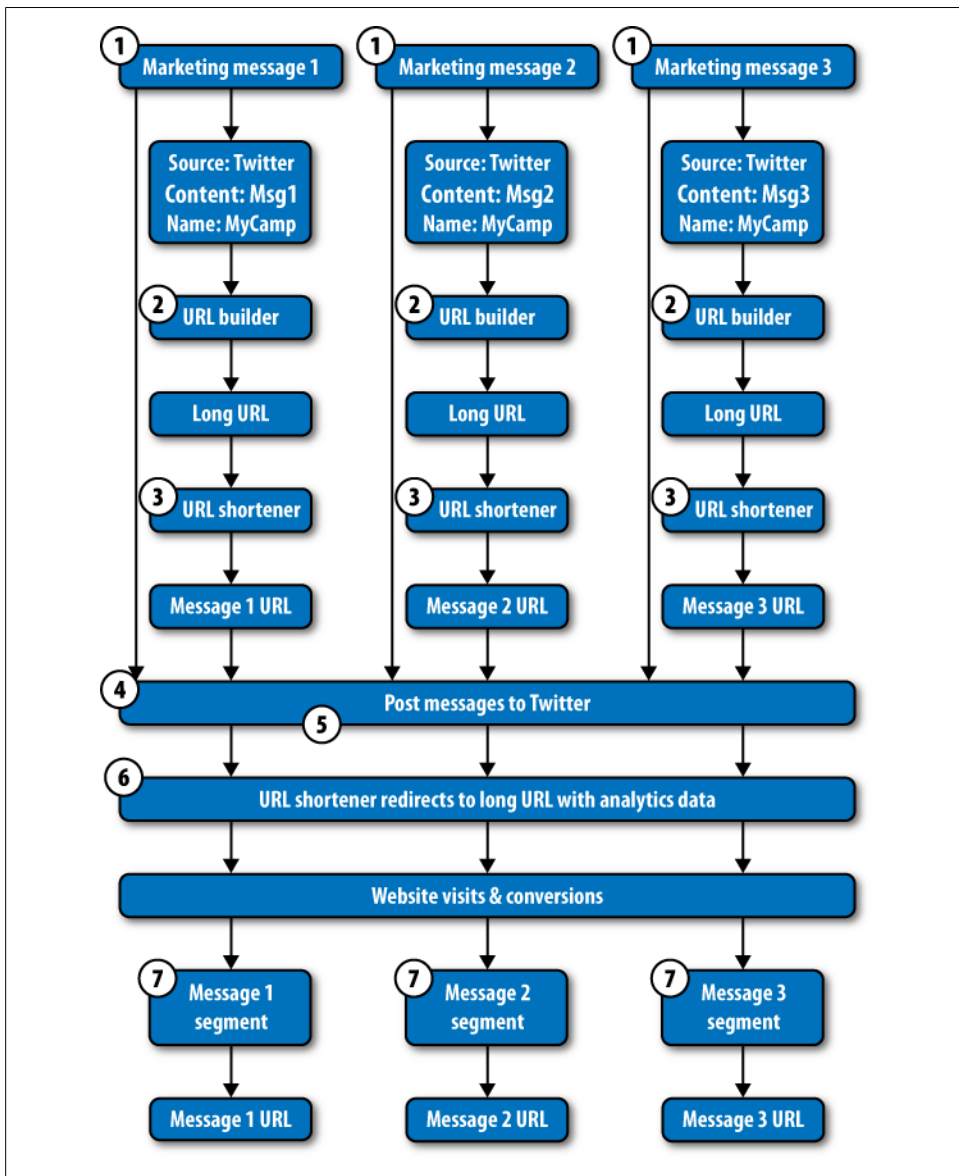


Figure 14-48. Test marketing multiple messages in microblogging with a URL shortener

7. When members are redirected to your content, the analytics tool will extract the campaign parameters and include them in your analytics reports as segments.

The best part is that this will happen whether members click on the link on the Web, within an email, or on a desktop client, so you can track viral spread across networks and tools.

This might seem a bit complex, so let's look at an example. Consider that we want to promote an upcoming event to our Twitter network. We first choose three message approaches:

- “Don’t miss out on this exciting event”
- “We need experts for the event”
- “Please come to the event; I miss you!”

For each of these, we create a URL using the Google URL builder (<http://www.google.com/support/googleanalytics/bin/answer.py?hl=en&answer=55578>), shown in [Figure 14-49](#). Note that we flag each of these messages with a unique “content” field, but we could use other fields for other segments, such as the social network we’re posting the message on or the medium (email, text message, blog comment) we’re using.

[Google Help](#) > [Analytics Help](#) > [Tracking central](#) > [Tracking basics](#) > Tool: URL Builder

### Tool: URL Builder

#### Google Analytics URL Builder

Fill in the form information and click the **Generate URL** button below. If you're new to tagging links or this is your first time using this tool, read [How do I tag my links?](#)

If your Google Analytics account has been linked to an active AdWords account, there's no need to tag your AdWords links - [auto-tagging](#) will do it for you automatically.

**Step 1:** Enter the URL of your website.

Website URL: \*   
(e.g. [http://www.urchin.com/download.html](#))

**Step 2:** Fill in the fields below. **Campaign Source**, **Campaign Medium** and **Campaign Name** should always be used.

Campaign Source: *	<input type="text" value="Twitter"/>	(referrer: google, citysearch, newsletter4)
Campaign Medium: *	<input type="text" value="Personal message"/>	(marketing medium: cpc, banner, email)
Campaign Term:	<input type="text"/>	(identify the paid keywords)
Campaign Content:	<input type="text" value="Personal appeal"/>	(use to differentiate ads)
Campaign Name*:	<input type="text" value="Cloudcamp09"/>	(product, promo code, or slogan)

**Step 3**

Figure 14-49. Using Google Analytics' URL Builder tool to generate a unique URL for the message

We now have a unique URL, which contains a web page and four parameters:

- <http://www.bitcurrent.com/interop-cloud-camp-2009/> is the destination page to which the link will lead.
- `?utm_source=Twitter` is the source of the traffic (where we intend to post the message).
- `&utm_medium=Personal%2Bmessage` is the format. If we'd used a banner or an email message, we'd mention it here.
- `&utm_content=Personal%2Bappeal` is the tone of the message (in this case, "personal message," but we will also generate URLs for "Don't miss out" and "Need experts").
- `&utm_campaign=Cloudcamp09` is the name of the campaign, which we use to distinguish it from other campaigns.

At 154 characters, that URL's a bit unwieldy—far too long to fit in a Twitter message, plus our community might notice that we've embedded a lot of information in there to track its members. We can easily shorten it into just a few characters, however, as shown in [Figure 14-50](#).

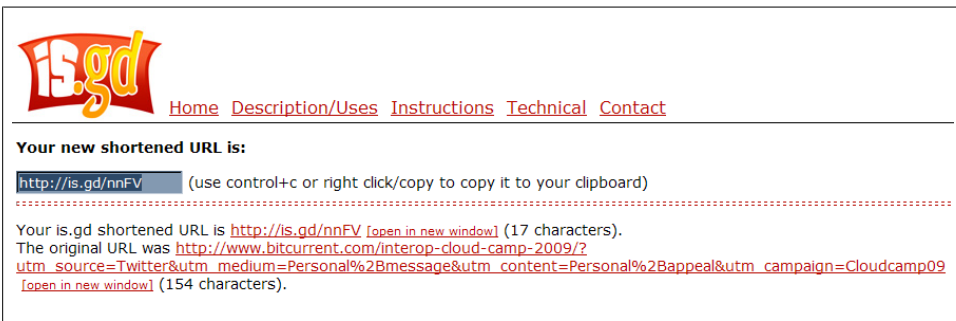


Figure 14-50. Creating the short URL for the campaign

Now we have a nice, short URL that we can embed in our message to Twitter ([Figure 14-51](#)). By combining the messages we've crafted with the shortened URLs that identify them, we can track which messages resonate best with our audience.

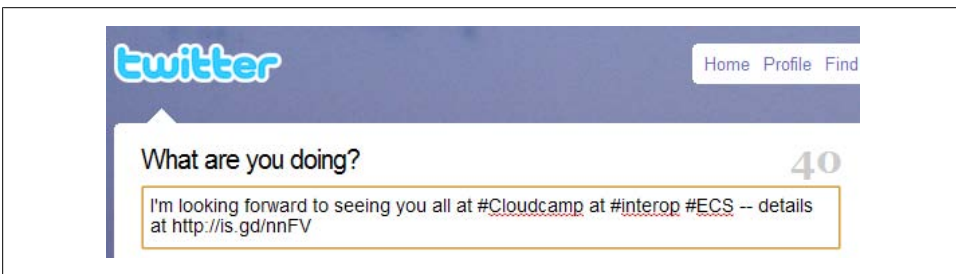


Figure 14-51. Posting the message in Twitter

After we run the campaign for a few days, we can see the results in Google Analytics, as shown in [Figure 14-52](#).

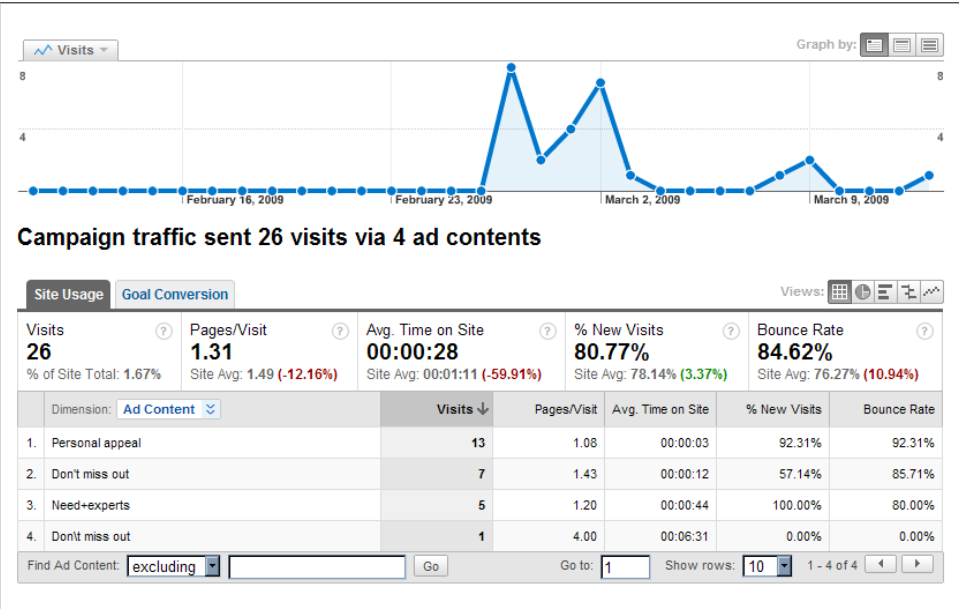


Figure 14-52. Examining the results in Google Analytics

Bear in mind that this technique is far from scientific: different messages go out at different times of day, when different numbers of visitors are online. More importantly, you'll get more attention for the first of the three messages you send, and a drop-off for subsequent messages, so try to pick a few tones of voice to stick with across several posts and transmit a different tone each time to remove this bias and learn what works best. Startup awe.sm and others are offering more automated, better integrated versions of this approach, too.

This is a good example of instrumenting the long funnel, tying together community analysis (discussions of a particular message) with analytical outcomes.

### Facebook and other multisite trackers

In November 2007, Facebook announced a system called Beacon that would track users across partners' websites. The marketing advantages were obvious: with better visibility across the Web, marketers could understand how and why visitors did what they did. Facebook members were less than enthusiastic, and the company quickly relented.

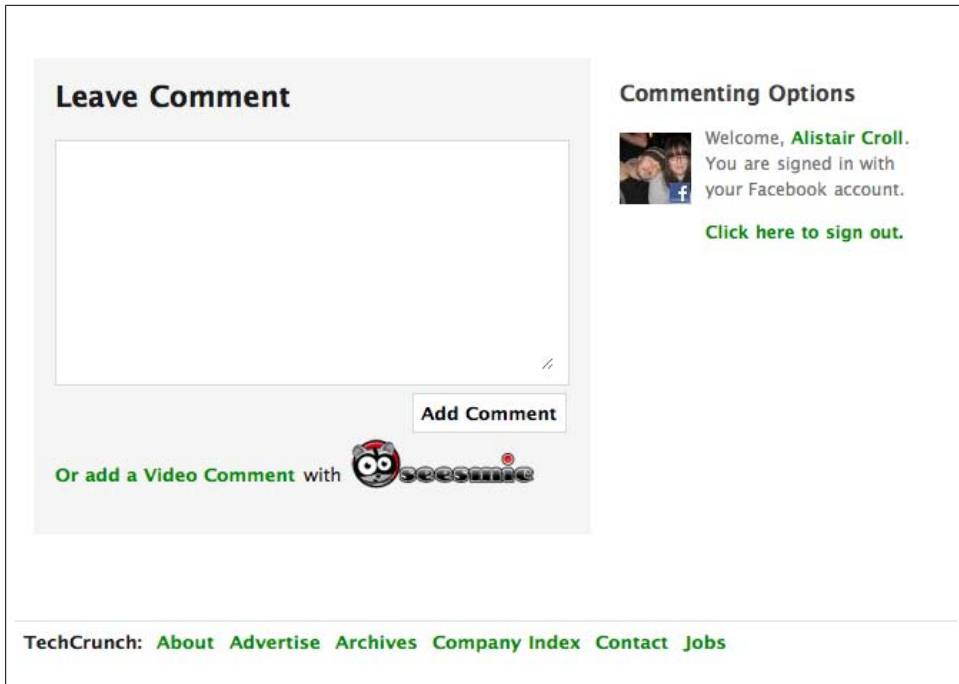


Figure 14-53. Facebook Connect integrated with TechCrunch; the system “follows” a user across the blogs on which they comment on an opt-in basis, then posts their comments to their Facebook lifestreams for their social graphs to read

Since then, Facebook has introduced a more moderate system, called Facebook Connect, that lets Facebook users flag content elsewhere online and tell their social graphs about it. As Figures 14-53 and 14-54 show, a Facebook Connect account spans several websites. When users log into blogs with their Facebook accounts, comments they leave are associated with their Facebook personas. That means comments are shared with their social graphs as part of their lifestreams.

Cross-community sites like Facebook Connect are likely to become more popular, since they provide an easy way for community members to track their contributions and share with their friends. Sites that participate in Facebook Connect get increased visibility through the sharing of those comments, and Facebook gets analytical data it can monetize.

## Responding to the Community

Every time you see a conversation online that’s relevant to your organization, you need to consider following up. You might join the conversation, amplify what’s being said, or directly engage with those who are talking. Your response may be as simple as visiting

an overview of#mesho9 and some of the great speakers we've got,  
including@mayormiller:http://bit.ly/YoNAV  
mathewi at 2009-03-12T18:37:58Z: I've got a post up with an overview of  
#mesho9 and some of the great speakers we've got, including @mayormiller:  
http://bit.ly/YoNAV

**Add New Comment** Login

 Logged in as **Alistair Croll**  
using Facebook Connect (Logout)

Type your comment here.

Post Comment

Figure 14-54. Facebook Connect integrated with [www.mathewingram.com/blog](http://www.mathewingram.com/blog)

a site and posting a comment or as complicated as mounting a dedicated marketing campaign to address a problem.

The U.S. Air Force developed an assessment tool, pictured in [Figure 14-55](#), that helps it decide whether it should respond to comments and blogs. It's an excellent example of turning community policies into clear, easy-to-follow directives that community managers can follow.

## Join the Conversation

The most obvious step to take in a community is to join the conversation. You can do this to steer opinions, to meet others so that you can engage them personally, or simply to improve your reputation within the community.

If you're adding information and opinions, make sure they're independently verifiable and avoid spamming. Make sure you look at previous remarks to be sure you aren't simply restating what someone has already said. Most importantly, start by listening: it's a good idea to lurk for a while before jumping in.

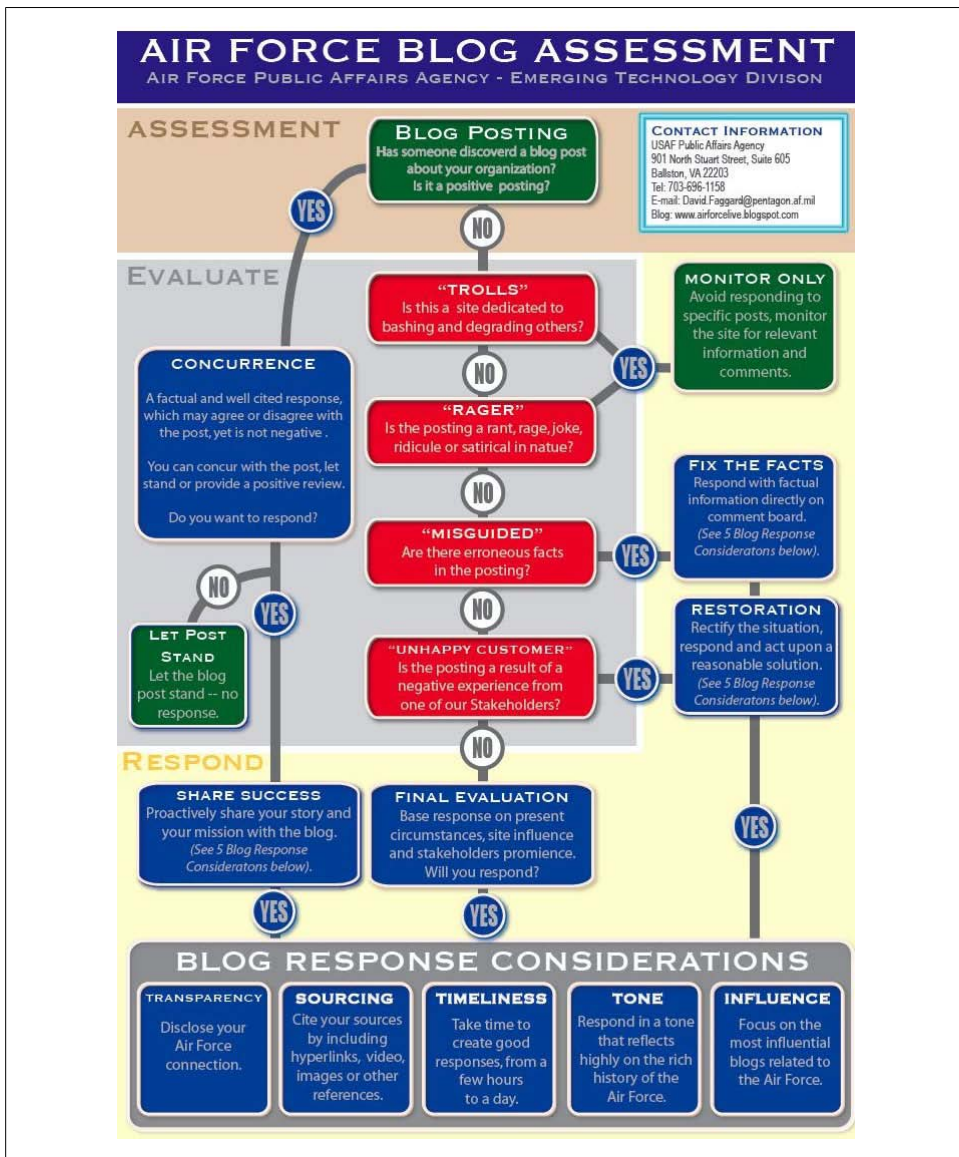


Figure 14-55. The Air Force blog assessment flowchart created by Capt. David Faggard, Chief of Emerging Technology at the Air Force Public Affairs Agency in the Pentagon

Remember that if you join the conversation, you may be unable to leave. Companies like Comcast and JetBlue, both of which have started to use micromessaging for support, have found that their communities expect them to always be around. These firms

have had to share the community management workload across many employees to provide the always-on presence their customers now demand.

## Amplify the Conversation

A second way to follow up on a mention is to try to encourage the spread of a message. If someone else says something about you, it will be far more genuine than if you say it yourself. But that shouldn't stop you from reinforcing it and pointing others to it, or from mentioning it on a blog or website.

Remember that there's a fine line between telling your friends about something and flat-out spamming. Many social news aggregators have sophisticated algorithms to try to weed out automated upvoting, but there are companies that specialize in pay-for-voting, resulting in an "arms race" between communities and gray-market, paid popularity. Jason Calcanis's offer to pay the top 50 users on any social networking site \$1,000 a month further stirred up this controversy (<http://calacanis.com/2006/07/18/everyones-gotta-eat-or-1-000-a-month-for-doing-what-youre/>).

The backlash you'll face for these kinds of practices often outweigh the value of the one-time burst of visits that occurs when your site makes it to the front page of reddit or Digg.

## Make the Conversation Personal

You may want to follow up directly with someone who mentions you online. Many social sites will let you see what else a person commented on or submitted, and you can follow that person's messages on a variety of social sites, from FriendFeed to Twitter to Plaxo.

Avoid being creepy. If the person with whom you're interested in connecting has an online presence that you can interact with, target this first, rather than using a more personal medium like email. For example, if a user has a blog, start by commenting on something she has written and build the relationship gradually. Gentle cultivation of community members is an important skill.

Remember, too, that just because everyone has a voice on the Web doesn't mean everyone has a veto. These early days of community management have many marketers walking on eggshells, terrified of even the smallest sound from a plaintive customer. Communities will soon reach a more sustainable equilibrium that distinguishes between legitimate customer grievances and the whining of partisan entitlement.

## Community Listening Platforms

Clearly, there's a tremendous amount of work required to find, engage, and report on communities. If you just consider the eight communities we've seen, and the four ways



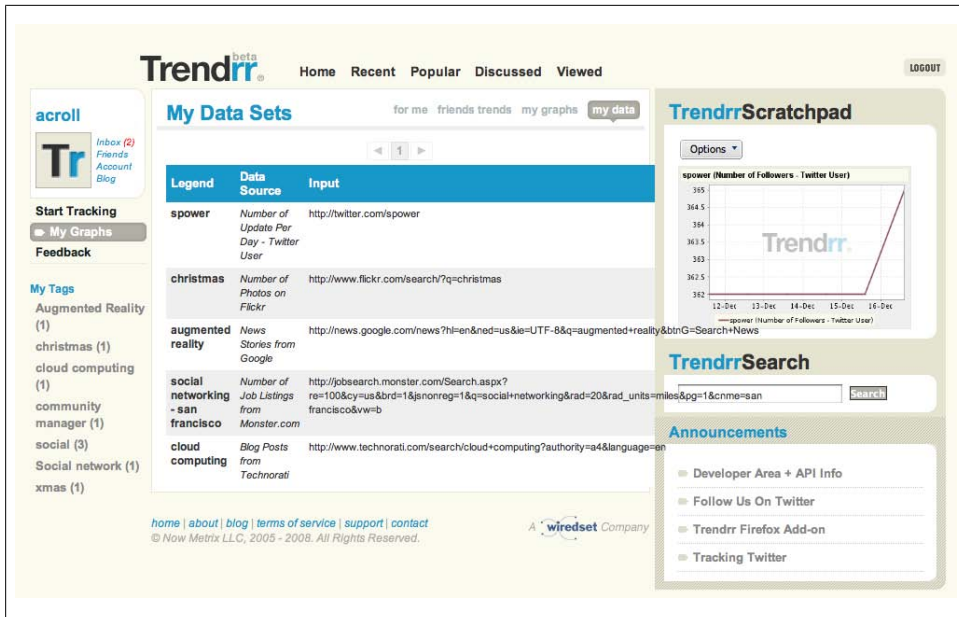


Figure 14-56. Trendrr is a general-purpose tool for tracking multiple sources of online data in a single place

you can get involved (search, enrollment, moderation, and running them), and the options to report, join, amplify, or contact members within a community, you have a daunting number of possibilities. Multiply this by the thousands of messages you might receive, and there's no way to keep up manually.

There are several tools and techniques for tracking your community presence. First of all, you can use consolidated reporting sites like Trendrr, shown in Figure 14-56, to collect and report on a variety of metrics.

Many commercial platforms, such as Radian6, Techrigy, ScoutLabs, BuzzMetrics, Sysonomos, and BuzzLogic, provide reports and automate many of the search techniques we've outlined here (we're convinced that the term "buzz" has been forever claimed by this segment of the monitoring industry). Still other services and tools focus on bulk messaging and automated responses to handle a growing number of direct interactions between a company and its communities.

Regardless of the vendor, community monitoring tools do three basic things:

- *Find conversations* across many community platforms and websites
- *Aggregate the content* into digestible chunks and visualizations, sometimes detecting sentiment, demographics, and attitudes
- *Manage your responses* to what's happening in the community

## How Listening Tools Find the Conversations

To collect all of the conversations happening on the Internet, social listening platforms use a variety of techniques. While they might rely on search engines like Google, many prefer their own collection systems. They run web crawlers that index websites and look for keywords. They subscribe to RSS feeds. And they harvest data from the APIs of micromessaging services like Twitter. Ultimately, they'll collect conversations and content from the many platforms we've looked at, but will usually stop short of joining those platforms on your behalf.

The use of crawlers for community monitoring is a controversial subject. To provide a real-time understanding of what communities are up to, these crawlers index sites more frequently than those of large search engines like Google, which can put additional load on servers. What really earns them the wrath of some website operators, however, is the fact that they crawl aggressively: they want to find any mention of a conversation, anywhere on the Web, even when told not to.

Most websites tell crawlers which pages on the site can be indexed through a file called *robots.txt*. This file specifies the pages that the crawler may traverse and those that it should leave alone. There are several reasons a site operator may not want crawlers to index the site. The site operator may want to force visitors to come to his sites to find content rather than finding it through an external search engine; he may not want the additional load; or the pages may be dynamic and subject to frequent changes.

Some social listening platforms' crawlers have been accused of ignoring *robots.txt*, or of using user agents that hide their true identities (<http://web-robot-abuse.blogspot.com/2008/05/wwwradian6comcrawler.html>). This is a violation of the site's terms of use, and some site operators go to lengths to block such crawlers through other means.

## How Tools Aggregate the Content

Once they've collected the thousands of conversations that might matter to you, community monitoring tools try to make sense of them so that the important interactions rise to the top. This takes several forms, including rivers of news, keyword summaries, influencer scores, threading and drill-downs to individuals, and time graphs.

### River of news to tap into the feed

A river of news is a lifestream for the topics you've flagged as important, as shown in [Figure 14-57](#). You can segment the river by platform, most popular topics, or most influential participants. This serves as a good real-time visualization of how people are discussing your brand or topics online, as well as where those conversations are happening. You can often drill down to individuals participating in the conversations.

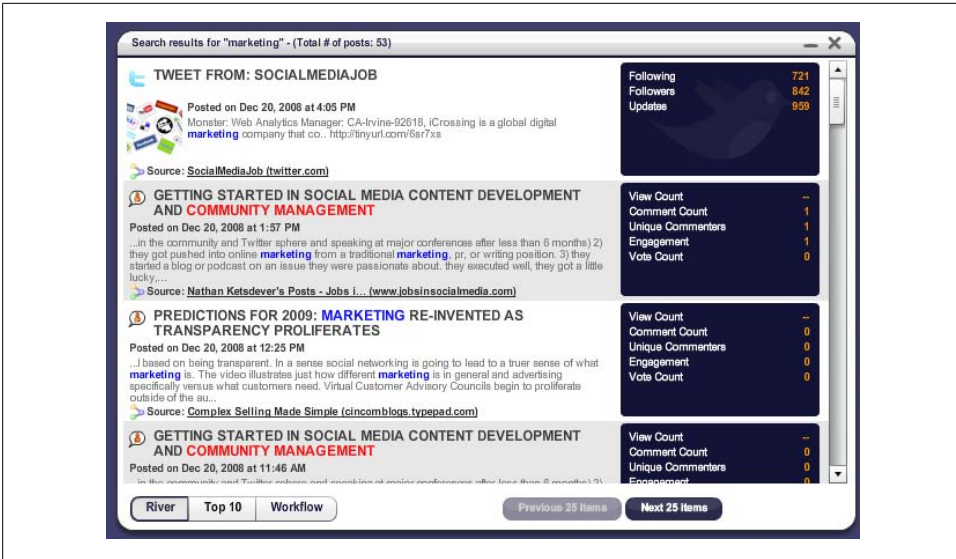


Figure 14-57. The river of news report in Radian6 shows recent activity on a topic

### Keyword summary to see what's being discussed

Rather than viewing community activity as a feed of news, you may want to aggregate conversations by keywords over time. You can compare specific terms within a certain topic for popularity, as shown in Figure 14-58.




Figure 14-58. An analysis of keywords across blog postings within Radian6




Figure 14-59. The Share of Voices report in Techrify’s SM2 product shows the number of times people have searched for a particular subject

Choosing the right words within a segment, such as product names, helps you estimate the relative popularity of particular topics or assess overall mood across communities. Be careful not to read too much into such results: without context, the results may be misleading. It’s much better to visit the places you’re being discussed and get a sense of what’s happening, rather than relying on how often a particular word is mentioned.

You can also view graphs that show you the distribution of keywords across social media channels (Figure 14-59).

### Influencer scores to see who’s got clout

Every mention on a community has someone behind it, and some community members are more influential. This is the “who’s talking” part of things. Perhaps they have more followers, a better Technorati ranking, more page views according to [Compete.com](#), more influence on [Klout.net](#), or a bigger social graph in OpenSocial. Whatever the case, you need to identify key influencers who are discussing you.

Figure 14-60 shows an influencer report that lets you compare the people contributing to topics you care about.

### Threading and drill-down to look at the people

Once you’ve identified an important influencer, regardless of his social network or the platform he’s using, you’ll want to learn more about him. Many of the community






Figure 14-62. Scoutlabs' report of keyword mentions on blogs

### Graphs over time to understand how much

It's not enough to know which topics are important and who's talking about them. You will also want to understand the trends of particular topics within a community or the blogosphere at large, as shown in [Figure 14-62](#).

Historical trend analysis of this kind can be overlaid with analytical data to see whether a spike in community traffic corresponds to a rise in site traffic or goal conversions. You can also use this kind of analysis to compare competitors and see who's pulling ahead.

### Sentiment analysis to understand tone and mood

While it's fine to see mentions of a topic, it's just as important to understand the context in which those mentions occurred: was the sentiment positive or negative?

Evaluating sentiment requires that a machine parse online mentions and assign it a score—positive, negative, or neutral. If a blog post reads, "[Complete Web Monitoring](#) is awesome!" it would be scored as a positive response. On the other hand, "I thought [Complete Web Monitoring](#) was rubbish!" would clearly be a negative response (and, of course, patently false).

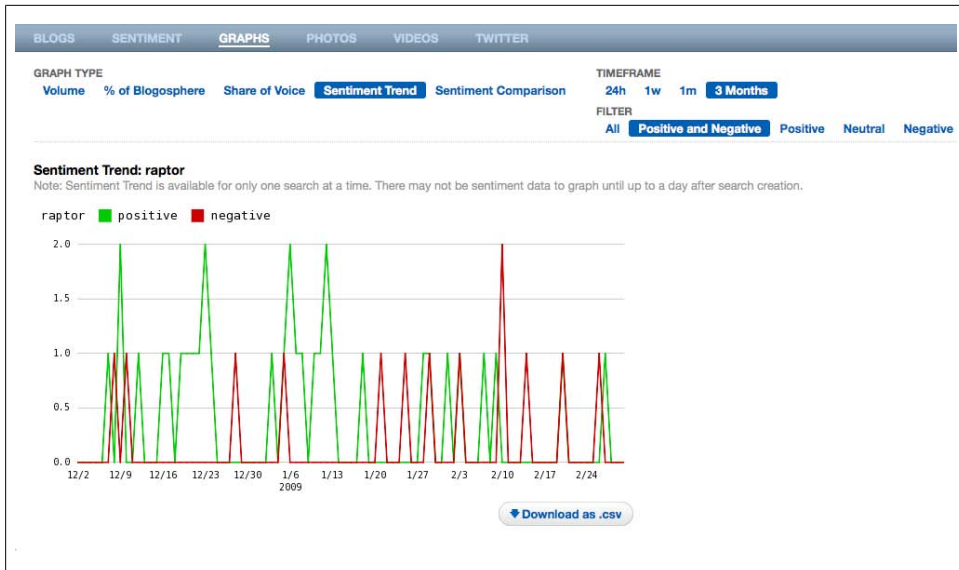


Figure 14-63. A sentiment graph in Scout Labs, based on a semantic analysis of language

But sentiment analysis isn't easy. Consider, for example, posts that contain irony, sarcasm, or colloquialisms. The sentence "Great, I'm so happy I just got Rickrolled again" isn't positive, but a computer would score it as such because of the tone of the keywords it contains.

Margaret Francis, VP of products at Scout Labs, a community management tool that includes sentiment analysis features says, "If [a sentiment analysis] algorithm ever gets good enough to [natively] classify that as negative, I will start stockpiling the explosives needed to take down Skynet"—in other words, such a tool would be frighteningly close to true artificial intelligence and might pass a Turing Test.

Understanding whether people like, dislike, or disregard certain features in your product helps you set product road map priority and anticipate support issues. [Figure 14-63](#) shows an example of a sentiment analysis report.

## How Tools Manage the Response

The third function of community management tools is handling responses. We've looked at how to do this on a personal level, and we believe that individual interactions from your market deserve individual responses—to do otherwise is disingenuous at best and disrespectful at most.

Some community management platforms help ensure you can follow up with community interactions by treating mentions as "tickets" that can be assigned to someone within your organization. CRM (customer relationship management) vendors, in particular, are integrating community interactions with customer support applications.

This allows you to track how a situation was resolved and determine which community interactions have yet to be addressed.

On the other hand, there are some tools that let you craft an automated response similar to a form letter and send it to followers based on what they say. Some companies may choose to adopt this functionality to reduce community management overhead or to deliver mass-marketing messages. You do so at your own peril, and you'll likely be called out for doing it.

## Community Monitoring Maturity Model

Here's how organizations develop their community monitoring as they become more comfortable with and sophisticated about communities. Note that this applies to external communities only—we'll look at a maturity model for internal communities in [Chapter 15](#).

	Level 1	Level 2	Level 3	Level 4	Level 5
<b>Focus</b>	Technical details	Minding your own house	Engaging the Internet	Building relationships	Web business strategy
<b>Emphasis</b>	Flags, alerts, and basic awareness	Brand and message update	Community mind-set, viral spread, and rankings	Nurturing communities, measuring ROI	Interactive marketing, multivariate community testing, moderation
<b>Questions</b>	Am I being discussed?	Do I have a presence? Are people engaging with me on my own site?	Are they listening & amplifying me? What's sentiment like? Which messages work best with which communities?	How does the "long funnel" affect my revenues? What's the lifetime engagement of a visitor across all social networks? Who are my most vocal/supportive community members?	What will the community buy? How do I automatically match the right messages to each audience? How does my community help itself? Is virality a part of business planning?