

Index

- 0-1 loss, 104, 276
- Absolute value rectification, 192
- Accuracy, 426
- Activation function, 170
- Active constraint, 95
- AdaGrad, 307
- ADALINE, see adaptive linear element
- Adam, 308, 428
- Adaptive linear element, 15, 24, 27
- Adversarial example, 267
- Adversarial training, 268, 270, 533
- Affine, 110
- AIS, see annealed importance sampling
- Almost everywhere, 71
- Almost sure convergence, 130
- Ancestral sampling, 583, 598
- ANN, see Artificial neural network
- Annealed importance sampling, 628, 670, 719
- Approximate Bayesian computation, 718
- Approximate inference, 586
- Artificial intelligence, 1
- Artificial neural network, see Neural network
- ASR, see automatic speech recognition
- Asymptotically unbiased, 124
- Audio, 102, 361, 461
- Autoencoder, 4, 357, 505
- Automatic speech recognition, 461
- Back-propagation, 203
- Back-propagation through time, 385
- Backprop, see back-propagation
- Bag of words, 474
- Bagging, 255
- Batch normalization, 266, 428
- Bayes error, 117
- Bayes' rule, 70
- Bayesian hyperparameter optimization, 439
- Bayesian network, see directed graphical model
- Bayesian probability, 55
- Bayesian statistics, 135
- Belief network, see directed graphical model
- Bernoulli distribution, 62
- BFGS, 316
- Bias, 124, 229
- Bias parameter, 110
- Biased importance sampling, 596
- Bigram, 465
- Binary relation, 485
- Block Gibbs sampling, 602
- Boltzmann distribution, 573
- Boltzmann machine, 573, 656
- BPTT, see back-propagation through time
- Broadcasting, 34
- Burn-in, 600
- CAE, see contractive autoencoder
- Calculus of variations, 179
- Categorical distribution, see multinoulli distribution
- CD, see contrastive divergence
- Centering trick (DBM), 675
- Central limit theorem, 63
- Chain rule (calculus), 206
- Chain rule of probability, 59

- Chess, 2
- Chord, 582
- Chordal graph, 582
- Class-based language models, 466
- Classical dynamical system, 376
- Classification, 100
- Clique potential, see factor (graphical model)
- CNN, see convolutional neural network
- Collaborative Filtering, 481
- Collider, see explaining away
- Color images, 361
- Complex cell, 366
- Computational graph, 204
- Computer vision, 455
- Concept drift, 541
- Condition number, 279
- Conditional computation, see dynamic structure
- Conditional independence, xiii, 60
- Conditional probability, 59
- Conditional RBM, 687
- Connectionism, 17, 446
- Connectionist temporal classification, 463
- Consistency, 130, 516
- Constrained optimization, 93, 237
- Content-based addressing, 422
- Content-based recommender systems, 483
- Context-specific independence, 576
- Contextual bandits, 483
- Continuation methods, 328
- Contractive autoencoder, 524
- Contrast, 457
- Contrastive divergence, 291, 613, 674
- Convex optimization, 141
- Convolution, 331, 685
- Convolutional network, 16
- Convolutional neural network, 252, 331, 428, 463
- Coordinate descent, 322, 673
- Correlation, 61
- Cost function, see objective function
- Covariance, xiii, 61
- Covariance matrix, 62
- Coverage, 427
- Critical temperature, 606
- Cross-correlation, 333
- Cross-entropy, 75, 132
- Cross-validation, 122
- CTC, see connectionist temporal classification
- Curriculum learning, 329
- Curse of dimensionality, 154
- Cyc, 2
- D-separation, 575
- DAE, see denoising autoencoder
- Data generating distribution, 111, 131
- Data generating process, 111
- Data parallelism, 450
- Dataset, 105
- Dataset augmentation, 270, 460
- DBM, see deep Boltzmann machine
- DCGAN, 554, 555, 703
- Decision tree, 145, 551
- Decoder, 4
- Deep belief network, 27, 532, 634, 659, 662, 686, 694
- Deep Blue, 2
- Deep Boltzmann machine, 24, 27, 532, 634, 655, 659, 665, 674, 686
- Deep feedforward network, 167, 428
- Deep learning, 2, 5
- Denoising autoencoder, 513, 691
- Denoising score matching, 622
- Density estimation, 103
- Derivative, xiii, 83
- Design matrix, 106
- Detector layer, 340
- Determinant, xii
- Diagonal matrix, 41
- Differential entropy, 74, 649
- Dirac delta function, 65
- Directed graphical model, 77, 510, 566, 694
- Directional derivative, 85
- Discriminative fine-tuning, see supervised fine-tuning
- Discriminative RBM, 688
- Distributed representation, 17, 150, 549
- Domain adaptation, 539

- Dot product, 34, 141
- Double backprop, 270
- Doubly block circulant matrix, 334
- Dream sleep, 612, 655
- DropConnect, 265
- Dropout, 257, 428, 433, 434, 674, 691
- Dynamic structure, 451, 452
- E-step, 637
- Early stopping, 246, 249, 272, 273, 428
- EBM, see energy-based model
- Echo state network, 24, 27, 406
- Effective capacity, 114
- Eigendecomposition, 42
- Eigenvalue, 42
- Eigenvector, 42
- ELBO, see evidence lower bound
- Element-wise product, see Hadamard product, see Hadamard product
- EM, see expectation maximization
- Embedding, 519
- Empirical distribution, 66
- Empirical risk, 276
- Empirical risk minimization, 276
- Encoder, 4
- Energy function, 572
- Energy-based model, 572, 598, 656, 665
- Ensemble methods, 255
- Epoch, 247
- Equality constraint, 94
- Equivariance, 339
- Error function, see objective function
- ESN, see echo state network
- Euclidean norm, 39
- Euler-Lagrange equation, 649
- Evidence lower bound, 636, 663
- Example, 99
- Expectation, 60
- Expectation maximization, 637
- Expected value, see expectation
- Explaining away, 577, 634, 647
- Exploitation, 484
- Exploration, 484
- Exponential distribution, 65
- F-score, 426
- Factor (graphical model), 570
- Factor analysis, 493
- Factor graph, 582
- Factors of variation, 4
- Feature, 99
- Feature selection, 236
- Feedforward neural network, 167
- Fine-tuning, 324
- Finite differences, 442
- Forget gate, 306
- Forward propagation, 203
- Fourier transform, 361, 363
- Fovea, 367
- FPCD, 617
- Free energy, 574, 682
- Freebase, 486
- Frequentist probability, 55
- Frequentist statistics, 135
- Frobenius norm, 46
- Fully-visible Bayes network, 707
- Functional derivatives, 648
- FVBN, see fully-visible Bayes network
- Gabor function, 369
- GANs, see generative adversarial networks
- Gated recurrent unit, 428
- Gaussian distribution, see normal distribution
- Gaussian kernel, 142
- Gaussian mixture, 67, 188
- GCN, see global contrast normalization
- GeneOntology, 486
- Generalization, 110
- Generalized Lagrange function, see generalized Lagrangian
- Generalized Lagrangian, 94
- Generative adversarial networks, 691, 702
- Generative moment matching networks, 705
- Generator network, 695
- Gibbs distribution, 571
- Gibbs sampling, 584, 602
- Global contrast normalization, 457
- GPU, see graphics processing unit
- Gradient, 84

- Gradient clipping, 289, 417
- Gradient descent, 83, 85
- Graph, xii
- Graphical model, see structured probabilistic model
- Graphics processing unit, 447
- Greedy algorithm, 324
- Greedy layer-wise unsupervised pretraining, 531
- Greedy supervised pretraining, 324
- Grid search, 435
- Hadamard product, xii, 34
- Hard tanh, 196
- Harmonium, see restricted Boltzmann machine
- Harmony theory, 574
- Helmholtz free energy, see evidence lower bound
- Hessian, 223
- Hessian matrix, xiii, 87
- Heteroscedastic, 187
- Hidden layer, 6, 167
- Hill climbing, 86
- Hyperparameter optimization, 435
- Hyperparameters, 120, 433
- Hypothesis space, 112, 118
- i.i.d. assumptions, 111, 122, 267
- Identity matrix, 36
- ILSVRC, see ImageNet Large-Scale Visual Recognition Challenge
- ImageNet Large-Scale Visual Recognition Challenge, 23
- Immortality, 580
- Importance sampling, 595, 627, 700
- Importance weighted autoencoder, 700
- Independence, xiii, 60
- Independent and identically distributed, see i.i.d. assumptions
- Independent component analysis, 494
- Independent subspace analysis, 496
- Inequality constraint, 94
- Inference, 565, 586, 634, 636, 638, 641, 651, 653
- Information retrieval, 528
- Initialization, 301
- Integral, xiii
- Invariance, 343
- Isotropic, 65
- Jacobian matrix, xiii, 72, 86
- Joint probability, 57
- k-means, 365, 549
- k-nearest neighbors, 143, 551
- Karush-Kuhn-Tucker conditions, 95, 237
- Karush–Kuhn–Tucker, 94
- Kernel (convolution), 332, 333
- Kernel machine, 551
- Kernel trick, 141
- KKT, see Karush–Kuhn–Tucker
- KKT conditions, see Karush-Kuhn-Tucker conditions
- KL divergence, see Kullback-Leibler divergence
- Knowledge base, 2, 486
- Krylov methods, 224
- Kullback-Leibler divergence, xiii, 74
- Label smoothing, 243
- Lagrange multipliers, 94, 649
- Lagrangian, see generalized Lagrangian
- LAPGAN, 704
- Laplace distribution, 65, 499, 500
- Latent variable, 67
- Layer (neural network), 167
- LCN, see local contrast normalization
- Leaky ReLU, 192
- Leaky units, 409
- Learning rate, 85
- Line search, 85, 86, 93
- Linear combination, 37
- Linear dependence, 38
- Linear factor models, 492
- Linear regression, 107, 110, 140
- Link prediction, 487
- Lipschitz constant, 92
- Lipschitz continuous, 92
- Liquid state machine, 406

- Local conditional probability distribution, 567
- Local contrast normalization, 459
- Logistic regression, 3, 140, 140
- Logistic sigmoid, 7, 67
- Long short-term memory, 18, 25, 306, 411, 428
- Loop, 582
- Loopy belief propagation, 588
- Loss function, see objective function
- L^p norm, 39
- LSTM, see long short-term memory
- M-step, 637
- Machine learning, 2
- Machine translation, 101
- Main diagonal, 33
- Manifold, 160
- Manifold hypothesis, 161
- Manifold learning, 161
- Manifold tangent classifier, 270
- MAP approximation, 138, 508
- Marginal probability, 58
- Markov chain, 598
- Markov chain Monte Carlo, 598
- Markov network, see undirected model
- Markov random field, see undirected model
- Matrix, xi, xii, 32
- Matrix inverse, 36
- Matrix product, 34
- Max norm, 40
- Max pooling, 340
- Maximum likelihood, 131
- Maxout, 192, 428
- MCMC, see Markov chain Monte Carlo
- Mean field, 641, 642, 674
- Mean squared error, 108
- Measure theory, 71
- Measure zero, 71
- Memory network, 419, 421
- Method of steepest descent, see gradient descent
- Minibatch, 279
- Missing inputs, 100
- Mixing (Markov chain), 604
- Mixture density networks, 188
- Mixture distribution, 66
- Mixture model, 188, 513
- Mixture of experts, 453, 551
- MLP, see multilayer perception
- MNIST, 21, 22, 674
- Model averaging, 255
- Model compression, 451
- Model identifiability, 284
- Model parallelism, 450
- Moment matching, 705
- Moore-Penrose pseudoinverse, 45, 240
- Moralized graph, 580
- MP-DBM, see multi-prediction DBM
- MRF (Markov Random Field), see undirected model
- MSE, see mean squared error
- Multi-modal learning, 542
- Multi-prediction DBM, 676
- Multi-task learning, 245, 541
- Multilayer perception, 5
- Multilayer perceptron, 27
- Multinomial distribution, 62
- Multinoulli distribution, 62
- n-gram, 464
- NADE, 710
- Naive Bayes, 3
- Nat, 73
- Natural image, 562
- Natural language processing, 464
- Nearest neighbor regression, 115
- Negative definite, 89
- Negative phase, 473, 609, 611
- Neocognitron, 16, 24, 27, 368
- Nesterov momentum, 300
- Netflix Grand Prize, 256, 482
- Neural language model, 466, 479
- Neural network, 13
- Neural Turing machine, 421
- Neuroscience, 15
- Newton's method, 89, 310
- NLM, see neural language model
- NLP, see natural language processing
- No free lunch theorem, 116

- Noise-contrastive estimation, 623
- Non-parametric model, 114
- Norm, xiv, 39
- Normal distribution, 63, 64, 125
- Normal equations, 109, 109, 112, 234
- Normalized initialization, 303
- Numerical differentiation, see finite differences

- Object detection, 456
- Object recognition, 456
- Objective function, 82
- OMP-k, see orthogonal matching pursuit
- One-shot learning, 541
- Operation, 204
- Optimization, 80, 82
- Orthodox statistics, see frequentist statistics
- Orthogonal matching pursuit, 27, 254
- Orthogonal matrix, 42
- Orthogonality, 41
- Output layer, 167

- Parallel distributed processing, 17
- Parameter initialization, 301, 408
- Parameter sharing, 251, 336, 374, 376, 389
- Parameter tying, see Parameter sharing
- Parametric model, 114
- Parametric ReLU, 192
- Partial derivative, 84
- Partition function, 571, 608, 671
- PCA, see principal components analysis
- PCD, see stochastic maximum likelihood
- Perceptron, 15, 27
- Persistent contrastive divergence, see stochastic maximum likelihood
- Perturbation analysis, see reparametrization trick
- Point estimator, 122
- Policy, 483
- Pooling, 331, 685
- Positive definite, 89
- Positive phase, 473, 609, 611, 658, 670
- Precision, 426
- Precision (of a normal distribution), 63, 65
- Predictive sparse decomposition, 526

- Preprocessing, 456
- Pretraining, 324, 531
- Primary visual cortex, 366
- Principal components analysis, 48, 146–148, 493, 634
- Prior probability distribution, 135
- Probabilistic max pooling, 685
- Probabilistic PCA, 493, 494, 635
- Probability density function, 58
- Probability distribution, 56
- Probability mass function, 56
- Probability mass function estimation, 103
- Product of experts, 573
- Product rule of probability, see chain rule of probability
- PSD, see predictive sparse decomposition
- Pseudolikelihood, 618

- Quadrature pair, 370
- Quasi-Newton condition, 316
- Quasi-Newton methods, 316

- Radial basis function, 196
- Random search, 437
- Random variable, 56
- Ratio matching, 621
- RBF, 196
- RBM, see restricted Boltzmann machine
- Recall, 426
- Receptive field, 338
- Recommender Systems, 481
- Rectified linear unit, 171, 192, 428, 510
- Recurrent network, 27
- Recurrent neural network, 379
- Regression, 101
- Regularization, 120, 120, 177, 228, 433
- Regularizer, 119
- REINFORCE, 691
- Reinforcement learning, 25, 106, 483, 691
- Relational database, 486
- Reparametrization trick, 690
- Representation learning, 3
- Representational capacity, 114
- Restricted Boltzmann machine, 357, 462, 482, 590, 634, 658, 659, 674, 678,

- 680, 683, 685
- Ridge regression, see weight decay
- Risk, 275
- RNN-RBM, 688
- Saddle points, 285
- Sample mean, 125
- Scalar, xi, xii, 31
- Score matching, 516, 620
- Secant condition, 316
- Second derivative, 86
- Second derivative test, 89
- Self-information, 73
- Semantic hashing, 528
- Semi-supervised learning, 244
- Separable convolution, 363
- Separation (probabilistic modeling), 575
- Set, xii
- SGD, see stochastic gradient descent
- Shannon entropy, xiii, 73
- Shortlist, 469
- Sigmoid, xiv, see logistic sigmoid
- Sigmoid belief network, 27
- Simple cell, 366
- Singular value, see singular value decomposition
- Singular value decomposition, 44, 148, 482
- Singular vector, see singular value decomposition
- Slow feature analysis, 496
- SML, see stochastic maximum likelihood
- Softmax, 183, 421, 453
- Softplus, xiv, 68, 196
- Spam detection, 3
- Sparse coding, 322, 357, 499, 634, 694
- Sparse initialization, 304, 408
- Sparse representation, 146, 226, 253, 508, 559
- Spearman, 439
- Spectral radius, 407
- Speech recognition, see automatic speech recognition
- Sphering, see whitening
- Spike and slab restricted Boltzmann machine, 683
- SPN, see sum-product network
- Square matrix, 38
- ssRBM, see spike and slab restricted Boltzmann machine
- Standard deviation, 61
- Standard error, 127
- Standard error of the mean, 128, 278
- Statistic, 122
- Statistical learning theory, 110
- Steepest descent, see gradient descent
- Stochastic back-propagation, see reparametrization trick
- Stochastic gradient descent, 15, 150, 279, 294, 674
- Stochastic maximum likelihood, 615, 674
- Stochastic pooling, 265
- Structure learning, 585
- Structured output, 101, 687
- Structured probabilistic model, 77, 561
- Sum rule of probability, 58
- Sum-product network, 556
- Supervised fine-tuning, 532, 664
- Supervised learning, 105
- Support vector machine, 140
- Surrogate loss function, 276
- SVD, see singular value decomposition
- Symmetric matrix, 41, 43
- Tangent distance, 269
- Tangent plane, 519
- Tangent prop, 269
- TDNN, see time-delay neural network
- Teacher forcing, 383, 384
- Tempering, 606
- Template matching, 141
- Tensor, xi, xii, 33
- Test set, 110
- Tikhonov regularization, see weight decay
- Tiled convolution, 353
- Time-delay neural network, 369, 375
- Toeplitz matrix, 334
- Topographic ICA, 496
- Trace operator, 46
- Training error, 110
- Transcription, 101

INDEX

- Transfer learning, 539
- Transpose, xii, 33
- Triangle inequality, 39
- Triangulated graph, see chordal graph
- Trigram, 465

- Unbiased, 124
- Undirected graphical model, 77, 510
- Undirected model, 569
- Uniform distribution, 57
- Unigram, 465
- Unit norm, 41
- Unit vector, 41
- Universal approximation theorem, 197
- Universal approximator, 556
- Unnormalized probability distribution, 570
- Unsupervised learning, 105, 146
- Unsupervised pretraining, 462, 531

- V-structure, see explaining away
- V1, 366
- VAE, see variational autoencoder
- Vapnik-Chervonenkis dimension, 114
- Variance, xiii, 61, 229
- Variational autoencoder, 691, 698
- Variational derivatives, see functional derivatives
- Variational free energy, see evidence lower bound
- VC dimension, see Vapnik-Chervonenkis dimension
- Vector, xi, xii, 32
- Virtual adversarial examples, 268
- Visible layer, 6
- Volumetric data, 361

- Wake-sleep, 654, 663
- Weight decay, 118, 177, 231, 434
- Weight space symmetry, 284
- Weights, 15, 107
- Whitening, 458
- Wikibase, 486
- Wikibase, 486
- Word embedding, 467
- Word-sense disambiguation, 487

- WordNet, 486
- Zero-data learning, see zero-shot learning
- Zero-shot learning, 541